



SQL & PL/SQL

Your First Step into Databases



Vaibhav Rajasha

Table of Contents:

Part I – SQL Fundamentals

1. Introduction to Data and Databases
 - Data, Database, DBMS, RDBMS
 - What is SQL
 - Schema
2. Variables and Data Types
 - CHAR, VARCHAR2, NUMBER, DATE, BOOLEAN, NULL
3. Operators
 - Arithmetic Operators
 - Comparison Operators
 - Logical Operators
 - Set Operators
 - Special Operators (LIKE, BETWEEN, IN, EXISTS, NOT IN, IS NULL, IS NOT NULL)
4. Dual Table
 - Purpose of DUAL Table in Oracle
5. Comments and Aliases
 - SQL Comments
 - Column and Table Aliases

Part II – SQL Constraints & Commands

6. Constraints and Keys
 - NOT NULL, UNIQUE, CHECK, DEFAULT
 - PRIMARY KEY, FOREIGN KEY
 - Composite Keys
 - Super, Candidate, Alternate Keys
7. Data Definition Language (DDL)
 - CREATE, ALTER, DROP, TRUNCATE, RENAME
8. Data Manipulation Language (DML)
 - INSERT, UPDATE, DELETE
9. Data Query Language (DQL)
 - SELECT, WHERE, ORDER BY, GROUP BY, HAVING, DISTINCT

Part III – Transactions, Security & Querying

10. Transaction Control Language (TCL)
 - COMMIT, ROLLBACK, SAVEPOINT, SET TRANSACTION

11. Data Control Language (DCL)

- GRANT, REVOKE

12. Joins

- INNER, LEFT, RIGHT, FULL, CROSS, SELF JOIN

13. Set Operators

- UNION, UNION ALL, INTERSECT, MINUS

14. Subqueries

- Nested, Correlated, Multi-row
- Subqueries with EXISTS, IN

15. SQL Functions

- Aggregate, Scalar, Conversion, String, Numeric, Date, Conditional
- Windows, System, JSON, XML
- CASE, COALESCE, DECODE, NULLIF

16. Views and Temporary Tables

- Types of Views
- WITH CHECK OPTION

Part IV – Functions & Advanced SQL

17. Temporary Table (Advanced SQL concepts)

- CTE and WITH Clause
- Recursive CTE
- PIVOT and UNPIVOT
- MERGE Statement (UPsert)
- Regular Expressions (REGEXP)

18. Indexes

- Normal and Composite Indexes

19. Optimization Concepts

- Query Execution Plan
- Indexing Tips
- Performance Tuning Basics

20. Data Export and Import (Basics)

- SPOOL, SQL*Loader, INSERT INTO Techniques

21. Real-Time SQL Use Cases

- ETL Queries
- Report Generation
- Data Validation

22. OLTP vs OLAP

Part V – PL/SQL Programming

23. Introduction to PL/SQL

- What is PL/SQL
- SQL vs PL/SQL
- Anonymous & Named Blocks
- PL/SQL Syntax

24. Variables and Data Types

- Constants, Scalar Types
- Anchored Types (%TYPE, %ROWTYPE)
- Constants and Scope

25. Control Structures

- IF, CASE
- LOOP, FOR, WHILE, Nested Loops
- Operators

26. Procedures and Functions

- IN, OUT, IN OUT Parameters
- Subprograms
- RETURN Values

27. Cursors

- Implicit / Explicit
- Cursor FOR Loop
- Parameterized / REF Cursors

28. Exception Handling

- Predefined Exceptions
- User-defined Exceptions
- RAISE / RAISE_APPLICATION_ERROR

29. Packages

- Package Specification / Body
- Public vs Private Members

30. Triggers

- BEFORE / AFTER / INSTEAD OF
- Row-Level / Statement-Level
- Compound Triggers

31. Collections

- RECORD, Nested Table, VARRAY
- Associative Arrays

32. Dynamic SQL

- EXECUTE IMMEDIATE, DBMS_SQL

33. Bulk Operations

- BULK COLLECT, FORALL, SAVE EXCEPTIONS

34. Advanced PL/SQL Topics

- UTL_FILE, DBMS_SCHEDULER, DBMS_JOB
- Autonomous Transactions, Sequences, Synonyms
- Object Types

35. Debugging and Logging

- DBMS_OUTPUT
- SQL Developer Debugger
- Logging Techniques

36. Real-Time PL/SQL Applications

- Reusable Packages
- Backend Jobs
- Complex Business Logic
- Project Case Studies

Additional Concepts

37. Data Modeling Concepts

- Normalization
- ER Diagrams

38. Performance Tuning Best Practices

39. Metadata Tables

- USER_TABLES, ALL_OBJECTS, DBA Views

40. PL/SQL Code Reusability Best Practices

SQL

◆ Introduction to Data :

◆ What is Data?

- Data means raw facts or figures.
💡 Examples: Names, phone numbers, sales records.
- In SQL/database terms, data refers to information stored in a structured format, usually in tables.
- A table is made up of:
 - Rows (also called records)
 - Columns (each holds a specific type of data)

💡 Example Table :

ID	Name	Age	Department
1	Raj	28	HR
2	Simran	30	IT

- Properly stored and organized data is valuable for analysis and decision-making.

◆ Difference Between Data and Information :

- Data = Raw, unprocessed facts.
💡 Example : 85, Raj, IT
- Information = Processed and meaningful data.
💡 Example : Raj is 28 years old and works in the HR department.

◆ What is a Database?

- A database is a place where data is stored in an organized way so it can be easily accessed, managed and updated.
- It helps reduce paperwork and store large amounts of data digitally.

💡 Example : A phone contact list or a student record system.

💻 Syntax to Create a Database :

`CREATE DATABASE database_name;`

💡 Example :

`CREATE DATABASE StudentDB;`

◆ Why do we need a Database?

- To store data securely without paper files.
- To reduce duplication, errors, and data loss.
- To easily retrieve and analyze data whenever needed.

◆ Types of Databases :

- Relational Database (RDBMS) :
 - Stores data in tables (rows and columns).
 - Uses SQL to manage data.

💡 Example : MySQL, Oracle, PostgreSQL.
- NoSQL (Non-relational Database) :
 - Stores data in non-tabular formats: key-value pairs, documents, graphs.
 - Used for big data and real-time apps.

💡 Example : MongoDB, Cassandra, Redis.
- Hierarchical Database :
 - Stores data in a tree-like structure (like folders and subfolders).
 - Used in old systems.

💡 Example : IBM IMS.
- Network Database :
 - Similar to hierarchical, but allows multiple parent-child relationships.
 - More flexible but complex.

💡 Example : Integrated Data Store (IDS), IDMS (Integrated Database Management System) by CA Technologies.
- Object-Oriented Database :
 - Stores data as objects, just like in object-oriented programming.

💡 Example : db4o, ObjectDB.

◆ Advantages of Using Databases :

- Stores large amounts of data easily
- Prevents duplicate data (reduces redundancy)
- Allows fast search, updates, and reporting
- Provides security and access control
- Supports multi-user access

◆ Disadvantages of Databases :

- Requires software and hardware resources
- Needs skilled people to manage (DBA, developers)
- Complex to set up and design
- If not managed well, can cause data corruption or loss

❖ Key Database Terms (Must-Know) :

- ◆ What is a Database Object?
 - Any item created in the database is called a database object.
 - 💡 Examples: Table, View, Index, Sequence, Procedure, etc.
- ◆ What are Schema Objects?
 - These are objects inside a schema (a user's area in the database).
 - 💡 Example: Tables, Views, Indexes, Triggers, Synonyms, etc.
 - 📌 *Think of a schema as a folder, and objects as files in that folder.*
- ◆ Table :
 - Stores data in rows and columns.
 - Every table has a name and a set of columns with data types.
- ◆ View :
 - A virtual table based on a SELECT query.
 - Doesn't store data, only displays it from other tables.
- ◆ Index :
 - Speeds up data search (like a book index).
 - Improves performance while searching records.
- ◆ Sequence :
 - Generates a series of numbers, often used to create unique IDs (like auto-increment).
- ◆ Synonym :
 - A shortcut name for another object.
 - Useful to simplify access to tables or views from other schemas.
- ◆ Stored Procedure :
 - A saved set of SQL commands.
 - Helps to automate repeated tasks.
- ◆ Trigger :
 - A set of instructions that runs automatically when an event happens in a table (like insert, update, delete).
- ◆ Cursor :
 - Used to go through each row one by one.
 - Mostly used in stored procedures when row-by-row operations are needed.

📌 *Note : These topics are explained in detail in the upcoming chapters*

◆ What is DBMS (Database Management System)?

- DBMS is software used to store, manage, and retrieve data in a database.
 - It makes sure the data is safe, consistent, and easily accessible.
- 💡 Examples : MySQL, Oracle, SQL Server, PostgreSQL.
-

◆ What is an RDBMS (Relational Database Management System)?

- RDBMS stores data in tables with rows and columns.
 - Each table has a unique identifier (Primary Key).
 - RDBMS supports relationships between tables (like links between student and marks tables).
- 💡 Example : MySQL, Oracle, PostgreSQL.
-

◆ What is SQL (Structured Query Language)?

- It is used to interact with databases – to insert, update, delete, or retrieve data.
 - SQL is a standard language supported by all major Relational Database Management Systems (RDBMS) like Oracle, MySQL, SQL Server, PostgreSQL, etc.
 - SQL is both human-readable and machine-readable, making it easy to learn and powerful to use.
 - It works on structured data (organized in rows and columns using tables).
- ◊ How does SQL help?
- You write SQL commands (queries) to talk to the database.
 - SQL helps in :
 - Creating database objects like tables, views, indexes, procedures, etc.
 - Inserting new records into tables.
 - Fetching reports or data using SELECT queries.
 - Filtering, sorting, and grouping data to get meaningful insights.
 - Updating or deleting old/unwanted data.
 - Giving or restricting access using privileges and roles.
 - Supporting data integrity, consistency, and security in databases.
-

◆ Schema in a Database

- A schema is like a folder in the database that contains related objects like tables, views, indexes, and procedures.
- It helps organize data and control access (who can see or use what).
- Each user in a database usually has their own default schema.

❖ Think of a schema as :

👉 A *library (database) and library section* (schema) that contains *books* (tables, views, etc.).

- ◆ SQL Command to Create a Schema (in databases that support it like PostgreSQL, SQL Server, Oracle):

💻 Syntax :

```
CREATE SCHEMA schema_name;
```

- Then, you can create tables inside the schema like this :

```
CREATE TABLE schema_name.table_name (Col_name1 data_type, Col_name2  
data_type,...n);
```

💡 Example :

```
CREATE SCHEMA hr;
```

```
CREATE TABLE hr.employees (id INT, name VARCHAR(50));
```

❖ In MySQL, schemas and databases are the same thing.

❖ But in Oracle, SQL Server, PostgreSQL, schemas are like folders inside a database suppose database is the main container.

◆ SQL Variables

- Variables are temporary storage locations that store values while a program or query runs.
- Variables are NOT supported in standard SQL (like the basic SQL you use in MySQL, PostgreSQL, etc.) for general queries.
- SQL is a *declarative language*, mainly used for querying and modifying data — not for procedural logic like loops, conditions, or variables.
- Variables ARE supported — especially in procedural extensions like:
 - PL/SQL (Oracle)
 - T-SQL (SQL Server)
 - Stored Procedures / Functions in MySQL/PostgreSQL

❖ In standard SQL, you write queries like:

```
SELECT * FROM employees WHERE department_id = 10;
```

→ No variable used here.

- ◆ Why use variables?

- To hold intermediate values
- To reuse values in multiple statements
- To make procedures or functions dynamic

❖ You use variables to store temporary data, control flow, loops, conditions, etc.

◆ What is a Bind Variable?

- A bind variable is a placeholder in an SQL query where you don't hard-code the value but instead pass the value later when the query is executed.
- Best Practice → Use meaningful names so it's easy to understand later.
- It does NOT have to match the column name, but should clearly describe the value.
- → :id, :param1, :userInput, etc.
- It helps in:
 - Preventing SQL Injection
 - Reusing the same SQL query with different values
 - Improving performance (by reducing parsing overhead)

💡 Example : Without Bind Variable (Bad Practice)

```
SELECT *
FROM employees
WHERE employee_id = 101;
```

- 👉 Here, 101 is hard-coded.
- Problem: Every time the value changes, the query needs to be parsed again → Performance hits.

💡 Example : With Bind Variable (Good Practice)

```
SELECT *
FROM employees
WHERE employee_id = :emp_id;
```

- 👉 :emp_id is a bind variable (placeholder).
- At runtime, you can pass different values (like 101, 102, 103, etc.) without changing the query itself.

◆ Why use :empid?

- Because in applications (like Oracle Forms, PL/SQL blocks, or other front-end programs) you don't hardcode the employee number.
- Instead, you pass it dynamically.
- For example, if you want to show department history for employee 101, you pass:
:empid = 101
- If you want to show for employee 205, you pass:
:empid = 205
- ✅ So the query becomes:

```
SELECT e.emp_no,
       d.dept_name,
       de.from_date,
       de.to_date
  FROM department d
 JOIN depet_no de ON d.dept_no = de.dept_no
 JOIN employee e ON e.emp_no = de.emp_no
 WHERE e.emp_no = :empid;
```

- When you run this in an application, the system asks you to enter the value for empid, or your program will supply it automatically.

◆ SQL Data Types

- SQL Data Types define what kind of data a column can store, like numbers, text, or dates.
- They help in :
 - Data Accuracy (right type of data)
 - Efficient Storage (use minimum space)
 - Better Performance (faster queries)

💡 Think of it like:

Data types are rules that say:

“This column will store numbers only” or “This column will store text only.”

◦ SQL Data Types :

◦ 1. Numeric Data Types :

- Numeric data types are used to store numbers like whole numbers and decimals.
- They support math operations (+, -, ×, ÷, %), so they are useful in finance, science, and data analysis.

Data Type	Description	Supported Databases
INT / INTEGER	Stores whole numbers	Oracle, MySQL, SQL Server
NUMBER(p,s)	General-purpose numeric type with precision	Oracle
DECIMAL(p,s)	Exact decimal values (finance)	MySQL, SQL Server, Oracle
FLOAT	Approximate numeric values	MySQL, SQL Server, Oracle

◦ 2. Character and String Data Types :

- Character data types are used to store text or character-based data.
- The choice between fixed-length and variable-length data types depends on the nature of your data.

Data Type	Description	Supported Database
CHAR(n)	Fixed-length character string	Oracle, MySQL, SQL Server
VARCHAR(n)	Variable-length string	MySQL, SQL Server
VARCHAR2(n)	Oracle-specific variable-length string	Oracle
Text	Large text data	MySQL, SQL Server <i>(deprecated)</i>

◊ 3. Date and Time Data Type :

- SQL provides several data types for storing date and time information.
- They are essential for managing timestamps, events, and time-based queries. These are given in the below table.

Data Type	Description	Supported Databases
DATE	stores the data of date (year, month, day)	Oracle, MySQL, SQL Server
TIMESTAMP	Date & time with high precision	Oracle, MySQL, SQL Server
DATETIME	store both the data and time (year, month, day, hour, minute, second)	MySQL, SQL Server

◊ 4. Binary Data Types in SQL :

- Binary data types are used to store binary data such as images, videos, or other file types. These include:

Data Type	Description	Max Length
Binary	Fixed-length binary data.	8000 bytes
VarBinary	Variable-length binary data.	8000 bytes
Image	Stores binary data as images.	2,147,483,647 bytes

◊ 5. Boolean Data Type in SQL :

- The BOOLEAN data types are used to store logical values, typically TRUE or FALSE. It's commonly used for flag fields or binary conditions.

Data Type	Description	Supported Databases
BOOLEAN	Stores a logical value (TRUE/FALSE).	MySQL, PL/SQL only

💡 Example of Oracle SQL Data Types :

```
CREATE TABLE Employees(
    EmpID NUMBER,          -- Numeric data type (whole number)
    FirstName VARCHAR2(50), -- Variable-length character string
    LastName VARCHAR2(50), -- Variable-length character string
    Salary NUMBER(10, 2),   -- For precise numeric values like salary
    JoinDate DATE,         -- Date datatype (includes time by default in Oracle)
    IsActive CHAR(1)        -- Used for Boolean-like flag ('Y' or 'N')
);
```

◆ SQL Operators

- SQL Operators are special symbols or keywords used to perform operations on data values in SQL queries.
- They are mostly used in **WHERE**, **HAVING**, and **SELECT** clauses to filter, compare, and manipulate data.

◆ Types of SQL Operators :

- Arithmetic Operator
- Comparison Operator
- Logical Operator
- Set Operators
- Special Operators

◇ 1. Arithmetic Operators :

- Used to perform basic mathematical operations and return a numeric result.

Operator	Description	Example
+	Addition	SELECT 5 + 2 FROM DUAL; → 7
-	Subtraction	SELECT 10 - 4 FROM DUAL; → 6
*	Multiplication	SELECT 3 * 3 FROM DUAL; → 9
/	Division	SELECT 10 / 2 FROM DUAL; → 5
% (<i>Modulus</i>)	Remainder (some RDBMS only)	SELECT MOD(10, 3) FROM DUAL; → 1

◊ 2. Comparison Operators :

- Used to compare two values. Mostly used in **WHERE** clause.
- Returns a Boolean result (TRUE/FALSE) internally for each row.
- **!=** and **<>** both mean "not equal". Use based on DBMS compatibility.

Operator	Description	Example
=	Equal to	WHERE salary = 50000
!= or <>	Not equal to	WHERE city != 'Pune'
>	Greater than	WHERE age > 25
<	Less than	WHERE salary < 60000
>=	Greater than or equal	WHERE marks >= 35
<=	Less than or equal	WHERE age <= 30

◊ 3. Logical Operators :

- Used to combine multiple conditions in the **WHERE** clause.

Operator	Description	Example
AND	Returns TRUE if all conditions are true	WHERE age > 25 AND city = 'Pune'
OR	Returns TRUE if at least one condition is true	WHERE dept = 'HR' OR dept = 'IT'
NOT	Reverses the condition	WHERE NOT city = 'Mumbai'

◊ 4. Set Operators :

- Used to combine results from two or more SELECT queries.
- These require both queries to have the same number of columns and compatible data types.

Operator	Description	Usage
UNION	Combines unique rows from both queries	SELECT city FROM A UNION SELECT city FROM B
UNION ALL	Includes duplicates too	SELECT city FROM A UNION ALL SELECT city FROM B
INTERSECT	Returns common rows	SELECT city FROM A INTERSECT SELECT city FROM B
MINUS	Returns rows from first query not in second	SELECT city FROM A MINUS SELECT city FROM B

- ◊ 5. Special Operators :
 - These are helpful for advanced filtering in WHERE clauses.

Operator	Description	Example
BETWEEN	Checks within a range	WHERE age BETWEEN 18 AND 30
IN	Checks if value exists in list	WHERE city IN ('Pune', 'Mumbai')
NOT IN	Checks if value doesn't exist	WHERE dept NOT IN ('HR', 'IT')
LIKE	Pattern matching (with %, _)	WHERE name LIKE 'V%'
IS NULL	Checks for NULL	WHERE salary IS NULL
IS NOT NULL	Checks for non-NUL	WHERE city IS NOT NULL
EXISTS	Checks if subquery returns rows	WHERE EXISTS (SELECT * FROM dept WHERE id = 10)

💡 Example :

```

SELECT
    employee_id,
    first_name || ' ' || last_name AS full_name,
    salary,
    salary * 12 AS annual_salary,
    Department_id
FROM employees
WHERE (salary BETWEEN 30000 AND 80000)
    AND department_id IN (10, 20, 30)
    AND last_name LIKE 'S%'
    AND manager_id IS NOT NULL
    AND salary + 5000 > 40000;

```

◆ What is the **DUAL** Table in Oracle?

- The **DUAL** table is a special, built-in table in Oracle Database that is used to perform calculations or return a value without querying any real table.

◆ Key Points About **DUAL** Table :

- ◆ 1. Single Row and Single Column :
 - The **DUAL** table has only one row and one column called **DUMMY**.
 - This column contains a single value: '**X**'.
- ◆ Example :

```
SELECT * FROM DUAL;
```

-- Output: X
- ◆ 2. Used for Expressions and Functions :
 - It's commonly used when you want to evaluate an expression, function, or calculation without using any actual table.
- ◆ Example :

```
SELECT 10 + 5 FROM DUAL;
```

-- Output: 15
- ◆ 3. System Table :
 - **DUAL** is part of the data dictionary, created by Oracle.
 - You don't need to create or manage it. It's always available.
- ◆ 4. Use with Functions
 - You can test SQL functions easily using **DUAL**.
- ◆ Example :

```
SELECT SYSDATE FROM DUAL;
```

-- Returns current system date
- ◆ 5. Use with Pseudo Columns :
 - You can also use pseudo-columns like **SYSDATE**, **USER**, **ROWNUM** with **DUAL**.
- ◆ Example :

```
SELECT USER FROM DUAL;
```

-- Returns current username
- ◆ 6. Why Not Use Real Tables?
 - Using **DUAL** avoids unnecessary reads from real data tables when:
 - You just want to display a constant
 - You want to test a formula or function
 - You are writing a simple script for checking syntax or results
- ◆ 7. Other Databases :
 - **DUAL** is specific to Oracle.
 - In MySQL, you can write **SELECT 5+2;** without **DUAL**.
 - In SQL Server, use **SELECT 5+2;** as well.
- ◆ 8. Can You Insert or Modify **DUAL**?
 - No. You should not modify the **DUAL** table.
 - Oracle maintains it, and any change can break system behavior.

💡 Examples Using DUAL :

Purpose	Query
Simple arithmetic	<code>SELECT 100 * 2 FROM DUAL;</code>
Current system date	<code>SELECT SYSDATE FROM DUAL;</code>
String manipulation	<code>SELECT INITCAP('vaibhav') FROM DUAL;</code>
Rounding numbers	<code>SELECT ROUND(10.567, 2) FROM DUAL;</code>
Concatenation	<code>'SELECT 'Vaibhav'</code>

🔔 Real-World Scenario :

- Imagine you're testing logic in a stored procedure or PL/SQL script. Before applying on real data :

💡 Example :

-- Check discount formula
`SELECT 1000 - (1000 * 10 / 100) AS DiscountedPrice FROM DUAL;`
→ 900

◆ SQL COMMENTS:

❖ What are SQL Comments?

- SQL Comments are non-executable statements used to explain and document your SQL code.
- They help make your code more readable and easier to understand — especially when revisiting later or sharing with others.

💻 Types of Comments in SQL:

- SQL supports two types of comments:

Type	Syntax	Description
❖ 1. Single-line comment	-- comment	Comment that spans one line
❖ 2. Multi-line comment	/* comment block */	Comment that can span multiple lines

💡 Examples:

💻 Single-line Comment:

-- This query fetches all employees
`SELECT * FROM employees;`

💻 Multi-line Comment:

/* This query fetches all employees
who have a salary more than 50000 */
`SELECT * FROM employees
WHERE salary > 50000;`

📝 Important Notes:

- Comments are ignored by the SQL engine — they do not affect query results.
- Use comments to explain:
 - Complex logic
 - Temporary conditions
 - Testing sections
- Over-commenting can make code messy. Use wisely!

◆ SQL ALIASES :

❖ What is an Alias?

- An alias is a temporary nickname given to a column or table in a SQL query to make the output more readable.
- It helps you:
 - Shorten long table/column names
 - Format output headers
 - Improve code readability

💻 Syntax of Column Alias

```
SELECT column_name AS alias_name  
FROM table_name;
```

📝 You can also omit AS keyword:

```
SELECT column_name alias_name  
FROM table_name;
```

💡 Example – Column Alias

```
SELECT first_name AS Name, salary AS Income  
FROM employees;
```

💻 Output:

Name		Income
John		50000

💻 Syntax of Table Alias

```
SELECT t.column_name  
FROM table_name AS t;
```

📝 OR (without AS):

```
SELECT t.column_name  
FROM table_name t;
```

💡 Example – Table Alias

```
SELECT e.first_name, d.department_name  
FROM employees e  
JOIN departments d ON e.department_id = d.department_id;
```

- ◆ Here, **e** and **d** are aliases for **employees** and **departments**.

📝 Important Notes:

- Aliases are temporary — they exist only during that query execution.
- Aliases are helpful in:
 - Joins (when multiple tables used)
 - Aggregations (when using functions like COUNT, AVG)
 - Making column headers user-friendly
- You cannot use aliases in WHERE clause — because WHERE is evaluated before SELECT.

Summary Table:

Feature	Purpose	Used On	Syntax Style
Comments	Make code readable	Any SQL section	-- or /* ... */
Aliases	Rename temporarily (for clarity)	Columns, Tables	AS alias or just alias

SQL Constraints

- Constraints are rules applied on columns in a table to enforce data integrity (i.e., make sure only valid data is stored in the database).
- They automatically prevent invalid data from being inserted, updated, or deleted.
- Constraints can be added during table creation or later using **ALTER TABLE**.
- If constraint is violated during **INSERT** or **UPDATE**, the SQL query fails with an error.
- All constraints can be named for better management.
CONSTRAINT unique_email UNIQUE (email)

Syntax for View Existing Constraints :

```
SELECT constraint_name, constraint_type, table_name  
FROM user_constraints  
WHERE table_name = 'YOUR_TABLE_NAME';
```

 You only need to write the table name, but here are important things to remember:

- Table name should be in UPPERCASE
- Don't include schema name

Example How to View Constraints :

-- For all constraints

```
SELECT * FROM user_constraints WHERE table_name = 'EMPLOYEES';
```

-- For columns with constraints

```
SELECT * FROM user_cons_columns WHERE table_name = 'EMPLOYEES';
```

Types of SQL Constraints :

Constraint	Description
NOT NULL	Ensures that a column cannot have a NULL value and It can only be applied at the column level, not at the table level.
UNIQUE	Ensures all values in a column are unique (no duplicates).
PRIMARY KEY	Uniquely identifies each row. Cannot be NULL + must be unique.
FOREIGN KEY	Connect two tables. Maintains referential integrity between tables.
CHECK	Limits the values based on a condition.
DEFAULT	Sets a default value if no value is given for that column.

◊ 1. **NOT NULL** Constraint :

- The **NOT NULL** constraint ensures that a column cannot have a **NONE** value.
- It guarantees that every row must contain a valid (non-empty) value in this column.

💡 Example :

```
CREATE TABLE employees(
    emp_id INT NOT NULL,
    name VARCHAR(50) NOT NULL
);
```

◊ 2. **UNIQUE** Constraint :

- Ensures values in the column are all different / Unique.
- Can be used at both column and table level.
- Multiple **UNIQUE** constraints allowed in a table.
- **NONE** is allowed once per column with **UNIQUE**.

💡 Example :

```
CREATE TABLE students(
    student_id INT UNIQUE,
    email VARCHAR(100) UNIQUE
);
```

◊ 3. **PRIMARY KEY** Constraint :

- Uniquely identifies each record in a table. Combines **NOT NULL** + **UNIQUE**.
- Only one **PRIMARY KEY** allowed per table.
- Can span multiple columns → called Composite Key.

💡 Example :

```
CREATE TABLE customers(
    cust_id INT PRIMARY KEY,
    name VARCHAR(100)
);
```

📝 You can also use composite keys :

💻 Syntax :

```
PRIMARY KEY (col1, col2)
```

💡 Composite PK Example :

```
CREATE TABLE project_assignment (
    emp_id INT,
    project_id INT,
    PRIMARY KEY (emp_id, project_id)
);
```

◊ 4. **FOREIGN KEY** Constraint :

- Links two tables together. The foreign key refers to the primary key of another table.
- Enforces referential integrity.
- You cannot insert a value unless it exists in the referenced table.
- You cannot delete a parent row if children exist, unless **ON DELETE CASCADE** is used.

💡 Example :

```
CREATE TABLE orders (
    order_id INT PRIMARY KEY,
    cust_id INT,
```

```
    FOREIGN KEY (cust_id) REFERENCES customers(cust_id)
);
```

☞ Prevents deleting a customer from `customers` if orders exist in `orders`.

❑ Basic FK Syntax :

```
FOREIGN KEY (dept_id) REFERENCES departments(dept_id)
```

💡 Example :

```
FOREIGN KEY (cust_id) REFERENCES customers(cust_id)
ON DELETE CASCADE
```

☞ `ON DELETE CASCADE` → deletes child rows when parent is deleted.

☞ `ON DELETE SET NULL` → sets child FK to NULL when parent is deleted.

◊ 5. **CHECK** Constraint :

- Validates values using a condition.
- Used to limit column values based on logic.
- Can be applied at column or table level.

💡 Example :

```
CREATE TABLE products (
product_id INT PRIMARY KEY,
price DECIMAL(10, 2) CHECK (price > 0)
);
```

- Supports multiple conditions.

💡 Example :

```
CHECK (salary > 0 AND salary < 100000)
```

👉 Some databases (like older MySQL versions) may ignore CHECK constraints unless explicitly enabled.

◊ 6. **DEFAULT** Constraint :

- Sets a default value if none is provided.
- Automatically applies a value if the user does not provide one.

💡 Example :

```
CREATE TABLE accounts (
acc_id INT PRIMARY KEY,
status VARCHAR(10) DEFAULT 'Active'
);
```

❑ Summary Table : CONSTRAINTS

Constraint	Description	Can be applied on	Default behavior
NOT NULL	The column must have a value. Can't be left empty.	Column only	Cannot insert or update with NULL
UNIQUE	Each value in the column (or set of columns) must be different.	Column/Table	Allows one NULL (Oracle)
PRIMARY KEY	Unique + Not Null. Only one per table. Used to uniquely identify rows.	Column/Table	Cannot contain NULL, always unique

FOREIGN KEY	Links one table's column to another table's primary key. Enforces referential integrity.	Column/Table	Rejects insert/update if reference not found
CHECK	Validates data against a logical condition.	Column/Table	Rejects rows failing the condition
DEFAULT	Assigns a default value when no value is provided.	Column only	Automatically inserts the default

◆ SQL KEYS :

- SQL keys are used to uniquely identify rows, establish relationships, enforce uniqueness and to maintain data integrity
- ◊ 1. Primary Key :
 - Main key to identify each record uniquely.
 - Only one primary key per table.
 - Cannot be NULL.
 - Often used with **AUTO_INCREMENT** (MySQL) or **SEQUENCE** (Oracle).

💻 Syntax :

```
CREATE TABLE students (
    student_id NUMBER PRIMARY KEY,
    name VARCHAR2(100)
);
```

☞ In Oracle, you can also name it:

```
CONSTRAINT pk_students PRIMARY KEY(student_id)
```

- ◊ 2. Unique Key :
 - Ensures all values in a column (or group of columns) are unique.
 - Allows one NULL (Oracle).
 - Multiple unique keys can exist in one table.

💻 Syntax :

```
CREATE TABLE employees (
    emp_id NUMBER PRIMARY KEY,
    email VARCHAR2(100) UNIQUE
);
```

- ◊ 3. Foreign Key :
 - Used to link two tables together.
 - Refers to the Primary key in another table.
 - Prevents deletion/updates if linked data exists.

💻 Syntax :

```
CREATE TABLE departments (
    dept_id NUMBER PRIMARY KEY,
    dept_name VARCHAR2(50)
);
```

```
CREATE TABLE employees (
```

```

emp_id NUMBER PRIMARY KEY,
name VARCHAR2(100),
dept_id NUMBER,
CONSTRAINT fk_dept FOREIGN KEY (dept_id) REFERENCES departments(dept_id)
);

```

◊ 4. Composite Key :

- Primary Key made up of multiple columns.
- Ensures that the combination of values is unique.

 Syntax:

```

CREATE TABLE course_enrollments (
    student_id NUMBER,
    course_id NUMBER,
    enrollment_date DATE,
    PRIMARY KEY (student_id, course_id)
);

```

◊ 5. Candidate Key :

- A column (or group of columns) that can qualify as a Primary Key.
- A table can have many candidate keys.
- One of them is chosen as the Primary Key, others remain as alternate keys.
- All candidate keys are unique and not null.

 Example :

```

-- Assume both emp_id and email can uniquely identify employees
emp_id -- candidate key (chosen as primary)
email -- candidate key (becomes alternate)

```

 One candidate key is selected as primary, others become alternate keys.

◊ 6. Alternate Key :

- A candidate key that was not selected as the primary key.
- Still unique and NOT NULL.

 Example :

```

CREATE TABLE users (
    user_id NUMBER PRIMARY KEY,
    username VARCHAR2(50) UNIQUE -- alternate key
);

```

◊ 7. Super Key :

- Any set of columns that uniquely identifies a row.
- Includes Primary Key, Composite Key, Candidate Keys, etc.
- May contain extra columns (not minimal).

 Example:

```

-- student_id is primary key
-- student_id + name = super key (but not minimal)

```

 Example of Keys and Constraints All Together :

```

CREATE TABLE employees (
    emp_id INT PRIMARY KEY,
    name VARCHAR(100) NOT NULL,

```

```

email VARCHAR(100) UNIQUE,
dept_id INT,
salary DECIMAL(10,2) CHECK (salary > 0),
status VARCHAR(10) DEFAULT 'Active',
FOREIGN KEY (dept_id) REFERENCES departments(dept_id)
);

```

 Important Notes:

- Always define primary and foreign keys for data integrity.
- Use unique keys to enforce business rules (e.g., unique email).
- Foreign keys help maintain relationships between tables.
- Choose minimal columns for Primary Keys (avoid extra columns).
- Oracle allows one NULL in **UNIQUE** constraints.

 Summary Table :

Key Type	Unique	NULL Allowed	Count per Table	Purpose
Primary Key	Yes	 No	One	Uniquely identifies a record
Unique Key	Yes	 One (Oracle)	Many	Enforce uniqueness
Foreign Key	No	 Yes	Many	Link to another table
Composite Key	Yes	 No	One	Unique combo of multiple columns
Candidate Key	Yes	 No	Many	Eligible for Primary Key
Alternate Key	Yes	 No	Many	Not chosen as Primary Key
Super Key	Yes	 No	Many	All keys that uniquely identify

◆ DDL (Data Definition Language) in SQL :

❖ What is DDL?

- DDL stands for Data Definition Language.
- It includes SQL commands that define, modify, or delete database objects such as tables, schemas, indexes, views, etc.

◊ DDL Commands :

DDL Command	Description
CREATE	Creates a new table or database object.
ALTER	Modifies an existing database object.
DROP	Deletes database objects permanently.
TRUNCATE	Removes all records from a table quickly (faster than DELETE).
RENAME	Renames a table or object.

◊ 1. CREATE TABLE :

- The **CREATE** command is used when you want to make something new in the database.
- This could be a table, a view, a sequence, a schema, etc.
- When you create a table, you define the structure — like column names, data types, and constraints.
- Once a table is created, it will exist until you delete it explicitly using **DROP**.



Syntax :

```
CREATE TABLE table_name (
    column1 datatype,
    column2 datatype,
    ...
);
```

💡 Example :

```
CREATE TABLE employees (
    emp_id NUMBER PRIMARY KEY,
    name VARCHAR2(50),
    salary NUMBER
);
```

📌 Note : Table is created with three columns. **emp_id** is marked as Primary Key.

◊ 2. ALTER TABLE :

- The **ALTER** command is used when you need to change the structure of an existing object like a table.
- You can add a new column, delete a column, rename a column, or change a column's data type.
- You can also use it to add or remove constraints.
- It's a powerful command used during design updates or changes to business needs.

 Syntax:

```
ALTER TABLE table_name  
ADD column_name datatype;
```

 Example:

```
ALTER TABLE employees  
ADD department VARCHAR2(30);
```

 Note: Adds a new column `department` to the existing table.

◊ 3. DROP TABLE :

- The `DROP` command is used to completely delete a database object like a table, view, or sequence.
- Once something is dropped, it's gone permanently — along with all of its data.
- You should use it carefully because it cannot be rolled back (undone).

 Syntax:

```
DROP TABLE table_name;
```

 Example:

```
DROP TABLE employees;
```

 Note: Deletes the table and all of its data permanently.

◊ 4. TRUNCATE TABLE :

- The `TRUNCATE` command is used to quickly delete all data from a table, but keep the table structure intact.
- Unlike `DELETE`, it does not log each row deletion, so it's faster.
- However, it cannot be rolled back, and you can't use a `WHERE` clause with it.

 Syntax:

```
TRUNCATE TABLE table_name;
```

 Example:

```
TRUNCATE TABLE employees;
```

 Note: Removes all records from the table but does not delete the table.

 Tip: `TRUNCATE` is faster than `DELETE` and cannot be rolled back.

◊ 5. RENAME :

- The `RENAME` command is used to change the name of a table or other database object.
- It's useful if you want your naming to reflect updated business logic.
- The structure and data remain the same — only the object's name changes.

 Syntax :

```
RENAME old_table_name TO new_table_name;
```

 Example :

```
RENAME employees TO emp_details;
```

 Note: Renames the table from `employees` to `emp_details`.

◆ DML (Data Manipulation Language) in SQL :

❖ What is DML :

- DML stands for Data Manipulation Language.
- It is used to add, update, delete, and retrieve data from database tables.
- These commands do not change the structure of the table — they only work with the data inside.

◊ 1. INSERT – Add New Data

- This command is used when you want to add new records (rows) into a table.
- Think of it like filling out a new entry in a register — every detail is fresh and complete.
- You use it when new employees join, when you get a new order, or when new products are added.
- All required (NOT NULL) fields must be filled when inserting a new row.

💻 Syntax:

```
INSERT INTO table_name (column1, column2, ...)  
VALUES (value1, value2, ...);
```

💡 Example:

```
INSERT INTO employees (emp_id, name, salary)  
VALUES (101, 'Vaibhav', 55000);
```

📌 Note:

- All NOT NULL columns must be included in INSERT.
- You can insert multiple rows using **INSERT ALL** (Oracle) or bulk methods.

◊ 2. UPDATE – Modify Existing Data

- This command helps you modify existing data in your table.
- For example, if someone's address changes, or an employee gets a salary hike, you use UPDATE to reflect those changes.
- It's important to carefully choose which records to update, or you might change data in multiple rows by accident.
- Used when data is outdated or incorrect and needs fixing.

💻 Syntax:

```
UPDATE table_name  
SET column1 = value1, column2 = value2, ...  
WHERE condition;
```

💡 Example:

```
UPDATE employees  
SET salary = 60000  
WHERE emp_id = 101;
```

📌 Note: Always use a **WHERE** clause to avoid updating all records by mistake.

📝 Tip: Without **WHERE**, all rows will be updated.

◊ 3. DELETE – Remove Data

- This command is used to remove specific rows from a table.
- You might delete a customer record if they've unsubscribed or remove an order that was canceled.

- Like UPDATE, it should be used with conditions to avoid deleting everything by mistake.
- Unlike TRUNCATE (a DDL command), DELETE allows you to remove only specific records.

 Syntax:

```
DELETE FROM table_name
WHERE condition;
```

 Example:

```
DELETE FROM employees
WHERE emp_id = 101;
```

 Note:

- Like UPDATE, ALWAYS use a WHERE clause.
- If WHERE is skipped, all rows will be deleted.

 Important:

- DELETE is slower than TRUNCATE because it logs individual row deletions and can be rolled back.

◊ 4. SELECT – Read Data

- This is the most commonly used command. It's used to retrieve data from a table — you're just reading the data, not changing it.
- It helps you answer questions like: "Who are the top 5 customers?", "What are the total sales this month?", or "Which products are low in stock?"
- Used to retrieve specific columns, filter rows, group results, and sort data.
- It's the foundation for reports, dashboards, and analysis.

 Syntax:

```
SELECT column1, column2, ...
FROM table_name
WHERE condition;
```

 Example:

```
SELECT name, salary
FROM employees
WHERE salary > 50000;
```

 Note:

- Use * to select all columns: `SELECT * FROM employees;`
- You can use WHERE, ORDER BY, GROUP BY, HAVING, etc., for advanced queries.

 Important Notes:

- DML commands affect the data, not the table structure.
- INSERT, UPDATE, and DELETE can be rolled back if something goes wrong (if not yet committed).
- SELECT is read-only and does not change data.
- These commands are often used in applications, reports, dashboards, and any system that interacts with a database.

Summary of DML Commands:

Command	Action	Can Be Rolled Back?
INSERT	Add new records	<input checked="" type="checkbox"/> Yes
UPDATE	Modify existing records	<input checked="" type="checkbox"/> Yes
DELETE	Remove specific records	<input checked="" type="checkbox"/> Yes
SELECT	Retrieve data (read-only)	<input checked="" type="checkbox"/> Not needed

◆ DQL (Data Query Language) :

❖ What is DQL?

- DQL stands for Data Query Language.
- It is used to fetch (retrieving) data from the database.
- It doesn't change the structure or content of the data
- They are read-only operations — they don't insert, update, or delete anything.

❖ Notes:

- It plays a major role in data analysis, reporting, and business intelligence.
- Other SQL command types (DDL, DML, DCL, TCL) may prepare or modify data, but DQL is used to understand and explore data.
- ✓ There is mainly one command in DQL: SELECT

❖ DQL Commands :

❖ 1. SELECT – Retrieve Data from Tables

- Used to get data from one or more tables in a database..
- You can select specific columns, all columns, or even calculated values.
- It supports a wide range of options like filtering, sorting, grouping, joining, and limiting data.
- It is a read-only command and doesn't modify data.

❖ Real-Life Use Cases:

- Viewing all customers from a specific city.
- Getting a list of employees with salaries above a certain amount.
- Checking daily sales reports.
- Generating dashboards or analytics reports.
- Reading data for machine learning models or data analysis.

❖ Notes: Works closely with WHERE, GROUP BY, ORDER BY, HAVING, LIMIT.

❖ 1. Basic SELECT

💻 Syntax:

```
SELECT column1, column2 FROM table_name;
```

💡 Example:

```
SELECT first_name, last_name FROM employees;
```

❖ Notes:

- Use * to select all columns: `SELECT * FROM employees;`

❖ 2. DISTINCT – Remove Duplicates

- Used to return only unique values from a column.
- It removes all duplicate rows in the result set.

💻 Syntax:

```
SELECT DISTINCT column_name FROM table_name;
```

💡 Example:

```
SELECT DISTINCT department FROM employees;
```

❖ Notes:

- Useful when you want unique values from a column.
- Applies only to selected columns.

◊ 3. WHERE – Filter Rows

- Used to filter records based on a given condition.
- Only rows that satisfy the condition will be selected.

💻 Syntax:

```
SELECT column1 FROM table_name WHERE condition;
```

💡 Example:

```
SELECT * FROM employees WHERE salary > 50000;
```

📌 Notes:

- Supports operators: `=, >, <, >=, <=, !=, BETWEEN, LIKE, IN, IS NULL`
- Use `AND, OR, NOT` for multiple conditions.

◊ 4. ORDER BY – Sort Results

- Used to sort the result set by one or more columns.
- Sorting can be in ascending (ASC) or descending (DESC) order.

💻 Syntax:

```
SELECT * FROM table_name ORDER BY column1 [ASC|DESC];
```

💡 Example:

```
SELECT * FROM employees ORDER BY salary DESC;
```

📌 Notes:

- The default is `ASC` (ascending).
- Can sort by multiple columns.

◊ 5. GROUP BY – Group Rows for Aggregation

- Used to group rows that have the same values in specified columns.
- Often used with aggregate functions like `SUM()`, `COUNT()`, `AVG()`, etc.

💻 Syntax:

```
SELECT column, AGG_FUNC(column2) FROM table_name GROUP BY column;
```

💡 Example:

```
SELECT department, COUNT(*) FROM employees GROUP BY department;
```

📌 Notes:

- Used with aggregation functions like `COUNT()`, `SUM()`, `AVG()`, `MAX()`, `MIN()`
- All columns in `SELECT` must either be in `GROUP BY` or an aggregated result.

◊ 6. HAVING – Filter After Grouping

- Used to filter groups created by `GROUP BY`.
- It works like `WHERE`, but is applied after grouping.

💻 Syntax:

```
SELECT column, AGG_FUNC(col) FROM table GROUP BY column HAVING  
condition;
```

💡 Example:

```
SELECT department, COUNT(*) FROM employees GROUP BY department  
HAVING COUNT(*) > 5;
```

📌 Notes: `WHERE` filters before grouping. & `HAVING` filters after grouping.

◊ 7. LIMIT / TOP / ROWNUM

- Restrict Number of Rows
 - MySQL/PostgreSQL: `LIMIT`
 - SQL Server: `TOP`
 - Oracle: `ROWNUM`

💡 Example (MySQL):

```
SELECT * FROM employees LIMIT 5;
```

💡 Example (SQL Server):

```
SELECT TOP 3 * FROM employees;
```

💡 Example (Oracle):

```
SELECT * FROM employees WHERE ROWNUM <= 3;
```

◊ 8. JOINS – Combine Data from Multiple Tables

- Used to combine rows from two or more tables based on a related column.

💻 Syntax (INNER JOIN):

```
SELECT a.col1, b.col2 FROM table1 a  
JOIN table2 b ON a.common_column = b.common_column;
```

💡 Example:

```
SELECT employees.name, departments.dept_name  
FROM employees  
JOIN departments ON employees.dept_id = departments.id;
```

◆ Types of Joins:

- `INNER JOIN`: Only matching rows
- `LEFT JOIN`: All rows from left, matched from right
- `RIGHT JOIN`: All from right, matched from left
- `FULL JOIN`: All rows from both tables
- `SELF JOIN`: Join table with itself
- `CROSS JOIN`: Cartesian product

🔍 Difference between WHERE vs HAVING Clause in SQL :

❖ WHERE Clause	❖ HAVING Clause
Filters rows before grouping	Filters groups after aggregation
Used with <code>SELECT, UPDATE, DELETE</code>	Used only with <code>SELECT + GROUP BY</code>
Cannot use aggregate functions (<code>SUM()</code> , etc.)	Can use aggregate functions like <code>SUM()</code> , <code>AVG()</code> , <code>COUNT()</code>
Applied before rows are grouped	Applied after <code>GROUP BY</code> has grouped rows
Works on individual records	Works on grouped/aggregated data
Used when no <code>GROUP BY</code> is present	Used when <code>GROUP BY</code> is present
Improves query performance (early filtering)	Typically used for reporting on aggregated results

 Simple Analogy:

- **WHERE** is like filtering raw vegetables before cooking.
- **HAVING** is like tasting the cooked dish and deciding whether to serve it.

 Example Showing Both:

-- Select departments with more than 5 employees and salary > 50000

SELECT department, COUNT(*) as emp_count

FROM employees

WHERE salary > 50000 -- Filters individual records

GROUP BY department

HAVING COUNT(*) > 5; -- Filters grouped results

 Important Notes:

- SQL is case-insensitive, but it's a good habit to write keywords in UPPERCASE.
- Always use aliases for clarity when joining tables.
- Use GROUP BY with aggregate functions, not WHERE.
- **WHERE** filters individual rows, **HAVING** filters aggregated results.

 Summary Table :

Clause	Purpose	Used With
SELECT	Retrieve data	All columns/tables
DISTINCT	Remove duplicate rows	After SELECT
WHERE	Filter raw data	Before grouping
ORDER BY	Sort the result set	At the end
GROUP BY	Group data for aggregation	With aggregate funcs
HAVING	Filter grouped data	After GROUP BY
LIMIT/TOP	Restrict number of rows returned	Last clause
JOIN	Combine rows from tables	ON condition

◆ Transaction Control Language (TCL) :

❖ What is TCL?

- TCL stands for Transaction Control Language.
- TCL commands are used to manage transactions in a database.
- A transaction is a group of operations (like INSERT, UPDATE, DELETE) that are executed as a single logical unit of work.
- They ensure data integrity and allow you to accept or discard changes made to the database during a transaction.

◊ TCL Commands :

◊ 1. COMMIT :

- Permanently saves all changes made during the current transaction to the database.
- Once you commit, you cannot undo those changes.

 When it's used:

- After a successful set of operations.
- To finalize a transaction so the changes are visible to other users.

 Think of COMMIT as saving your work in a Word file.

 Syntax:

COMMIT;

 Example:

```
INSERT INTO employees (id, name, salary)
VALUES (101, 'Vaibhav', 50000);
COMMIT;
```

 This will permanently save the inserted row in the `employees` table.

 Note: After COMMIT, changes are visible to other users and cannot be rolled back.

◊ 2. ROLLBACK :

- Cancels all changes made in the current transaction.
- Returns the database to the state it was in before the transaction began.

 When it's used:

- If something goes wrong (error, wrong data inserted, failed condition).
- To undo any uncommitted changes.

 Think of ROLLBACK as pressing 'Undo' or closing a file without saving.

 Syntax:

ROLLBACK;

 Example:

```
UPDATE employees
SET salary = 70000
WHERE id = 101;
```

ROLLBACK;

 This will undo the salary update for employee 101.

 Note: You can only ROLLBACK if you haven't committed yet.

◊ 3. SAVEPOINT :

- Creates a bookmark or checkpoint within a transaction.
- Allows partial rollback to that point instead of rolling back the entire transaction.

 When it's used:

- In complex transactions where you want to save progress in steps.
- Useful for selectively undoing certain parts of a transaction.

 Think of SAVEPOINT like checkpoints in a game — you can return there if something goes wrong later.

 Syntax:

```
SAVEPOINT savepoint_name;
```

 Example:

```
SAVEPOINT sp1;
```

```
INSERT INTO employees (id, name, salary)  
VALUES (102, 'Rahul', 60000);
```

```
SAVEPOINT sp2;
```

```
UPDATE employees  
SET salary = 65000  
WHERE id = 102;
```

```
ROLLBACK TO sp1;
```

 This will undo everything after SAVEPOINT sp1, including the update and insert.

 Note: SAVEPOINTS help in partial rollbacks within a larger transaction.

◊ 4. SET TRANSACTION (less commonly used) :

- Configures properties of the current transaction.
- Can specify things like isolation level or read-only mode.

 When it's used:

- When you need fine control over how a transaction behaves (e.g., in multi-user environments).

 Think of it like setting rules for how your transaction should behave.

 Syntax:

```
SET TRANSACTION [READ ONLY | READ WRITE];
```

 Example:

```
SET TRANSACTION READ ONLY;  
SELECT * FROM employees;  
COMMIT;
```

 This transaction can only read data, not modify it.

 Note: **SET TRANSACTION** is advanced and rarely used in beginner-level projects. It's mainly useful in multi-user environments.

 Important Notes:

- TCL commands only work with DML commands (**INSERT**, **UPDATE**, **DELETE**) — not with DDL.
- If AUTO-COMMIT is ON (like in some tools), then COMMIT happens automatically after every command.
- You should use TCL to maintain data integrity and error recovery in your applications.

 Summary of TCL Commands:

Command	Purpose
COMMIT	Save all changes permanently
ROLLBACK	Undo all changes made in the current session
SAVEPOINT	Create a restore point within a transaction
SET TRANSACTION	Set properties for the current transaction

◆ What is DCL (Data Control Language)?

❖ What is DCL ?

- DCL stands for Data Control Language.
- DCL commands are used to control access to data in the database.
- They help the database administrator or user with high privileges to grant or revoke rights (like SELECT, INSERT, UPDATE, etc.) on database objects (like tables, views, etc.).

◆ DCL commands:

◊ 1. GRANT :

- **GRANT** command is used to give specific privileges to users.
- For example: allow someone to SELECT or INSERT into a table.

💻 Syntax:

```
GRANT privilege_name  
ON object_name  
TO user_name;
```

💡 Example:

```
GRANT SELECT, INSERT  
ON employees  
TO rahul;
```

- ✓ This means the user **rahul** is allowed to view & insert records into the **employees** table.

📌 Note:

- You can grant multiple privileges like **SELECT, INSERT, UPDATE, DELETE**, etc.
- Objects can be tables, views, procedures, etc.
- In Oracle, use **WITH GRANT OPTION** to allow the user to grant the same permission to others.

◊ 2. REVOKE :

- **REVOKE** command is used to remove previously granted privileges from a user.

💻 Syntax:

```
REVOKE privilege_name  
ON object_name  
FROM user_name;
```

💡 Example:

```
REVOKE INSERT  
ON employees  
FROM rahul;
```

- ✓ This removes the **INSERT** privilege from the user **rahul**. Now he can only **SELECT** (if that privilege remains).

📌 Note:

- If you revoke all privileges, the user cannot access the table anymore.
- Always revoke specific permissions when a user leaves a project or role.

❖ Real Life Analogy:

- **GRANT** is like giving keys to someone for a room.
- **REVOKE** is like taking the keys back when they no longer need access.

📝 Important Notes:

- **GRANT** and **REVOKE** apply to specific database objects like tables, views, procedures, etc.
- DCL ensures data security by limiting access to only authorized users.
- Only database administrators (DBAs) or users with sufficient privileges can use DCL commands.
- Be careful when granting UPDATE, DELETE access as it allows modifying or deleting data.

📋 Summary Table:

Command	Action	Syntax Pattern
GRANT	Give privileges to a user	GRANT SELECT ON table TO user;
REVOKE	Take away privileges from user	REVOKE SELECT ON table FROM user;

◆ SQL JOINS:

- A SQL JOIN is used to combine data from two or more tables based on a common column that links them.
- In real-world databases, data is split across multiple related tables (like employees and departments).
- JOINs allow you to fetch meaningful, combined information from these tables.

❖ Types of SQL JOINS :

Join Type	Returns
◆ INNER JOIN	Only matching rows from both tables
◆ LEFT JOIN	All rows from left table + matching rows from right table
◆ RIGHT JOIN	All rows from right table + matching from left table
◆ FULL OUTER JOIN	All rows from both tables, matched where possible
◆ CROSS JOIN	Cartesian product – every row from A with every row from B
◆ SELF JOIN	Joining a table with itself

◆ Sample Tables Used for Examples :

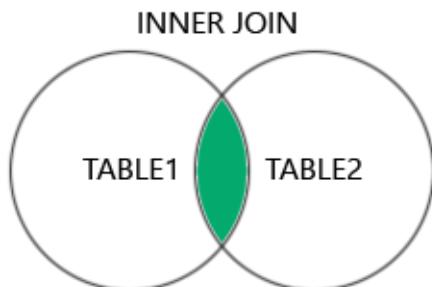
◊ Employees Table:

emp_id	emp_name	dept_id
1	John	101
2	Sara	102
3	Mike	103
4	Lily	NULL

◊ Departments Table:

dept_id	dept_name
101	HR
102	IT
104	Finance

- ◊ 1. INNER JOIN :
 - It gives only the matching rows from both tables.
 - If there's no match between the tables, the row is not included in the final result.
 - Think of it as: “*Give me only the data where both tables have a valid match.*”



❑ Syntax:

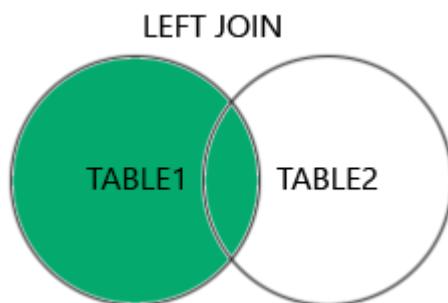
```
SELECT e.emp_name, d.dept_name
FROM employees e
INNER JOIN departments d
ON e.dept_id = d.dept_id;
```

❑ Output:

emp_name	dept_name
John	HR
Sara	IT

❑ Note: Mike and Lily are excluded because no matching `dept_id` exists.

- ◊ 2. LEFT JOIN (LEFT OUTER JOIN):
 - Returns all rows from the left table.
 - Also returns the matching rows from the right table.
 - If there is no match, the result still includes the left table row and shows NULL for the right table columns.
 - Think of it as: “*Keep everything from the left, fill in the right when you can.*”



❑ Syntax:

```
SELECT e.emp_name, d.dept_name
FROM employees e
LEFT JOIN departments d
ON e.dept_id = d.dept_id;
```

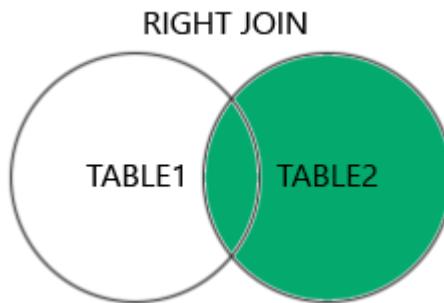
💡 Output:

Emp_name	dept_name
John	HR
Sara	IT
Mike	NULL
Lily	NULL

📝 Note: Unmatched rows from the left side appear with **NULL** for right table columns.

◊ 3. RIGHT JOIN (RIGHT OUTER JOIN):

- The opposite of LEFT JOIN.
- Returns all rows from the right table.
- Also includes the matching rows from the left table.
- If there is no match, it still keeps the right table data and fills NULL for the left.



💻 Syntax:

```
SELECT e.emp_name, d.dept_name  
FROM employees e  
RIGHT JOIN departments d  
ON e.dept_id = d.dept_id;
```

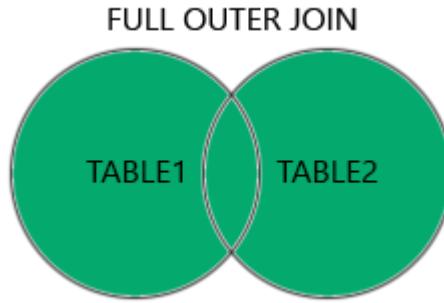
💡 Output:

emp_name	dept_name
John	HR
Sara	IT
NULL	Finance

📝 Note: Returns unmatched rows from the right side with **NULL** from left table.

◊ 4. FULL OUTER JOIN:

- Combines the result of both LEFT and RIGHT JOIN.
- Returns all rows from both tables, with NULLs in places where no match is found.
- Think of it as: “*Show me everything from both tables, matched or not.*”



❑ Syntax:

```
SELECT e.emp_name, d.dept_name
FROM employees e
FULL OUTER JOIN departments d
ON e.dept_id = d.dept_id;
```

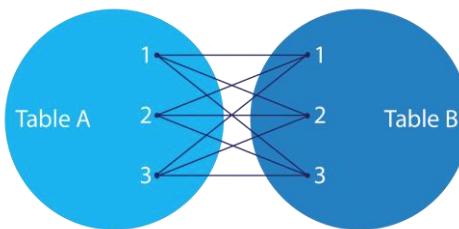
💡 Output:

emp_name dept_name	
----- -----	
John	HR
Sara	IT
Mike	NULL
Lily	NULL
NULL	Finance

✍ Note: This is a combination of LEFT and RIGHT JOIN.

◊ 5. CROSS JOIN:

- Returns the cartesian product of both tables.
- That means every row of the first table is combined with every row of the second.
- Used rarely, and typically when we want all possible combinations.



❑ Syntax:

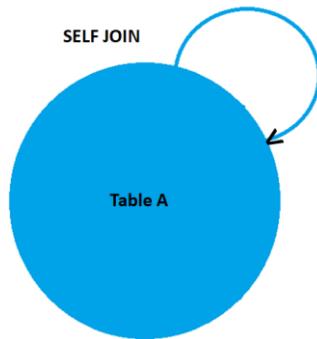
```
SELECT e.emp_name, d.dept_name
FROM employees e
CROSS JOIN departments d;
```

✍ Be careful! If `employees` have 4 rows and `departments` have 3, the result will have $4 \times 3 = 12$ rows.

◊ 6. SELF JOIN :

- A Self JOIN is a regular join where a table is joined with itself.

- It's useful when the rows in the same table are related to each other, such as an employee reporting to a manager — both are part of the same table.



◆ Key Points About Self Join :

- The same table is used twice, but with different aliases to differentiate them.
- It is often used to compare rows within the same table.
- Helps in building hierarchical relationships, such as:
 - Employees and their managers.
 - Categories and sub-categories.
 - Friends-of-friends in social networks.

Example Scenarios Where Self Join is Useful:

- In an employee table, if each employee has a manager (also listed as an employee), you can use self join to display:
 - Employee Name → Manager Name.
- In a table of parts and subparts, you can link a part with its parent part using self join.

Important Notes:

- Aliases are necessary in self joins to avoid confusion between column references.
- Self joins help uncover relationships within a single entity.
- It's logically similar to joining two different tables but actually joins different rows of the same table.

Example:

```
SELECT A.emp_name AS Employee, B.emp_name AS Manager
FROM employees A
JOIN employees B ON A.manager_id = B.emp_id;
```

Useful for hierarchical relationships like employees and managers.

When to Use LEFT JOIN Over INNER JOIN?

- Use LEFT JOIN when:
 - You want to see all records from the left table, even if they don't have a match in the right table.
 - You are okay with seeing NULLs for unmatched data from the right.
 - Example scenarios:
 - Showing customers and their orders (even if some customers didn't order anything).
 - Listing employees and their departments (even if some employees are not assigned).

- Use INNER JOIN when:
 - You want only the matched data from both tables.
 - You are only interested in records that exist on both sides.

 Important Notes:

- Always use aliases when working with multiple tables – improves clarity.
- JOIN conditions should always be defined using **ON**, not **WHERE**, especially for OUTER joins.
- Performance tip: Use **INNER JOIN** for most use-cases unless nulls are needed.
- Use **IS NULL** to detect unmatched rows in OUTER joins.
- Not all databases support **FULL OUTER JOIN** directly – in such cases use:
LEFT JOIN ...
UNION
RIGHT JOIN ...

 Differences Between JOINS:

Feature	INNER JOIN	LEFT JOIN	RIGHT JOIN	FULL OUTER JOIN
Matching Rows	✓	✓	✓	✓
Unmatched Left Rows	✗	✓	✗	✓
Unmatched Right Rows	✗	✗	✓	✓
NULLS in result	✗	For unmatched right	For unmatched left	For both sides
Common Usage	To fetch only related records	When all records from left table are needed	When all right side data is needed	When complete data from both sides is needed

◆ SQL SET OPERATORS :

- SQL Set Operators are used to combine the results of two or more SELECT queries.
- They work like set theory in mathematics — they merge, compare, or filter rows across two result sets.
- To use them:
 - Both queries must have the same number of columns.
 - The columns must have compatible data types.
 - The column order should match.

❖ Types of SQL Set Operators:

- There are four main Set Operators in SQL:

Operator	Description
UNION	Combines results of two SELECTs, removes duplicates.
UNION ALL	Combines results of two SELECTs, keeps all duplicates.
INTERSECT	Returns only common records in both SELECTs.
MINUS (Oracle) / EXCEPT (SQL Server, PostgreSQL)	Returns records from the first SELECT that aren't in the second.

❖ 1. UNION:

- Combines result sets and removes duplicate rows.
- Used When: You want to merge data from two queries but remove duplicates.

💡 Example Use Cases:

- Combining customer data from two regions and showing only unique customers.
- Merging employees from two departments, excluding duplicates.

💻 Syntax:

```
SELECT column1, column2 FROM table1
UNION
SELECT column1, column2 FROM table2;
```

💡 Example:

◆ Tables:

◆ Employees_2023

Emp_ID	Name
101	Rahul
102	Priya

◆ Employees_2024

Emp_ID	Name
102	Priya
103	Sameer

◆ Example:

```
SELECT Emp_ID, Name FROM Employees_2023  
UNION  
SELECT Emp_ID, Name FROM Employees_2024;
```

◆ Output:

Emp_ID	Name
101	Rahul
102	Priya
103	Sameer

✓ Removes duplicate entry 102 - Priya

📌 Important Notes:

- Duplicates are removed.
- Automatically sorts records in some DBs (like Oracle).
- Can be slower than UNION ALL due to the sorting process.

◊ 2. UNION ALL:

- Same as UNION, but does not remove duplicates.
- Used When: You want to retain duplicates for accurate record-keeping or counts.
- Faster than UNION since no duplicate-checking is done.
- Use when duplicates matter, like analytics and reporting.

💡 Example Use Cases:

- Merging multiple order records across months for total revenue tracking.
- Getting all attendance logs, even if users clock in/out multiple times.

📌 Notes: May need to use GROUP BY or DISTINCT manually if uniqueness is needed.

📘 Syntax:

```
SELECT column1, column2 FROM table1  
UNION ALL  
SELECT column1, column2 FROM table2;
```

💡 Example:

◆ Same tables as above

```
SELECT Emp_ID, Name FROM Employees_2023  
UNION ALL  
SELECT Emp_ID, Name FROM Employees_2024;
```

◆ Output:

Emp_ID	Name
101	Rahul

102	Priya
102	Priya
103	Sameer

✓ Shows all entries, including duplicate 102 - Priya

◊ 3. INTERSECT:

- Returns only the common rows between two SELECT statements.
- Used When: You need shared data points across sources.

💡 Example Use Cases:

- Find customers who are both buyers and newsletter subscribers.
- Identify students enrolled in both courses A and B.

📌 Notes:

- Duplicates are removed by default.
- Not supported in MySQL (workaround needed using JOIN + GROUP BY).
- Useful in comparative queries, e.g., common customers, shared access.

💻 Syntax:

```
SELECT column1, column2 FROM table1
INTERSECT
SELECT column1, column2 FROM table2;
```

💡 Example:

◆ Customers_Mumbai

Cust_ID	Name
1	Anil
2	Ramesh

◆ Customers_Delhi

Cust_ID	Name
2	Ramesh
3	Suresh

💡 Example:

```
SELECT Cust_ID, Name FROM Customers_Mumbai
INTERSECT
SELECT Cust_ID, Name FROM Customers_Delhi;
```

◆ Output:

Cust_ID	Name
2	Ramesh

✓ Only the common row is returned.

◊ 4. MINUS / EXCEPT:

- Returns rows from the first query that are not present in the second.
- Oracle uses **MINUS**, while SQL Server & PostgreSQL use **EXCEPT**.

💡 Example Use Cases:

- Find employees in HR who are not in the company email system.
- Show products listed in the catalog that have no sales.

📌 Important Notes:

- Order matters: **A MINUS B ≠ B MINUS A**.
- Removes duplicates automatically.
- MySQL doesn't support **MINUS** or **EXCEPT** natively (requires workaround).
- Great for finding unmatched data, e.g., not enrolled, not purchased.

💻 Syntax (Oracle):

```
SELECT column1, column2 FROM table1  
MINUS  
SELECT column1, column2 FROM table2;
```

💻 Syntax (SQL Server/PostgreSQL):

```
SELECT column1, column2 FROM table1  
EXCEPT  
SELECT column1, column2 FROM table2;
```

💡 Example:

◆ All_Students

Roll_No	Name
101	Riya
102	Amit
103	Soham

◆ Library_Members

Roll_No	Name
101	Riya
103	Soham

 Example:

```
SELECT Roll_No, Name FROM All_Students  
MINUS  
SELECT Roll_No, Name FROM Library_Members;
```

◆ Output:

Roll_No	Name
102	Amit

✓ Only shows those not found in the second table.

 Important Notes:

- ◆ All Set Operators:
 - Must return the same number of columns.
 - Same data type and sequence of columns.
 - Column names are taken from the first SELECT statement.
 - Results can be ordered using **ORDER BY** after the final SELECT block.
- ◆ Preferred Practices:
 - Use **UNION** when uniqueness is important.
 - Use **UNION ALL** for performance-sensitive applications where duplicates are valid.
 - Use **INTERSECT** and **MINUS/EXCEPT** for comparison-focused operations.
- ◆ Common Mistakes:
 - Mismatched column count or data types.
 - Placing **ORDER BY** in each individual SELECT (should only be in the last SELECT).

 Summary: SQL Set Operators:

Operator	Purpose	Removes Duplicates	Common Use Case
UNION	Combines results of two queries and removes duplicates	 Yes	Merging datasets where unique results are needed
UNION ALL	Combines results of two queries without removing duplicates	 No	Faster merging when all data is important
INTERSECT	Returns only common rows between two queries	 Yes	Finding shared data (e.g., common customers)
MINUS / EXCEPT	Returns rows in first query not found in the second (MINUS = Oracle, EXCEPT = SQL Server/PostgreSQL)	 Yes	Identifying unmatched or missing data

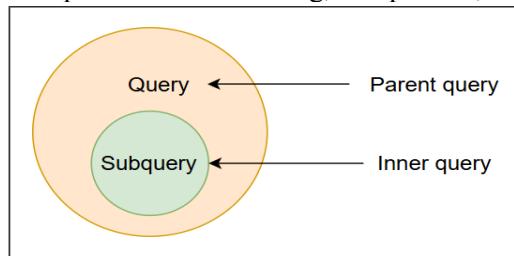
◆ SQL Subqueries:

❖ What is a Subquery?

- A subquery is a query that is written inside another query (usually inside the `SELECT`, `FROM`, or `WHERE` clause).
- It is also called a nested query or inner query.
- It helps break down complex problems by running a query inside another.

◆ Why use Subqueries?

- To divide complex SQL logic into manageable parts.
- To use results from one query as input to another.
- To perform operations like filtering, comparison, or aggregation in one step.



❖ Types of Subqueries:

◊ 1. Nested Subquery (Simple Subquery)

- A Nested Subquery, also called a Simple Subquery or Non-Correlated Subquery.
- Is a query inside another query
- It runs independently of the outer query and its result is passed to the main query.
- Common Use: Inside `WHERE`, `IN`, `=`, `<`, `>`, etc.



Syntax Examples:

```
SELECT name FROM employees  
WHERE salary > (SELECT AVG(salary) FROM employees);
```



Notes:

- Executes independently.
- Can return a single value or a list (in case of `IN`), a row, or multiple rows depending on the context.
- Can be placed in the `SELECT`, `FROM`, or `WHERE` clause.

◊ 2. Correlated Subquery:

- A Correlated Subquery is a subquery that depends on a value from the outer query.
- It runs once for each row processed by the outer query.



Syntax Examples:

```
SELECT name FROM employees e1  
WHERE salary > (SELECT AVG(salary)  
                  FROM employees e2  
                  WHERE e1.department_id = e2.department_id);
```



Notes:

- Slower than nested queries due to multiple executions.— once for each row in the outer query.
- Useful when the inner query needs a value from the outer query row.

- Often used in WHERE EXISTS or WHERE clause but also with IN, =, <, etc.
 - It cannot run independently because it uses columns from the outer query.
- ◊ 3. Multi-row Subquery
- A subquery that returns more than one row. You use it with operators like IN, ANY, ALL.
 - You can use IN, ANY, ALL, or even a JOIN to handle multi-row results.
 - If you're told "multi-row subquery," they usually mean:
 - A subquery intended to return multiple rows
 - Not a single-value result
 - And not one used with comparison operators like = (which fails with multiple rows)

 Syntax Examples:

```
SELECT employee_name
FROM employees
WHERE department_id IN (
    SELECT department_id
    FROM departments
    WHERE location = 'New York'
);
```

◆ Notes:

- Returns a list of values.
- Cannot be used with =, <, > unless combined with ANY or ALL.

◊ 4. Subquery with IN:

- Used when you want to filter the outer query using a list of results from inner query.
- Basic concept: A subquery that returns multiple values (more than one row) to the outer query.
- The outer query uses IN to check if a value exists in the list returned by the subquery.
- So technically, *every* subquery with IN is a multi-row subquery — as long as it can return multiple rows.

 Syntax Examples:

```
SELECT name
FROM employees
WHERE department_id IN (
    SELECT department_id FROM departments WHERE location = 'New York'
);
```

◆ Notes:

- Can be replaced by a JOIN in some cases (for performance)
- Safer with multiple results.

◊ 5. Subquery with EXISTS:

- Checks whether at least one row is returned by the subquery. Returns TRUE or FALSE.

 Syntax Examples:

```
SELECT d.department_name
FROM departments d
WHERE EXISTS (SELECT 1 FROM employees e WHERE e.department_id =
d.department_id);
```

◆ Notes:

- Great for performance with large data.

- Returns TRUE if a single row exists—doesn't retrieve data.
- Often used with correlated subqueries.

 **Important Notes:**

- You cannot use a subquery in the `SELECT` clause to return multiple values (only one).
- Subqueries can be used in:
 - `SELECT` clause
 - `FROM` clause
 - `WHERE` clause
- Correlated Subqueries are generally slower than Joins or nested subqueries due to row-by-row evaluation.
- Use `EXISTS` when you only need to check the presence of rows, not the actual values.
- Prefer `JOINS` if you want to fetch actual data from multiple tables instead of filtering based on it.
- If using `ORDER BY` in subqueries, you must use `TOP`, `LIMIT`, or `ROWNUM` depending on your SQL dialect.

 **Summary: Subqueries in SQL:**

Type	Description	Common Use Case
Nested Subquery	A standalone query inside another query	Comparing values with an aggregate result
Correlated	Depends on outer query, runs for each row	Row-wise comparison within groups
Multi-row	Returns multiple values; uses <code>IN</code> , <code>ANY</code> , <code>ALL</code>	Filtering by a list of results
IN Subquery	Filters values found in the subquery result list	Check membership in another set
<code>EXISTS</code> Subquery	Checks if at least one row exists in the subquery result	Presence check, performance-friendly

 **SELECT 1 FROM employees what does it means select 1?**

`SELECT 1 FROM employees;`

- This query does not retrieve any column from the `employees` table.
- Instead, it simply returns the value `1` for each row in the `employees` table.
- So if your `employees` table has 5 rows, the result will be:

```
1
1
1
1
1
1
```

 **Why is it used like that?**

- This pattern is often used when:
 - You don't care about the actual data, just want to check existence or row count.
 - Used inside `EXISTS` clauses.

⌚ Example with EXISTS:

```
IF EXISTS (SELECT 1 FROM employees WHERE department_id = 10) THEN  
    -- Do something  
END IF;
```

Here, **SELECT 1** is just a placeholder — it's used because Oracle requires a **SELECT** inside **EXISTS**, but we don't need actual data — we only care if at least one row exists.

◆ SQL Functions:

- Functions in SQL are predefined programs that perform operations on data and return a result.
- They help you manipulate data, transform it, or make decisions based on it.

❖ Types of SQL Functions:

▫ Aggregate Functions: **SUM, AVG, COUNT, MIN, MAX**

▫ Scalar Functions:

- String: **CONCAT, SUBSTRING, LENGTH, TRIM, REPLACE,...etc**
- Numeric: **ABS, CEIL, FLOOR, ROUND, TRUNC, MOD, POWER,...etc**
- Date: **SYSDATE, GETDATE, NOW, TO_DATE, LAST_DAY,...etc**
- Conversion: **TO_CHAR, TO_DATE, CAST, CONVERT, TO_NUMBER**
- Conditional: **CASE, COALESCE, NULLIF, DECODE**

▫ Analytical / Window Functions: **ROW_NUMBER, RANK, DENSE_RANK,..etc**

▫ System Functions: **USER, SESSION_USER, DATABASE, VERSION,..etc**

▫ JSON Functions: **JSON_VALUE, JSON_QUERY, JSON_OBJECT, JSON_ARRAY**

▫ XML Functions: **EXTRACTVALUE, XMLTABLE, XMLAGG, FOR XML PATH**

▫ 1. Aggregate Functions:

- Aggregate Functions perform operations on a group of rows and return a single value as a result.
- They are most often used with the **GROUP BY** clause but can also be used without it.

◆ List of Aggregate Functions:

1. **SUM()**: Returns the total of a numeric column

📊 Table:

employee_id	name	salary
101	Vaibhav	50000
102	Rohan	60000
103	Priya	40000

💻 Syntax:

```
SELECT SUM(salary) FROM employees;
```

💻 Output:

SUM(salary)

150000

2. AVG(): Returns the average of a numeric column

💻 Syntax:

SELECT AVG(salary) FROM employees;

🖨️ Output:

AVG(salary)

50000

3. COUNT(): Returns the number of rows

💻 Syntax:

SELECT COUNT(*) FROM employees;

🖨️ Output:

COUNT(*)

3

👉 You can also count specific values:

SELECT COUNT(salary) FROM employees;

4. MAX(): Finds highest value

💻 Syntax:

SELECT MAX(salary) FROM employees;

🖨️ Output:

MAX(salary)

60000

5. MIN() : Finds lowest value

💻 Syntax:

SELECT MIN(salary) FROM employees;

🖨️ Output:

MIN(salary)

40000

📝 Important Notes:

- All aggregate functions ignore NULLs by default.
- You can combine aggregate functions with **GROUP BY** to get summaries per category (e.g., per department).
- These functions do not modify data, only read and summarize it.

◊ 2. Scalar Functions:

- A Scalar Function is a function in SQL that returns a single value (a scalar) for each row in a result set.
- These functions operate on individual values, unlike aggregate functions (which operate on groups of rows).

◆ Categories of Scalar Functions:

Function Type	Description
String Functions	Modify or return information about text values
Numeric Functions	Perform operations on numbers
Date Functions	Work with date and time values
Conversion Functions	Convert data types
Conditional Functions	Return values based on conditions

◊ 1. String Functions:

1. **UPPER()**: Converts all characters in a string to uppercase.



Syntax: `UPPER(string)`



Example: `SELECT UPPER('vaibhav') FROM dual;`



Output: `VAIBHAV`

2. **LOWER()**: Converts all characters in a string to lowercase.



Syntax: `LOWER(string)`



Example: `SELECT LOWER('SQL Expert') FROM dual;`



Output: `sql expert`

3. **INITCAP()**: Converts the first letter of each word in a string to uppercase.



Syntax: `INITCAP(string)`



Example: `SELECT INITCAP('vaibhav dhakulkar') FROM dual;`



Output: `Vaibhav Dhakulkar`

4. **LENGTH()**: Returns the number of characters in a string.



Syntax: `LENGTH(string)`



Example: `SELECT LENGTH('SQL') FROM dual;`



Output: `3`

5. **SUBSTR()**: Extracts a substring from a string starting at a specific position for a given length.



Syntax: `SUBSTR(string, start_position, length)`



Example: `SELECT SUBSTR('DATABASE', 1, 4) FROM dual;`



Output: `DATA`

6. **INSTR()**: Returns the position of the first occurrence of a substring in a string.



Syntax: `INSTR(string, substring)`



Example: `SELECT INSTR('Hello World', 'o') FROM dual;`



Output: `5`

7. **TRIM()**: Removes leading and trailing spaces (or specified characters) from a string.

 Syntax: **TRIM(string)**

 Example: **SELECT TRIM(' Vaibhav ') FROM dual;**

 Output: **Vaibhav**

8. **LTRIM()**: Removes spaces from the beginning (left side) of the string.

 Syntax: **LTRIM(string)**

 Example: **SELECT LTRIM(' SQL') FROM dual;**

 Output: **SQL**

9. **RTRIM()**: Removes spaces from the end (right side) of the string.

 Syntax: **RTRIM(string)**

 Example: **SELECT RTRIM('SQL ') FROM dual;**

 Output: **SQL**

10. **REPLACE()**: Replaces a part of the string with another substring.

 Syntax: **REPLACE(string, search_string, replace_with)**

 Example: **SELECT REPLACE('SQL and PL/SQL', 'SQL', 'Oracle') FROM dual;**

 Output: **Oracle and PL/Oracle**

11. **CONCAT()**: Concatenates (joins) two strings. Only works with 2 strings in Oracle.

 Syntax: **CONCAT(string1, string2)**

 Example: **SELECT CONCAT('Data', 'Base') FROM dual;**

 Output: **DataBase**

12. **|| (Concatenation Operator)**: Concatenates multiple strings. Preferred in Oracle over **CONCAT()** for more than 2 strings.

 Syntax: **string1 || string2 || string3 ...**

 Example: **SELECT 'Vaibhav' || '' || 'Dhakulkar' FROM dual;**

 Output: **Vaibhav Dhakulkar**

Important Notes:

- **UPPER**, **LOWER**, **INITCAP** are useful for formatting names and text.
- **LENGTH**, **INSTR**, **SUBSTR** help in string analysis.
- **||** is the preferred method in Oracle for string concatenation when combining more than 2 strings.
- **REPLACE** is useful in cleaning and formatting text data.
- **INSTR** returns the position of first occurrence.
- Use **SUBSTR()** to extract values like first name, last name, etc.
- **TRIM()**, **LTRIM()**, and **RTRIM()** are useful for removing extra spaces before inserting or comparing strings.

Summary: String Functions

Function	Description	Supported In
CONCAT()	Joins two or more strings	 All
SUBSTR()	Extracts part of a string	 SQL Server
LEFT()	Left part of a string	 Oracle,
RIGHT()	Right part of a string	 Oracle
LENGTH()	Length of string	 All
INSTR()	Position of substring	 SQL Server
REPLACE()	Replace part of a string	 All
UPPER()	Convert to uppercase	 All
LOWER()	Convert to lowercase	 All
TRIM()	Removes leading/trailing spaces	 All

◊ 2. Numeric Functions:

1. **ABS()**: Returns the absolute (positive) value of a number.

 Syntax: `ABS(number)`

 Example: `SELECT ABS(-42) FROM DUAL;`

 Output: `42`

2. **CEIL()**: Rounds a number *up* to the nearest integer.

 Syntax: `CEIL(number)`

 Example: `SELECT CEIL(3.14) FROM DUAL;`

 Output: `4`

3. **FLOOR()**: Rounds a number *down* to the nearest integer.

 Syntax: `FLOOR(number)`

 Example: `SELECT FLOOR(3.99) FROM DUAL;`

 Output: `3`

4. **MOD()**: Returns the remainder of division (modulus).

 Syntax: `MOD(dividend, divisor)`

 Example: `SELECT MOD(10, 3) FROM DUAL;`

 Output: `1`

5. **POWER()**: Raises a number to the power of another number.

 Syntax: `POWER(base, exponent)`

 Example: `SELECT POWER(2, 4) FROM DUAL;`

 Output: `16`

6. **ROUND()**: Rounds a number to a specified number of decimal places.

 Syntax: `ROUND(number, decimal_places)`

 Example: `SELECT ROUND(3.4567, 2) FROM DUAL;`

 Output: `3.46`

7. **TRUNC()**: Truncates a number to a specified number of decimal places (no rounding).

 Syntax: `TRUNC(number, decimal_places)`

 Example: `SELECT TRUNC(3.4567, 2) FROM DUAL;`

 Output: `3.45`

8. **SIGN()**: Returns `-1`, `0`, or `1` based on whether the number is negative, zero, or positive.

 Syntax: `SIGN(number)`

 Example: `SELECT SIGN(-20) FROM DUAL;`

 Output: `-1`

9. **SQRT()**: Returns the square root of a number.

 Syntax: `SQRT(number)`

 Example: `SELECT SQRT(16) FROM DUAL;`

 Output: `4`

10. **EXP()**: Returns e raised to the power of the specified number.

 Syntax: `EXP(number)`

 Example: `SELECT EXP(1) FROM DUAL;`

 Output: `2.7182818...`

 Important Notes:

- These functions work on numeric data types (like `NUMBER`, `INTEGER`, `FLOAT`, etc.).
- `DUAL` is a special dummy table in Oracle used for quick function testing
- `TRUNC()` and `ROUND()` are often used when formatting currency or precision-sensitive data.

Summary: Numeric Functions

Function	Description	Supported In
<code>ABS()</code>	Absolute value	 All
<code>CEIL() / CEILING()</code>	Smallest integer \geq number	 All
<code>FLOOR()</code>	Largest integer \leq number	 All
<code>MOD()</code>	Remainder after division	 SQL Server
<code>ROUND()</code>	Round number to decimals	 All
<code>TRUNC()</code>	Truncate decimal part	 SQL Server,  MySQL
<code>POWER()</code>	Raise number to a power	 All
<code>SQRT()</code>	Square root	 All
<code>EXP()</code>	Exponential value	 All
<code>LOG() / LN()</code>	Natural / base-10 logarithm	 SQL Server

◊ 3. Date Functions:

1. **`SYSDATE`**: Returns the current date and time of the database server.

 Syntax: `SYSDATE`

 Example: `SELECT SYSDATE FROM DUAL;`

 Output: `23-MAY-25 10:45:00` (your system date and time)

2. **`CURRENT_DATE`**: Returns the current date and time in the user's session time zone.

 Syntax: `CURRENT_DATE`

 Example: `SELECT CURRENT_DATE FROM DUAL;`

 Output: `23-MAY-25 10:45:00` (according to session time zone)

3. **`SYSTIMESTAMP`**: Returns the current date and time with the time zone of the system.

 Syntax: `SYSTIMESTAMP`

 Example: `SELECT SYSTIMESTAMP FROM DUAL;`

 Output: `23-MAY-25 10:45:00.123456 +05:30`

4. **`ADD_MONTHS()`**: Adds or subtracts a number of months to/from a date.

 Syntax: `ADD_MONTHS(date, number_of_months)`

 Example: `SELECT ADD_MONTHS(SYSDATE, 3) FROM DUAL;`

 Output: `23-AUG-25`

5. **`MONTHS_BETWEEN()`**: Returns the number of months between two dates.

 Syntax: `MONTHS_BETWEEN(date1, date2)`

 Example: `SELECT MONTHS_BETWEEN('23-MAY-25', '23-JAN-25') FROM DUAL;`
 Output: 4

6. **NEXT_DAY()**: Returns the next date of the specified weekday after the given date.

 Syntax: `NEXT_DAY(date, 'DAY_NAME')`
 Example: `SELECT NEXT_DAY(SYSDATE, 'MONDAY') FROM DUAL;`
 Output: 27-MAY-25 (next Monday)

7. **LAST_DAY()**: Returns the last day of the month of a given date.

 Syntax: `LAST_DAY(date)`
 Example: `SELECT LAST_DAY(SYSDATE) FROM DUAL;`
 Output: 31-MAY-25

8. **ROUND(date)**: Rounds the date to the nearest day, month, or year, depending on the format.

 Syntax: `ROUND(date, 'format')`
 Example: `SELECT ROUND(TO_DATE('23-MAY-25'), 'MONTH') FROM DUAL;`
 Output: 01-JUN-25

9. **TRUNC(date)**: Truncates the date to the start of a specific unit (day, month, year, etc.).

 Syntax: `TRUNC(date, 'format')`
 Example: `SELECT TRUNC(TO_DATE('23-MAY-25'), 'MONTH') FROM DUAL;`
 Output: 01-MAY-25

10. **EXTRACT()**: Extracts a specific part (YEAR, MONTH, DAY) from a date.

 Syntax: `EXTRACT(YEAR FROM date)`
 Example: `SELECT EXTRACT(YEAR FROM SYSDATE) FROM DUAL;`
 Output: 2025

Important Notes:

- Dates are stored in Oracle in `DD-MON-YY` format by default.
- Use `TO_DATE()` to explicitly convert a string to a date format.
- You can do arithmetic on dates (e.g., `SYSDATE - 7` gives last week's date).
- Oracle automatically handles leap years, months with 30/31 days, etc.

Summary: Date Functions

Function	Description	Supported In
<code>SYSDATE</code>	Current date and time	 Oracle
<code>CURRENT_DATE</code>	Current date	 All
<code>NOW()</code>	Current timestamp	 Oracle
<code>ADD_MONTHS()</code>	Add months to date	 Oracle
<code>MONTHS_BETWEEN()</code>	Months between two dates	 Oracle
<code>NEXT_DAY()</code>	Next weekday after given date	 Oracle
<code>LAST_DAY()</code>	Last day of month	 Oracle
<code>EXTRACT()</code>	Extract part of date (e.g., year)	 All
<code>DATE_PART()</code>	Extract part of date	 PostgreSQL
<code>DATEDIFF()</code>	Days between dates	 Oracle,  PostgreSQL

◊ 4. Conversion Functions:

- ◆ 1. `TO_CHAR()`: Convert Date or Number to String & Mainly used in Oracle DB

 Syntax: `TO_CHAR(value, 'format')`

 Example:

```
SELECT TO_CHAR(SYSDATE, 'DD-Mon-YYYY') AS formatted_date FROM DUAL;
```

 Output:

`FORMATTED_DATE`

23-May-2025

- ◆ 2. `TO_DATE()` – Convert String to Date and Oracle-specific function

 Syntax: `TO_DATE('string_date', 'format')`

 Example:

```
SELECT TO_DATE('23-05-2025', 'DD-MM-YYYY') AS actual_date FROM DUAL;
```

 Output:

`ACTUAL_DATE`

2025-05-23

- ◆ 3. `TO_NUMBER()` – Convert String to Number and Oracle-specific

 Syntax: `TO_NUMBER('string_number')`

 Example: `SELECT TO_NUMBER('100') + 50 AS result FROM DUAL;`

 Output:

RESULT

150

- ◆ 4. **CAST()** – General Purpose Conversion and work in all DBMS

💻 Syntax: `CAST(expression AS data_type)`

💡 Example: `SELECT CAST('2025-05-23' AS DATE) AS converted_date;`

💻 Output: (in MySQL/PostgreSQL):

`CONVERTED_DATE`

2025-05-23

- ◆ 5. **CONVERT()** – Convert Between Data Types

⚠ Mostly used in SQL Server

💻 Syntax: `CONVERT(data_type, expression [, style])`

💡 Example: `SELECT CONVERT(VARCHAR, GETDATE(), 103) AS formatted_date;`

💻 Output: (SQL Server):

`FORMATTED_DATE`

23/05/2025

📝 Important Notes:

- Use **CAST()** when writing portable SQL across different DBMSs.
- **TO_CHAR**, **TO_DATE**, **TO_NUMBER** are Oracle-specific but very powerful for formatting.
- **CONVERT()** is handy in SQL Server for date and string formatting.
- Always match format patterns (like '**YYYY-MM-DD**') to input/output for correct results.
- If formats are mismatched, conversion may fail or return NULL.

⌚ Summary: SQL Conversion Functions

Function	Converts What	Commonly Used In
TO_CHAR()	Date/Number → String	<input checked="" type="checkbox"/> Oracle
TO_DATE()	String → Date	<input checked="" type="checkbox"/> Oracle
TO_NUMBER()	String → Number	<input checked="" type="checkbox"/> Oracle
CAST()	Any type → Any type	<input checked="" type="checkbox"/> All
CONVERT()	Expression → Data type	<input checked="" type="checkbox"/> SQL Server

◊ 5. Conditional Functions :

◆ 1. CASE :

- General Conditional Expression
- Works like **IF-THEN-ELSE**, lets you evaluate conditions and return specific
- Works in all DBMSs (Oracle, SQL Server, MySQL, PostgreSQL)

 Syntax:

```
CASE
    WHEN condition1 THEN result1
    WHEN condition2 THEN result2
    ...
    ELSE default_result
END
```

 Example:

```
SELECT emp_name,
       salary,
       CASE
           WHEN salary >= 50000 THEN 'High'
           WHEN salary >= 30000 THEN 'Medium'
           ELSE 'Low'
       END AS salary_level
FROM employees;
```

 Output:

EMP_NAME	SALARY	SALARY_LEVEL
Raj	60000	High
Amit	35000	Medium
Ravi	20000	Low

◆ 2. COALESCE:

- Returns the first non-null value from the list.
- It checks each value in order and returns the first one that is not NULL.
- Works in all DBMSs

 Syntax: **COALESCE(value1, value2, ..., valueN)**

 Example:

```
SELECT emp_name,
       COALESCE(phone, email, 'No Contact') AS contact_info
FROM employees;
```

 Output:

EMP_NAME	CONTACT_INFO
Raj	9876543210
Amit	amit@mail.com
Ravi	No Contact

- ◆ Returns the first non-null value among the list.

- ◆ 3. NULLIF:
 - Returns NULL if two expressions are equal, else returns the first one.
 - Works in all DBMSs

 Syntax: `NULLIF(value1, value2)`

 Example:

```
SELECT emp_name,
       NULLIF(dept_id, 10) AS updated_dept
    FROM employees;
```

 Output:

EMP_NAME	UPDATED_DEPT
Raj	20
Amit	NULL
Ravi	30

- ◆ If both values are equal, returns `NULL`, else returns `value1`.

- ◆ 4. DECODE:

- Oracle-Specific CASE Alternative
- Acts like a simple `IF-THEN-ELSE`, similar to CASE but Oracle-specific.
- Only in Oracle

 Syntax:

```
DECODE(expression, search1, result1, search2, result2, ..., default)
```

 Example:

```
SELECT emp_name,
       DECODE(dept_id,
              10, 'Finance',
              20, 'HR',
              30, 'IT',
              'Unknown') AS department
    FROM employees;
```

 Output:

EMP_NAME	DEPARTMENT
Raj	HR
Amit	Finance
Ravi	IT

- ◆ Works like `CASE`, but shorter syntax. Limited to equality checks.

- ◆ 5. NVL Function (Oracle Only)

- Support in Oracle only.
- Replaces `NULL` with a specified value.

 Syntax: `NVL(expr1, expr2)`

 Example:

```
SELECT NVL(commission_pct, 0) AS commission FROM employees;
```

 Output: If commission is `NULL`, it will display `0`.

- ◆ 6. **NVL2** Function (Oracle Only):
 - Returns one value if the expression is NOT NULL, another if it is NULL.
 - Support in Oracle only

 Syntax: `NVL2(expr1, value_if_not_null, value_if_null)`

 Example:

```
SELECT NVL2(commission_pct, 'Has Bonus', 'No Bonus') FROM employees;
```

 Output:

commission_pct	result
NULL	No Bonus
0.3	Has Bonus

Important Notes:

- Use **CASE** for complex, multi-condition logic.
- Use **COALESCE** to replace NULLs with fallback values.
- Use **NULLIF** to suppress equal values by converting to NULL.
- Use **DECODE** for simpler conditions in Oracle, but prefer **CASE** for cross-platform compatibility.
- All are Scalar Functions — return one value per row.

Summary: Conditional Functions

Function	Purpose	Supported In DBMS
CASE	IF-ELSE style logic	 All
COALESCE	Return first non-NULL value	 All
NULLIF	Return NULL if both values are equal	 All
DECODE	Oracle-specific conditional replacement	 Oracle Only
NVL	Replace NULL with value	 Oracle Only
NVL2	Two results based on NULL or NOT NULL	 Oracle Only

◊ 3. Window (Analytical) Functions:

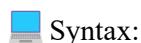
◆ What are Window Functions?

- Windows functions in Oracle SQL are powerful tools.
- perform calculations across a set of rows related to the current rows without collapsing the result. Set into a single value like aggregate functions do.
- Unlike aggregate functions, they do not collapse rows into a single result.
- Useful for ranking, running totals, moving averages, etc.

◆ What is a Window?

- A window is a defined set of rows that the function can access to perform its calculation.
- Specified using the `OVER()` clause.

- ◆ PARTITION BY Clause:
 - Divides the result set into groups (partitions).
 - Each partition is treated independently by the window function.
 - If we don't use Partition By the entire result set is treated as a single partition
- ◆ ORDER BY Clause (Inside OVER())
 - Determines the order in which rows are processed within each partition.
 - Essential for functions like RANK(), LEAD(), LAG(), etc.
- ◆ There is no official division of the SQL window functions into categories but high level we can divide into 3 types



Syntax:

```

function_name (expression) OVER (
  [PARTITION BY column]
  [ORDER BY column]
  [ROWS BETWEEN ...]
)
  
```

- Let's go one-by-one with examples, syntax, and outputs:
We'll use this sample `sales` table:

EMP_ID	EMP_NAME	DEPT	SALARY	DOJ
101	Aakash	HR	50000	2022-01-15
102	Neha	IT	60000	2021-06-10
103	Amit	HR	50000	2020-03-20
104	Riya	IT	75000	2022-11-25
105	Raj	SALES	62000	2019-09-01
106	Priya	SALES	62000	2020-01-10
107	Arjun	IT	60000	2023-04-03
108	Sneha	HR	80000	2021-02-17
109	Vinay	IT	60000	2022-07-19
110	Kiran	SALES	55000	2022-08-30

- Aggregate Functions:

- Windows aggregate functions are the set of functions that allows you to perform calculations across a set of rows that are related to the current row.
- Unlike regular aggregate functions which group data and return a single result per group, windows functions can return multiple results within a partition and ordered data set

💡 What's the difference between normal aggregate functions and window aggregate functions?

- Normal aggregates (like **GROUP BY**) collapse rows into 1 per group.
- Window aggregates return a result per row while calculating across a logical group.

1. SUM()

- Sums values within a window.
- Returns the total (sum) of values over the window of rows.
- Example: Total salary for each department.

💡 Example:

```
SELECT  
EMP_ID, EMP_NAME, DEPT, SALARY,  
SUM(SALARY) OVER (PARTITION BY DEPT) AS DEPT_TOTAL  
FROM EMPLOYEE;
```

2. AVG()

- Calculates the average value within a window.
- Useful for comparing each row with department average.

💡 Example:

```
SELECT  
EMP_ID, EMP_NAME, DEPT, SALARY,  
AVG(SALARY) OVER (PARTITION BY DEPT) AS DEPT_AVG  
FROM EMPLOYEE;
```

3. MIN()

- The MIN() function returns the minimum value from a specified column within the window of rows defined by the partition by and order by clauses.
- Used to find minimum marks, dates, etc., within a group.

💡 Example:

```
SELECT  
EMP_ID, EMP_NAME, DEPT, SALARY,  
MIN(SALARY) OVER (PARTITION BY DEPT) AS DEPT_MIN  
FROM EMPLOYEE;
```

4. MAX()

- The MAX() function returns the maximum value from a specified column within the window of rows defined by the partition by and order by clauses.
- Helps find max salary, score, or date in a group.

💡 Example:

```
SELECT  
EMP_ID, EMP_NAME, DEPT, SALARY,  
MAX(SALARY) OVER (PARTITION BY DEPT) AS DEPT_MAX  
FROM EMPLOYEE;
```

5. COUNT()

- The COUNT() window function counts the number of rows within the window of rows

💡 Example:

```

SELECT
EMP_ID, EMP_NAME, DEPT, SALARY,
COUNT(*) OVER (PARTITION BY DEPT) AS DEPT_COUNT
FROM EMPLOYEE;

```

 Output:

EMP_ID	EMP_NAME	DEPT	SALARY	DEPT_SUM	DEPT_AVG	DEPT_MIN	DEPT_MAX	DEPT_COUNT
101	Aakash	HR	50000	180000	60000	50000	80000	3
103	Amit	HR	50000	180000	60000	50000	80000	3
108	Sneha	HR	80000	180000	60000	50000	80000	3
102	Neha	IT	60000	255000	63750	60000	75000	4
104	Riya	IT	75000	255000	63750	60000	75000	4
107	Arjun	IT	60000	255000	63750	60000	75000	4
109	Vinay	IT	60000	255000	63750	60000	75000	4
105	Raj	SALES	62000	179000	59666.67	55000	62000	3
106	Priya	SALES	62000	179000	59666.67	55000	62000	3
110	Kiran	SALES	55000	179000	59666.67	55000	62000	3

◊ Ranking Functions:

- Window Ranking Functions are a type of analytic (window) function in SQL that assigns a rank or position to each row within a partition of the result set.
- Operate within a specified partition (group of rows) defined by **PARTITION BY**.
- Order rows using **ORDER BY** inside the **OVER()** clause.
- Unlike aggregate functions, they don't collapse rows — they return one row per input row with an additional calculated value (rank/position).

📦 Use Cases:

- Ranking employees by salary within each department.
- Finding top-N performers.
- Dividing data into quartiles or percentiles.
- Detecting ties in values and assigning ranks accordingly.

1. **ROW_NUMBER()** –

- Assigns a unique sequential number to each row in a partition of a result set, starting with 1 for the first row in each partition.
- Use: To get unique row IDs, top-N rows per group.

💡 Example:

```
SELECT
    EMP_ID, EMP_NAME, DEPT, SALARY,
    ROW_NUMBER() OVER (PARTITION BY DEPT ORDER BY SALARY DESC) AS ROW_NUM
FROM EMPLOYEE;
```

2. **RANK()** –

- Assigns a rank to each row within a partition of a result set.
- Rows with equal values (Duplicate) receive the same rank, and the next rank after a tie is incremented by the number of tied rows.
- Use: Ranking with gaps.
- As the name suggests the rank function assigns rank to all the rows within every partition rank is assigned such that rank 1 given to the first row and rows having the same values are assigned the same rank.
- For the next rank after two same rank values one rank value will be skipped. For instance if two rows share rank1 the next row gets rank 3 not 2

💡 Example:

```
SELECT
    EMP_ID, EMP_NAME, DEPT, SALARY,
    RANK() OVER (PARTITION BY DEPT ORDER BY SALARY DESC) AS RANKING
FROM EMPLOYEE;
```

3. **DENSE_RANK()** –

- Like **RANK()**, but no gaps after duplicates.
- This function is similar to the **RANK()** function, but it handles ties differently when assigning ranks.
- The key difference is that ‘**DENSE_RANK()**’ does not leave gaps in the ranking sequence when there are ties. Instead it continues with the next consecutive rank.
- Use: Compact ranking.

💡 Example:

```
SELECT
    EMP_ID, EMP_NAME, DEPT, SALARY,
    DENSE_RANK() OVER (PARTITION BY DEPT ORDER BY SALARY
    DESC) AS DENSE_RANKING
FROM EMPLOYEE;
```

4. NTILE(n) –

- Divides the rows in an ordered partition into ‘N’ groups as evenly as possible and assigns a unique group member to each row.
- Number of Tiles: Specify how many tiles (groups) you want the rows to be divided into
- Distribution: The function divides the rows in each partition into the specified number of tiles. If the number of rows does not divide evenly some tiles will have one more row than others.
- Use: Bucketizing for quartiles, deciles, etc.

💡 Example:

```
SELECT
    EMP_ID, EMP_NAME, DEPT, SALARY,
    NTILE(2) OVER (PARTITION BY DEPT ORDER BY SALARY DESC)
    AS NTILE_2
FROM EMPLOYEE;
```

5. PERCENT_RANK()

- Calculates the relative rank (0 to 1) of each row within a partition.
- Use: To find percentile-like values.

💡 Example:

```
SELECT
    EMP_ID, EMP_NAME, DEPT, SALARY,
    ROUND(PERCENT_RANK() OVER (PARTITION BY DEPT ORDER BY
    SALARY DESC), 2) AS PERCENT_RANKING
FROM EMPLOYEE;
```

💻 Output:

EMP_ID	EMP_NAME	DEPT	SALARY	ROW_NUM	RANKING	DENSE RANKING	NTILE_2	PERCENT RANKING
108	Sneha	HR	80000	1	1	1	1	0.00
101	Aakash	HR	50000	2	2	2	2	0.5
103	Amit	HR	50000	3	2	2	2	0.5
104	Riya	IT	75000	1	1	1	1	0.00
102	Neha	IT	60000	2	2	2	1	0.33
107	Arjun	IT	60000	3	2	2	2	0.33
109	Vinay	IT	60000	4	2	2	2	0.33
105	Raj	SALES	62000	1	1	1	1	0.00
106	Priya	SALES	62000	2	1	1	1	0.00

110	Kiran	SALES	55000	3	3	2	2	0.66
-----	-------	-------	-------	---	---	---	---	------

◊ Value Functions:

- Window Value Functions return a value from another row in the same result set, relative to the current row — without collapsing the data.
- They are very useful for comparing rows (e.g., previous, next, first, last) within a partition.

1.LEAD()

- Returns the value from the next row in the window.
- Useful for comparing the current row with the next one.

💡 Example:

```
SELECT
    EMP_ID, EMP_NAME, DEPT, SALARY,
    LEAD(SALARY) OVER (PARTITION BY DEPT ORDER BY SALARY)
    AS NEXT_SALARY
FROM employees;
```

2.LAG()

- Returns the value from the previous row.
- Used for difference calculations like sales vs. previous month.

💡 Example:

```
SELECT
    EMP_ID, EMP_NAME, DEPT, SALARY,
    LAG(SALARY) OVER (PARTITION BY DEPT ORDER BY SALARY)
    AS PREV_SALARY
FROM employees;
```

3.FIRST_VALUE()

- Returns the first value in the window frame.
- Helps track the starting value of a group.

💡 Example:

```
SELECT
    EMP_ID, EMP_NAME, DEPT, SALARY,
    FIRST_VALUE(SALARY) OVER (PARTITION BY DEPT ORDER BY
    SALARY) AS FIRST_SALARY
FROM employees;
```

4.LAST_VALUE()

- Returns the last value in the window frame.
- Tracks the ending value of a partition.

💡 Example:

```
SELECT
    EMP_ID, EMP_NAME, DEPT, SALARY,
    LAST_VALUE(SALARY) OVER (PARTITION BY DEPT ORDER BY
    SALARY)
```

```

ROWS BETWEEN UNBOUNDED PRECEDING AND UNBOUNDED
FOLLOWING) AS LAST_SALARY
FROM employees;

```

5.NTH_VALUE()

- Returns the nth value (e.g., 2nd, 3rd...) in the partition/order of the result set.
- If $n = 2$, it gives the second value in the frame.

 Example:

```

SELECT
    EMP_ID, EMP_NAME, DEPT, SALARY,
    NTH_VALUE(SALARY, 2) OVER(PARTITION BY DEPT ORDER BY SALARY
    ROWS BETWEEN UNBOUNDED PRECEDING AND UNBOUNDED
    FOLLOWING) AS SECOND_SALARY
FROM employees;

```

 Output:

EMP_ID	EMP_NAME	DEPT	SALARY	PREV_SALARY	NEXT_SALARY	FIRST_SALARY	LAST_SALARY	SECOND_SALARY
103	Amit	HR	50000	null	50000	50000	80000	50000
101	Aakash	HR	50000	50000	80000	50000	80000	50000
108	Sneha	HR	80000	50000	null	50000	80000	50000
102	Neha	IT	60000	null	60000	60000	75000	60000
107	Arjun	IT	60000	60000	60000	60000	75000	60000
109	Vinay	IT	60000	60000	75000	60000	75000	60000
104	Riya	IT	75000	60000	null	60000	75000	60000
110	Kiran	SALES	55000	null	62000	55000	62000	62000
105	Raj	SALES	62000	55000	62000	55000	62000	62000
106	Priya	SALES	62000	62000	null	55000	62000	62000

 Important Notes:

- **PARTITION BY**: Like GROUP BY inside a window. It splits the data into partitions to apply the function.
- **ORDER BY**: Sorts rows within each partition.
- **ROWS BETWEEN**: Specifies the range of rows in a frame.
- Window functions do not reduce rows, unlike GROUP BY. They add an extra column with calculated values.

Summary: Windows Functions

Function	Description
SUM()	Running total of values over a partition
AVG()	Running average of values over a partition
MIN()	Minimum value in the current window or partition
MAX()	Maximum value in the current window or partition
COUNT()	Counts rows or non-null values over a partition
ROW_NUMBER()	Unique sequential number for each row per partition
RANK()	Rank with gaps
DENSE_RANK()	Rank without gaps
NTILE(n)	Distributes rows into n buckets
PERCENT_RANK()	Percentile rank of the row in the partition
LEAD()	Next row value
LAG()	Previous row value
FIRST_VALUE()	First value in partition
LAST_VALUE()	Last value in partition
NTH_VALUE()	Nth row's value in the partition

◊ 4. System Functions:

- System functions in SQL are built-in scalar functions used to retrieve:
 - System-level details (like current user, DB name, server)
 - Session-related info
 - Metadata about database objects
- They do not need table data and often work without FROM clauses.

Common System Functions (with Support):

◆ 1. **CURRENT_USER**

- Returns: The database username that is executing the current SQL command.
- Useful for: Tracking who is executing which SQL logic in multi-user environments.

 Example: **SELECT CURRENT_USER;**

 Output: '**VAIBHAV_USER**'

◆ 2. **SESSION_USER**

- Returns: The name of the user who opened the session, which may differ from the user who is executing the query (**CURRENT_USER**).
- Useful for: Understanding session context vs. current query user.

 Example: **SELECT SESSION_USER;**

 Output: '**ADMIN**'

- ◆ 3. **SYSTEM_USER**
 - Returns: The system login name of the user running the SQL session (e.g., Windows or OS-level login).
 - Useful for: Security auditing and OS-level user tracking.

 Example: `SELECT SYSTEM_USER;`

 Output: 'DESKTOP-VAIBHAV\Vaibhav'

- ◆ 4. **CURRENT_DATE**
 - Returns: The current system date (only the date part).
 - Useful for: Logging, date comparisons, filtering current-day records.

 Example: `SELECT CURRENT_DATE;`

 Output: 2025-05-23

- ◆ 5. **CURRENT_TIME**
 - Returns: The current system time (only time part).
 - Useful for: Time-based operations like logging or session tracking.

 Example: `SELECT CURRENT_TIME;`

 Output: 15:07:58

- ◆ 6. **CURRENT_TIMESTAMP**
 - Returns: The current system date and time (full timestamp).
 - Useful for: Creating timestamps in audit tables or logging.

 Example: `SELECT CURRENT_TIMESTAMP;`

 Output: 2025-05-23 15:07:58.123

- ◆ 7. **SYSDATE()** (*Oracle/MySQL specific*)
 - Returns: The system's current date and time.
 - Useful for: Same use cases as **CURRENT_TIMESTAMP**, but behavior differs slightly across DBMS.

 Example: `SELECT SYSDATE FROM DUAL;` -- Oracle

-- or

 `SELECT SYSDATE();` -- MySQL

 Output: 2025-05-23 15:08:30

- ◆ 8. **DB_NAME()** (*SQL Server only*)
 - Returns: The name of the current database.
 - Useful for: Dynamic database name retrieval in scripts.

 Example: `SELECT DB_NAME();`

 Output: 'StudentDB'

- ◆ 9. **DATABASE()** (*MySQL/PostgreSQL*)
 - Returns: Name of the currently selected database.
 - Useful for: Knowing which DB you're working in, especially in dynamic environments.

 Example: `SELECT DATABASE();`

Output: 'hr_management'

◆ 10. **VERSION()**

- Returns: The version of the current SQL engine.
- Useful for: Debugging or compatibility checks across environments.

💡 Example: `SELECT VERSION();`

Output: 'PostgreSQL 15.2 (Ubuntu 20.04)'

◆ 11. **HOST_NAME()** (*SQL Server only*)

- Returns: The host machine name where the SQL Server is running.
- Useful for: Knowing where your DB engine is hosted.

💡 Example: `SELECT HOST_NAME();`

Output: 'DESKTOP-VAIBHAV'

◆ 12. **@@SERVERNAME** (*SQL Server only*)

- Returns: Name of the SQL Server instance.
- Useful for: Tracking server names in multi-server environments.

💡 Example: `SELECT @@SERVERNAME;`

Output: 'SQLPROD01'

📝 Important Notes:

- System functions help you interact with user, session, server, time, and environment details.
- No table is needed for these queries.
- Useful in auditing, logging, and debugging and understanding system behavior.
- SQL Server provides a lot of system functions using `@@`.

⌚ Summary : System Functions

Function	Purpose	DBMS Supported
<code>CURRENT_USER</code>	User running the query	All
<code>SESSION_USER</code>	Session creator	PostgreSQL, SQL Server
<code>SYSTEM_USER</code>	OS/network login	SQL Server, PostgreSQL
<code>CURRENT_DATE</code>	System date	All
<code>CURRENT_TIME</code>	Current time	All
<code>CURRENT_TIMESTAMP</code>	Date + time	All
<code>SYSDATE()</code>	Oracle/MySQL system date	Oracle, MySQL
<code>DATABASE()</code>	Name of the current DB	MySQL, PostgreSQL
<code>DB_NAME()</code>	Name of current DB	SQL Server
<code>VERSION()</code>	DB engine version	All

HOST_NAME()	Host machine name	SQL Server
@@SERVERNAME	Name of SQL server	SQL Server

◊ 5. XML Functions:

◆ What Are XML Functions in SQL?

◆ XML functions allow you to:

- Parse XML data.
- Extract values or nodes using XPath.
- Convert SQL data to XML format.
- Manipulate, transform, or store XML as structured data inside relational databases.

Use Case Example:

A company stores employee info in a column of type **XML**. You want to extract the **<name>** tag or split the **<skills>** into rows — XML functions help you do this.

◆ XML Functions in Oracle SQL:

- Oracle treats XML as a first-class data type using **XMLTYPE**.

◆ 1. **XMLTYPE()**:

- Converts a string (VARCHAR, CLOB) to XML format so it can be queried using XML functions.

💡 Example:

```
SELECT XMLTYPE('<emp><name>Vaibhav</name></emp>') AS xml_col FROM dual;
```

📌 *Why it's useful:* You can now use XPath queries on it.

◆ 2. **EXTRACT() / EXTRACTVALUE()**: Extracts a part of the XML using XPath.

💡 Example:

```
SELECT EXTRACT(XMLTYPE('<emp><name>Vaibhav</name></emp>'), '/emp/name/text()') AS emp_name FROM dual;
```

💻 Output: Vaibhav

📌 *Use it when:* You want a specific value (like a tag's text).

◆ 3. **XMLTABLE()**:

- Converts XML data into a virtual table so that you can SELECT rows/columns.

💡 Example:

```
SELECT *
FROM XMLTABLE('/employees/employee'
  PASSING
  XMLTYPE('<employees><employee><id>1</id></employee></employees>')
  COLUMNS emp_id NUMBER PATH 'id');
```

💻 Output:

Emp_id

1

📌 *Use it when:* You want to treat XML data as a regular table and extract multiple rows.

◆ 4. XMLSERIALIZE():

- Converts XMLTYPE back into VARCHAR or CLOB to view as plain text.

💡 Example:

```
SELECT XMLSERIALIZE(DOCUMENT XMLTYPE('<a>1</a>') AS VARCHAR2(50))
AS plain_xml FROM dual;
```

💻 Output: <a>1

◆ XML Functions in SQL Server:

- SQL Server stores XML in columns of type XML.

◆ 1. .value():

- Extract a single scalar value using XPath.

💡 Example:

```
SELECT xml_data.value('/employee/name')[1], 'VARCHAR(100)' AS emp_name
FROM employees;
```

💻 Output: Vaibhav

📌 Use it when: You want just one tag's value.

◆ 2. .query(): Returns an XML fragment.

💡 Example: `SELECT xml_data.query('/employee/name') AS name_xml
FROM employees;`

💻 Output: <name>Vaibhav</name>

◆ 3. .nodes(): Break XML into multiple rows (nodes).

💡 Example:

```
SELECT x.value('(name)[1]', 'VARCHAR(50)') AS emp_name
FROM employees
CROSS APPLY xml_data.nodes('/employees/employee') AS T(x);
```

💻 Output:

Emp_name

Vaibhav

Akash

📌 Use it when: You want to normalize XML into rows.

◆ 4. FOR XML: Converts SQL query results into an XML structure.

💡 Example: `SELECT emp_id, emp_name FROM employees FOR XML AUTO;`

💻 Output:

```
<employees emp_id="1" emp_name="Vaibhav" />
<employees emp_id="2" emp_name="Akash" />
```

📌 Use it when: You want to export data in XML format.

◊ XML Functions in PostgreSQL:

- PostgreSQL uses built-in XML data type and XPath queries.

- ◆ 1. `xpath()`: Returns an array of values matching an XPath query.

 Example:

```
SELECT xpath('/employee/name/text()', xml
            '<employee><name>Vaibhav</name></employee>');

```

 Output: {Vaibhav}

 Use it when: You want XPath results from an XML column.

Important Notes:

- XML is mostly used in Oracle, SQL Server, and PostgreSQL.
- MySQL does not support advanced XML functions — use JSON instead.
- XML is great for interoperability, especially in web APIs and metadata.
- Learn [XPath](#) to query XML more powerfully.
- XML data can be complex to process — it's often better stored as relational if possible.

Summary : XML Functions

Function	Oracle	SQL Server	PostgreSQL	Description
<code>XMLTYPE()</code>	✓	✗	✗	Converts to XML type
<code>EXTRACT()</code>	✓	✗	✗	Extracts XML node using XPath
<code>XMLTABLE()</code>	✓	✗	✗	Converts XML to rows
<code>XMLSERIALIZE()</code>	✓	✗	✗	Converts XML to string
<code>.value()</code>	✗	✓	✗	Gets a single value from XML
<code>.query()</code>	✗	✓	✗	Returns XML node
<code>.nodes()</code>	✗	✓	✗	Converts XML array into rows
<code>FOR XML</code>	✗	✓	✗	Converts result set to XML
<code>xpath()</code>	✗	✗	✓	Gets values using XPath in array

◆ VIEWS:

◆ What is a View in SQL?

- A View is a virtual table in SQL that is created by writing a **SELECT** query and saving it with a name.
- It doesn't store data physically — it just shows the result of the SELECT query when you use it.

◆ Real-life Analogy:

- Think of a view like a lens or a filter placed over your data.
- You can choose what data to display, how to format it, and even hide sensitive information — all without touching the original tables.

▀ Syntax

```
CREATE VIEW view_name AS  
    SELECT column1, column2  
    FROM table_name  
    WHERE condition;
```

❖ TYPES OF VIEWS (with Multi-line Explanations)

1. Simple View
2. Complex View
3. Materialized View
4. Updatable View
5. Non-Updatable View

❖ 1. Simple View

- A simple view is based on only one table.
- It does not contain any of the following: **GROUP BY**, **JOIN**, **DISTINCT**, **subqueries**, or aggregate functions.
- It is often used to limit access to specific columns.
- Can be updated easily if it includes only base columns.

💡 Example

```
CREATE VIEW vw_EmployeeNames AS  
    SELECT Emp_ID, Emp_Name FROM Employees;
```

❖ 2. Complex View

- A complex view is based on multiple tables and can include **JOINS**, **GROUP BY**, aggregate functions, and more.
- It is often used to combine data from different tables or summarize data.
- Usually not updatable, especially when it includes groupings or aggregates.

💡 Example

```
CREATE VIEW vw_DeptEmployeeCount AS  
    SELECT Dept_ID, COUNT(*) AS EmployeeCount  
    FROM Employees  
    GROUP BY Dept_ID;
```

❖ 3. Materialized View

- A materialized view is different from a normal view.
- It physically stores the query result, which improves performance for heavy aggregations.
- Data inside a materialized view does not automatically update. It must be refreshed manually or on a schedule.

- Mostly used in data warehouses for reporting and analytics.

 Example (Oracle/PostgreSQL):

```
CREATE MATERIALIZED VIEW mv_EmployeeSummary AS
SELECT Dept_ID, COUNT(*) AS EmpCount
FROM Employees
GROUP BY Dept_ID;
```

- ◆  To refresh:

```
REFRESH MATERIALIZED VIEW mv_EmployeeSummary;
```

◊ 4. Updatable View

- An updatable view is one that allows **INSERT**, **UPDATE**, and **DELETE** operations.
- Usually based on a single base table and contains no aggregates or joins.
- Columns in the view should map directly to columns in the table.

 Example:

```
CREATE VIEW vw_Salaries AS
SELECT Emp_ID, Emp_Salary FROM Employees;
```

You can then run:

```
UPDATE vw_Salaries SET Emp_Salary = 80000 WHERE Emp_ID = 101;
```

 This will update the base table.

◊ 5. Non-Updatable View

- A non-updatable view does not support DML operations.
- This includes views that:
 - Use **GROUP BY**, **DISTINCT**, **JOINS**
 - Include aggregate functions like **SUM()**, **AVG()**
- Such views are useful for read-only reports or dashboards.

 Example: (non-updatable):

```
CREATE VIEW vw_DeptAvgSalary AS
SELECT Dept_ID, AVG(Emp_Salary) AS AvgSalary
FROM Employees GROUP BY Dept_ID;
```

Trying to update this view will result in an error.

- ◆  Additional Concepts

- ◆  Modify a View

```
CREATE OR REPLACE VIEW vw_EmployeeNames AS
SELECT Emp_ID, Emp_Name FROM Employees;
```

❖  Drop a View

```
DROP VIEW vw_EmployeeNames;
```

- ◆ View Metadata (See View Definition)

- ◆ SQL Server:

```
sp_helptext 'vw_EmployeeNames';
```

- ◆ MySQL:

```
SHOW CREATE VIEW vw_EmployeeNames;
```

◆ Oracle:

```
SELECT TEXT FROM USER_VIEWS WHERE VIEW_NAME =  
'VW_EMPLOYEENAMES';
```

💡 Important Notes:

- Views help organize and simplify your SQL queries.
- They improve security by showing only selected columns.
- Views don't store data → they use base table data every time you query them.
- Use Materialized Views to improve performance for large data sets.
- Always check if a view is updatable before trying DML operations.
- Best practice: Use views to isolate business logic and protect data.

📋 Summary : Views in SQL

View Type	Description
Simple View	Based on one table, no aggregates or joins. Often updatable and simple.
Complex View	Based on multiple tables, often includes joins or aggregates. Read-only.
Materialized View	Stores actual data. Used for performance optimization. Needs to refresh.
Updatable View	Allows INSERT/UPDATE/DELETE. Based on one table, no aggregates.
Non-Updatable View	Read-only. Includes joins, groups, functions. Can't modify the base table.

◆ Temporary Tables

◆ What is a Temporary Table?

- A Temporary Table is a special kind of table that stores data temporarily during a session or transaction.
- It behaves just like a normal table, but it gets deleted automatically once the session ends or the transaction completes.
- These tables are useful in scenarios like:
 - Storing intermediate results
 - Complex calculations
 - Simplifying stored procedures
 - Reducing redundant calculations in large queries

❖ Types of Temporary Tables:

❖ 1. Local Temporary Table (SQL SERVER)

- It is specific to the current user session.
- Other users or sessions cannot access it.
- The table is automatically dropped once the session ends or connection is closed.
- In SQL Server, it starts with a single #.

💻 Syntax (SQL Server):

```
CREATE TABLE #LocalTemp (
    ID INT,
    Name VARCHAR(50)
);
```

💡 Example:

```
INSERT INTO #LocalTemp VALUES (1, 'Vaibhav'), (2, 'Amit');

SELECT * FROM #LocalTemp;
```

💻 Output:

ID	Name
1	Vaibhav
2	Amit

❖ 2. Global Temporary Table (SQL SERVER)

- It is accessible to all sessions/users.
- The table is automatically deleted only when all sessions that use it are closed.
- In SQL Server, it starts with double ##.

💻 Syntax (SQL Server):

```
CREATE TABLE ##GlobalTemp (
    ID INT,
    Status VARCHAR(20)
);
```

 Example:

```
INSERT INTO ##GlobalTemp VALUES (1, 'Active'), (2, 'Inactive');

SELECT * FROM ##GlobalTemp;
```

 Output:

ID	Status
1	Active
2	Inactive

 This table will be accessible from other connections too.

◊ 3. Temporary Table using **TEMP** or **TEMPORARY** Keyword (MySQL, PostgreSQL, Oracle)

- Works across many databases.
- Automatically dropped when session ends (similar to local temp).
- Use **CREATE TEMP TABLE** or **CREATE TEMPORARY TABLE**.

 Syntax (MySQL/PostgreSQL):

```
CREATE TEMPORARY TABLE TempOrders (
    OrderID INT,
    OrderAmount DECIMAL(10,2)
);
```

 Example:

```
INSERT INTO TempOrders VALUES (101, 500.00), (102, 750.50);

SELECT * FROM TempOrders;
```

 Output:

OrderID	OrderAmount
101	500.00
102	750.50

◊ 4. ON COMMIT DELETE ROWS / PRESERVE ROWS (PostgreSQL Specific)

- Used in PostgreSQL for better control over when rows are deleted from a temporary table.
- You can choose whether data should persist or be deleted after commit.

 Options:

- **ON COMMIT DELETE ROWS** → clears data after each commit.
- **ON COMMIT PRESERVE ROWS** → keeps data until session ends.

 Example:

```
CREATE TEMP TABLE temp_sales (
    sale_id INT,
    amount DECIMAL
) ON COMMIT DELETE ROWS;
```

- ◊ 5. Common Table Expressions (CTE) – Temporary Result Set
 - Not an actual table, but acts like a temporary table during the execution of a single query.
 - It is like a named subquery that improves readability and modularity of complex SQL queries.

 Syntax

```
WITH TempCTE AS (
    SELECT Emp_ID, Emp_Salary FROM Employees WHERE Dept_ID = 2
)
SELECT * FROM TempCTE WHERE Emp_Salary > 60000;
```

 Modify Temp Table Structure

- You can modify temp tables using standard **ALTER TABLE** commands if needed.
`ALTER TABLE #LocalTemp ADD Email VARCHAR(100);`

 Dropping Temporary Tables

- Temp tables are usually auto-dropped, but you can drop them manually too:
`DROP TABLE #LocalTemp;`

 Important Notes:

- Indexes, constraints, and primary keys can be used on temporary tables.
- Temporary tables are often used in stored procedures.
- You cannot use temp tables between sessions unless it's a global temp table.
- In PostgreSQL/MySQL, avoid naming temp tables the same as base tables — they can shadow them.
- Don't forget to drop temp tables manually in long-running apps to save memory.

 Summary : Regular Expression

Type	Description
Local Temp Table	Exists only during the current session. Named with #.
Global Temp Table	Visible to all sessions. Deleted when all sessions close. Named with ##.
TEMPORARY Table (MySQL)	Temporary table, auto-deleted at session end. Uses TEMP keyword.
ON COMMIT Options	PostgreSQL-specific control of row persistence in TEMP tables.
CTE	Temporary result set used only during query execution.

◆ CTE (Common Table Expression)

- A Common Table Expression (CTE) is a temporary result set defined within the execution scope of a single **SELECT**, **INSERT**, **UPDATE**, or **DELETE** statement.
 - It is like a named subquery that improves readability and modularity of complex SQL queries.
 - CTEs are defined using the **WITH** keyword followed by a query block, which you can reference just like a table in the main SQL query.
- ◆ Does CTE get stored in database?
- ✖ No.
- CTE does not store data (like a table).
 - CTE does not store query definition (like a view).
 - It's just like a temporary alias for a subquery.
 - Once the query finishes execution → the CTE disappears
- ◆ Key Features of CTE:

1. Improves Readability
 - Breaks down large queries into logical blocks.
 - Makes SQL cleaner and easier to understand.
2. Reusable in the Same Query
 - Once defined, a CTE can be used multiple times within the main query.
3. Supports Recursion
 - Allows writing recursive queries like hierarchical structures (e.g., org charts, folder paths).
4. Scoped to a Single Query
 - CTEs exist only during the execution of the query.
 - Not stored permanently like a view or table.

💻 Syntax:

```
WITH cte_name AS (
    SELECT ...
)
SELECT * FROM cte_name;
```

💡 Example:

```
WITH high_salary AS (
    SELECT employee_id, salary FROM employees WHERE salary > 50000
)
SELECT * FROM high_salary;
```

💻 Output: Returns employees with salaries above 50,000.

◆ WITH Clause:

- The **WITH** clause is used to define a CTE.
- It acts like a temporary table that can be reused in your query.
- It simplifies complex SQL queries by breaking them into simpler pieces.
- The **WITH** clause is not separate from CTE—it's the keyword used to define a CTE.

☞ Important Notes:

- Exists only during query execution.
- Can be recursive (for hierarchical data).

◇ PIVOT

- The **PIVOT** operator turns row data into columns.
- It is helpful for generating cross-tab reports.
- It helps to summarize and transform data for better reporting.
- You use PIVOT when you want to turn category values (like months, years, or status) into separate columns.

🧠 Think of it like:

- You're converting data from vertical (rows) into horizontal (columns).

💻 Syntax (Oracle):

```
SELECT * FROM (
    SELECT department, year, sales FROM sales_data
)
PIVOT (
    SUM(sales) FOR year IN (2022, 2023)
);
```

💡 Example:

Department	2022	2023
HR	5000	6000
IT	7000	9000

☞ Important Notes:

- Only supported in Oracle and SQL Server.
- Helps in reporting formats.

◇ UNPIVOT

- The **UNPIVOT** operator turns columns into rows.
- It is the reverse of **PIVOT**.
- It converts columns into rows — useful for normalizing wide tables.

🧠 Think of it like:

- You're taking horizontal data (columns) and turning it vertical (rows).

💻 Syntax (Oracle):

```
SELECT * FROM sales_data
UNPIVOT (
    sales FOR year IN (sales_2022 AS '2022', sales_2023 AS '2023')
);
```

💡 Example:

Department	Year	Sales
HR	2022	5000
HR	2023	6000

Important Notes:

- Very useful for data normalization or ETL jobs.
- All these functions (except UNPIVOT and PIVOT) are supported in Oracle, SQL Server, PostgreSQL, and MySQL 8+.
- Use analytical functions with `OVER()` clauses.
- `LEAD()` and `LAG()` are non-aggregate functions used with windowing.
- Recursive CTEs are supported in Oracle, SQL Server, PostgreSQL but not in MySQL (until version 8+).
- Pivoting is DBMS-specific—standard SQL doesn't support it natively.

Summary :

Feature	Description	Supported In
CTE / WITH	Temporary result set used in query	Oracle, SQL Server, PostgreSQL, MySQL 8+
PIVOT	Converts rows to columns	Oracle, SQL Server
UNPIVOT	Converts columns to rows	Oracle, SQL Server

◇ Recursive CTE in SQL

- A Recursive CTE is a CTE (Common Table Expression) that calls itself to return multiple levels of data.
 - A Recursive CTE is a powerful SQL feature that allows a query to refer to itself.
 - It's especially useful when working with hierarchical or tree-structured data (e.g., employees reporting to managers, category trees, etc.).
- ◆ Structure:
- A recursive CTE has two parts:
- Anchor Member: The base result (starting point)
 - Recursive Member: The part that references the CTE itself and builds on the result

💻 Syntax:

```
WITH RECURSIVE cte_name (column_list) AS (
    -- Anchor member
    SELECT ...
    FROM ...
    WHERE ...
    UNION ALL
    -- Recursive member
    SELECT ...
    FROM table_name
    JOIN cte_name ON ...
)
SELECT * FROM cte_name;
```

💡 Example: Employee Hierarchy

Let's say you have this employee table:

```
CREATE TABLE employees (
    emp_id INT,
    emp_name VARCHAR(50),
    manager_id INT
);
INSERT INTO employees VALUES
(1, 'CEO', NULL),
(2, 'Manager A', 1),
(3, 'Manager B', 1),
(4, 'Dev A1', 2),
(5, 'Dev A2', 2),
(6, 'Dev B1', 3);
```

⌚ Goal: Find the hierarchical structure of employees starting from the CEO.

💻 Recursive CTE Query:

💡 Example:

```
WITH RECURSIVE emp_hierarchy AS (
    -- Anchor member (CEO has no manager)
    SELECT emp_id, emp_name, manager_id, 1 AS level
    FROM employees
    WHERE manager_id IS NULL
)
SELECT emp_id, emp_name, manager_id, level
FROM emp_hierarchy
ORDER BY emp_id;
```

UNION ALL

```
-- Recursive member: find employees reporting to previous result
SELECT e.emp_id, e.emp_name, e.manager_id, h.level + 1
FROM employees e
JOIN emp_hierarchy h ON e.manager_id = h.emp_id
)
SELECT * FROM emp_hierarchy;
```

Output:

emp_id	emp_name	manager_id	level
1	CEO	NULL	1
2	Manager A	1	2
3	Manager B	1	2
4	Dev A1	2	3
5	Dev A2	2	3
6	Dev B1	3	3

Important Notes:

- **WITH RECURSIVE** is supported in PostgreSQL, MySQL 8+, SQL Server (using different style), and Oracle 11g+
- Use **UNION ALL** (not **UNION**) to avoid performance issues unless duplicates are a concern
- You must ensure the recursion terminates — else it causes infinite loop errors

Use Cases:

- Employee-manager hierarchy
- Bill of materials
- Category-subcategory tree
- Folder-directory structure
- Graphs (paths in networks)

Summary : Recursive CTE

Feature	Description
Anchor member	Base query (non-recursive part)
Recursive member	Refers back to the CTE to continue building result
Use case	Tree/hierarchy traversal
Limitation	Needs proper termination condition (avoid infinite loop)
SQL support	PostgreSQL ✓, MySQL 8+ ✓, SQL Server ✓ (CTE style), Oracle ✓

◆ MERGE Statement (also called UPSERT)

◆ What is MERGE?

- The **MERGE** statement is used to perform INSERT, UPDATE, or DELETE operations in a single SQL command, based on whether matching rows exist or not.
- It's also known as UPSERT = UPDATE + INSERT.
- Common in data warehousing, ETL, data sync, and slowly changing dimensions (SCDs).

◆ When to Use MERGE?

- You have a target table and a source table/data
- If a match is found → UPDATE existing row
- If no match is found → INSERT new row
- (Optional) If match is found and some condition → DELETE

Syntax

```
MERGE INTO target_table t
  USING source_table s
  ON (t.id = s.id)
  WHEN MATCHED THEN
    UPDATE SET t.col1 = s.col1, t.col2 = s.col2
  WHEN NOT MATCHED THEN
    INSERT (id, col1, col2)
    VALUES (s.id, s.col1, s.col2);
```

💡 Example

🎯 Scenario: We have two tables:

employees_target

emp_id	name	salary
101	Amit	30000
102	Vaibhav	40000

employees_source

emp_id	name	salary
102	Vaibhav R	45000
103	Rahul	35000

💡 Example:

```
MERGE INTO employees_target tgt
  USING employees_source src
  ON (tgt.emp_id = src.emp_id)
  WHEN MATCHED THEN
    UPDATE SET tgt.name = src.name, tgt.salary = src.salary
  WHEN NOT MATCHED THEN
    INSERT (emp_id, name, salary)
    VALUES (src.emp_id, src.name, src.salary);
```

Output:

 `employees_target` becomes:

emp_id	name	salary
101	Amit	30000
102	Vaibhav R	45000
103	Rahul	35000

Important Notes:

Feature	MERGE Behavior
Multiple operations	Yes – INSERT, UPDATE, DELETE in one go
Atomicity	Fully atomic – succeeds or fails together
Conditions	Can be added to both matched & unmatched
Performance	Better than writing separate queries

Notes:

- The `ON` clause is mandatory and defines how to match rows.
- `WHEN MATCHED` can have a condition like `AND salary <> src.salary`.
- `WHEN NOT MATCHED` only happens if no row is matched in the `target`.
- `MERGE` can update, insert, and even delete if conditions are added.
- Oracle, SQL Server, DB2 support `MERGE`.
- MySQL and PostgreSQL use `INSERT ... ON DUPLICATE KEY UPDATE` or `INSERT ... ON CONFLICT` instead.

Summary : MERGE Statement

Concept	Description
MERGE	Combines INSERT, UPDATE, DELETE based on matching
Use Case	Data sync, ETL jobs, deduplication
Synonyms	UPSERT
Supported DBs	<input checked="" type="checkbox"/> Oracle, SQL Server, DB2; <input checked="" type="checkbox"/> MySQL/PostgreSQL (use alternative syntax)

❖ UPSERT in PostgreSQL & MySQL

◆ PostgreSQL - `INSERT ... ON CONFLICT`

- PostgreSQL uses the `ON CONFLICT` clause to handle duplicates.

💻 Syntax:

```
INSERT INTO table_name (col1, col2, col3)
VALUES (val1, val2, val3)
ON CONFLICT (conflict_column)
DO UPDATE
SET col2 = EXCLUDED.col2,
    col3 = EXCLUDED.col3;
```

💡 Example:

Assume we have this table in PostgreSQL:

```
CREATE TABLE employees (
    emp_id INT PRIMARY KEY,
    name VARCHAR(50),
    salary INT
);
```

Then run:

```
INSERT INTO employees (emp_id, name, salary)
VALUES (102, 'Vaibhav R', 45000)
ON CONFLICT (emp_id)
DO UPDATE
SET name = EXCLUDED.name,
    salary = EXCLUDED.salary;
```

💻 Output:

emp_id	name	salary
101	Amit	30000
102	Vaibhav R	45000
103	Rahul	35000

◆ MySQL — `INSERT ... ON DUPLICATE KEY UPDATE`

- MySQL uses `ON DUPLICATE KEY UPDATE` to achieve the same functionality.

💻 Syntax:

```
INSERT INTO table_name (col1, col2, col3)
VALUES (val1, val2, val3)
ON DUPLICATE KEY UPDATE
col2 = VALUES(col2),
    col3 = VALUES(col3);
```

💡 Example:

Assume this MySQL table:

```
CREATE TABLE employees (
    emp_id INT PRIMARY KEY,
    name VARCHAR(50),
    salary INT
);
```

Then run:

```

INSERT INTO employees (emp_id, name, salary)
VALUES (102, 'Vaibhav R', 45000)
ON DUPLICATE KEY UPDATE
name = VALUES(name),
salary = VALUES(salary);

```

 Output:

Same output as before — updates existing or inserts new based on PRIMARY KEY.

 Important Notes:

Feature	PostgreSQL	MySQL
Clause	ON CONFLICT	ON DUPLICATE KEY UPDATE
Conflict Detection	Any UNIQUE or PRIMARY KEY	Only PRIMARY KEY or UNIQUE
Access to incoming values	Uses EXCLUDED.col_name	Uses VALUES(col_name) (deprecated, now use VALUES() in older versions)
Flexibility in conflict keys	Yes (you can choose conflict columns)	No (only keys defined in schema)

 Summary : Upsert

DBMS	UPSERT Syntax Equivalent
Oracle	MERGE INTO ...
SQL Server	MERGE INTO ...
PostgreSQL	INSERT ... ON CONFLICT ... DO UPDATE
MySQL	INSERT ... ON DUPLICATE KEY UPDATE

◆ Regular Expressions in SQL (REGEXP)

- ◆ What is a Regular Expression (Regex)?
 - Regular Expression (Regex) is a sequence of characters that defines a search pattern, often used for pattern matching, searching, and string manipulation.
 - help you perform advanced string matching beyond simple `LIKE` or `INSTR()`

💻 Syntax (General for MySQL & Oracle)

```
SELECT column  
FROM table  
WHERE column REGEXP 'pattern';
```

◆ Common REGEXP Operators:

Pattern	Meaning	Example
^	Start of string	<code>^A</code> (starts with A)
\$	End of string	<code>end\$</code> (ends with "end")
.	Any single character	<code>b.d</code> matches "bad", "bed", etc.
*	Zero or more of the preceding character	<code>ab*c</code> matches "ac", "abc", "abbc"
+	One or more of the preceding character	<code>ab+c</code> matches "abc", "abbc", but not "ac"
?	Zero or one of the preceding character	<code>ab?c</code> matches "abc" or "ac"
,		OR operator
[]	Match any character inside brackets	<code>[aeiou]</code> (vowels)
[^]	Not any character inside brackets	<code>[^0-9]</code> (not a digit)
{n}	Exactly n occurrences	<code>a{3}</code> matches "aaa"
{n,}	n or more occurrences	<code>a{2,}</code>
{n,m}	Between n and m occurrences	<code>a{2,4}</code>

- ◆ Examples (MySQL / Oracle)

1. Match Names Starting with "A"

💡 Example:

```
SELECT name  
FROM employees  
WHERE name REGEXP '^A';
```

💻 Output: Names like "Alice", "Andrew", etc.

2. Names Ending with "y"

💡 Example:

```
SELECT name  
FROM employees  
WHERE name REGEXP 'y$';
```

💻 Output: Names like "Mary", "Anthony"

3. Names Containing Digits

💡 Example:

```
SELECT name  
FROM employees  
WHERE name REGEXP '[0-9]';
```

💻 Output: "John123", "Emp007"

4. Names That Do Not Contain Vowels

💡 Example:

```
SELECT name  
FROM employees  
WHERE name REGEXP '^[^aeiouAEIOU]+$';
```

💻 Output: "Rrr", "Bcd"

5. Names with At Least Two Consecutive "a"

💡 Example:

```
SELECT name  
FROM employees  
WHERE name REGEXP 'a{2,}';
```

💻 Output: "Aarav", "Saaab"

- ◆ Oracle Specific REGEXP Functions
- ◆ REGEXP_LIKE(): Checks if string matches a pattern (boolean filter)

 Example:

```
SELECT emp_name
FROM employees
WHERE REGEXP_LIKE(emp_name, '^A');
```

- ◆ REGEXP_INSTR(): Returns position of pattern in string

 Example: `SELECT REGEXP_INSTR('abc123xyz', '[0-9]+') AS position FROM dual;`

 Output: `4` (position of first number)

- ◆ REGEXP_SUBSTR(): Returns substring matching a pattern

 Example: `SELECT REGEXP_SUBSTR('abc123xyz', '[0-9]+') AS number FROM dual;`

 Output: `123`

- ◆ REGEXP_REPLACE(): Replaces substring matching pattern with something else

 Example: `SELECT REGEXP_REPLACE('abc123xyz', '[0-9]', '#') AS masked FROM dual;`

 Output: `abc###xyz`

 Important Notes:

- Use REGEXP carefully, as it's powerful but slower than LIKE.
- Useful for data validation, cleaning, text filtering, and complex searches.
- MySQL and Oracle have better REGEXP support than SQL Server (which is more limited).
- SQL Server uses LIKE, CHARINDEX, or CLR integration for regex-like behavior.

☒ Summary : Regular Expression

Feature	Description
REGEXP_LIKE	Boolean test for regex pattern
REGEXP_INSTR	Position of regex match in string
REGEXP_SUBSTR	Extract substring based on regex
REGEXP_REPLACE	Replace substring using regex
Common use cases	Validation, formatting, pattern matching
SQL support	Oracle ✓ , MySQL ✓ , PostgreSQL ✓ , SQL Server ✗ (limited)

◆ Indexes

◆ What is an Index in SQL?

- An index in SQL is a performance-tuning method used to speed up the retrieval of data from a database table.
- It works like an index in a book — instead of searching each page (row), the index lets you jump directly to the page (row) containing the data.
- Indexes are especially useful for large tables and are automatically used by the SQL engine when you run queries.

◆ Why Use Indexes?

- To speed up SELECT queries and WHERE clause searches.
- To improve JOIN performance.
- Helps with ORDER BY, GROUP BY, and DISTINCT operations.

◆ How Index Works?

- Indexes are created on one or more columns.
- Behind the scenes, most databases use a B-Tree or Hash structure for indexing.
- When you query the table, the index allows faster lookup compared to scanning every row.

❖ Types of Indexes:

1. Single-Column Index
2. Composite Index (Multi-Column Index)
3. Unique Index
4. Clustered Index (✗ Oracle)
5. Non-Clustered Index
6. Full-Text Index (MySQL, SQL Server)
7. Bitmap Index (Oracle, Data Warehouses)
8. Filtered Index (SQL Server, ✗ Oracle)

❖ 1. Single-Column Index

- Indexes built on a single column of a table.

- Best for queries filtering or sorting by that specific column.
- Simple and fast to create.

 Syntax:

```
CREATE INDEX idx_employee_name ON Employees(Name);
```

 Example:

```
SELECT * FROM Employees WHERE Name = 'Vaibhav';
```

◊ 2. Composite Index (Multi-Column Index)

- An index on two or more columns.
- Improves performance for queries using multiple columns in WHERE or JOIN.
- The order of columns in the index matters.

 Syntax:

```
CREATE INDEX idx_emp_dept ON Employees(DepartmentID, Name);
```

 Example:

```
SELECT * FROM Employees WHERE DepartmentID = 3 AND Name = 'Vaibhav';
```

 Note: This index is effective only if DepartmentID is included in WHERE clause (left-to-right rule).

◊ 3. Unique Index

- Ensures that all values in the indexed column(s) are unique.
- Automatically created when you define a PRIMARY KEY or UNIQUE constraint.

 Syntax:

```
CREATE UNIQUE INDEX idx_unique_email ON Employees>Email);
```

 Example:

```
INSERT INTO Employees (Email) VALUES ('vaibhav@example.com'); -- Allowed
INSERT INTO Employees (Email) VALUES ('vaibhav@example.com'); -- Error
```

◊ 4. Clustered Index

- The actual data is stored in the order of the index.
- Each table can have only one clustered index because data can be sorted only one way.
- Automatically created with the PRIMARY KEY in many databases.

 Syntax (SQL Server):

```
CREATE CLUSTERED INDEX idx_emp_id ON Employees(EmployeeID);
```

 Example:

```
SELECT * FROM Employees ORDER BY EmployeeID; -- Fast
```

◊ 5. Non-Clustered Index

- The index contains pointers to the actual table rows.
- Can be multiple non-clustered indexes on a table.
- Faster than scanning the whole table but slower than clustered index for range queries.

 Syntax

```
CREATE NONCLUSTERED INDEX idx_emp_name ON Employees(Name);
```

 Example:

```
SELECT * FROM Employees WHERE Name = 'Vaibhav';
```

◊ 6. Full-Text Index (MySQL, SQL Server)

- Used for searching large blocks of text like blogs, articles, descriptions.
- Enables **MATCH ... AGAINST** or **CONTAINS** style queries.

 Syntax:

```
CREATE FULLTEXT INDEX idx_description ON Products(Description);
```

 Example:

```
SELECT * FROM Products WHERE MATCH(Description) AGAINST('wireless');
```

◊ 7. Bitmap Index (Oracle, Data Warehouses)

- Used for columns with low cardinality (few distinct values, like Gender, Yes/No).
- Very efficient for data warehouses and analytical queries.

 Example: (Oracle):

```
CREATE BITMAP INDEX idx_gender ON Employees(Gender);
```

◊ 8. Filtered Index (SQL Server)

- An index with a WHERE condition, so it only indexes a subset of rows.
- Useful when most values in a column are **NULL** or you query a specific value frequently.

 Syntax:

```
CREATE INDEX idx_active_emps ON Employees(Status) WHERE Status = 'Active';
```

 Example:

```
SELECT * FROM Employees WHERE Status = 'Active'; -- Fast using filtered index
```

 View Existing Indexes:

◆ SQL Server:

```
EXEC sp_helpindex 'Employees';
```

◆ PostgreSQL:

```
SELECT * FROM pg_indexes WHERE tablename = 'employees';
```

◆ MySQL:

```
SHOW INDEXES FROM Employees;
```

◆ 🗑 Dropping Indexes

📘 Syntax:

```
DROP INDEX idx_employee_name; -- SQL Server, Oracle  
DROP INDEX idx_employee_name ON Employees; -- MySQL, PostgreSQL
```

📝 Important Notes:

- Indexes increase read performance but can slow down insert/update/delete operations because indexes also need to be updated.
- Use indexes only when needed — don't over-index.
- For large-scale OLAP systems, bitmap indexes and filtered indexes are more efficient.
- Always check query execution plans to ensure indexes are being used.
- Use Composite Indexes carefully: follow the leftmost prefix rule.

⌚ Summary : SQL Index

Index Type	Description
Single-Column	Index on one column. Improves filtering/sorting.
Composite	Multi-column index. Useful for multi-condition queries.
Unique	Prevents duplicate values. Enforces data integrity.
Clustered	Data is physically stored in order. Only one per table.
Non-Clustered	Pointers to actual rows. Multiple allowed.
Full-Text	Optimized for text searches.
Bitmap	Efficient for few distinct values (analytics).
Filtered	Indexes a subset of rows. Good for NULLs or common values.

◆ SQL Optimization Concepts

◆ 🔎 1. Query Execution Plan

◆ What is a Query Execution Plan?

- A query execution plan (QEP) is like a map that tells you how your SQL query is executed internally by the database.
- It shows which indexes are used, which joins are performed, in what order, and how many rows are processed.
- It helps you understand why a query is slow and how to make it faster.

◆ Why is it Important?

- It helps identify bottlenecks (e.g., full table scans).
- You can fine-tune queries for better performance.
- It shows cost estimates for each operation.

◆ How to View the Execution Plan?

◆ 1. SQL Server:

```
-- Show estimated plan  
SET SHOWPLAN_ALL ON;  
GO  
SELECT * FROM Employees WHERE Name = 'Vaibhav';  
GO  
SET SHOWPLAN_ALL OFF;
```

OR (easier in SSMS):

-- Click on "Display Estimated Execution Plan" (Ctrl + L)

◆ 2. PostgreSQL:

```
EXPLAIN SELECT * FROM Employees WHERE Name = 'Vaibhav';
```

◆ 3. MySQL:

```
EXPLAIN SELECT * FROM Employees WHERE Name = 'Vaibhav';
```

✓ Sample Output (simplified):

id	select_type	table	type	key	rows	Extra
1	SIMPLE	Employees	index	idx_emp_name	10	Using index

- **key:** the index used.
- **type:** shows how the table was accessed (e.g., ALL, index, ref).
- **rows:** number of rows processed (lower is better).

- ◆ 2. Indexing Tips for Optimization
- ◆ How to Use Indexes Wisely
 1. Create indexes on columns used in:
 - o WHERE, JOIN, ORDER BY, GROUP BY.
 2. Use composite indexes when you query using multiple columns.
 3. Avoid indexing:
 - o Columns with too many NULLs or updates.
 - o Columns that are rarely used in filters.

Example : Slow vs Optimized Query

Without index (slow):

```
SELECT * FROM Employees WHERE Email = 'v@example.com';
```

Create index:

```
CREATE INDEX idx_email ON Employees>Email);
```

With index (fast):

```
SELECT * FROM Employees WHERE Email = 'v@example.com';
```

Execution plan will now show `idx_email` is used.

◆ 3. Performance Tuning Basics

- ◆ What is Performance Tuning?
 - It means improving speed and efficiency of your SQL queries.
 - You tune your queries, indexes, or schema to reduce resource usage (CPU, RAM, Disk I/O).
- ◆ Basic Tuning Techniques:

A. SELECT Only Required Columns

Bad:

```
SELECT * FROM Employees;
```

Good:

```
SELECT Name, Email FROM Employees;
```

Reason: `SELECT *` fetches all columns even if not needed.

B. Use WHERE Clause to Filter Early

Bad:

```
SELECT Name FROM Employees;
```

-- Retrieves entire table then filters in app

Good:

```
SELECT Name FROM Employees WHERE Department = 'IT';
```

C. Avoid Functions on Indexed Columns

Index won't work:

```
SELECT * FROM Employees WHERE UPPER(Name) = 'VAIBHAV';
```

Let index help:

```
SELECT * FROM Employees WHERE Name = 'Vaibhav';
```

D. Use JOINS Smartly

- Use appropriate JOINs (INNER, LEFT, etc.)
- Always join on indexed keys

E. Use LIMIT or TOP if You Need Few Rows

 Example:

```
SELECT TOP 10 * FROM Employees ORDER BY Salary DESC;  
-- SQL Server  
SELECT * FROM Employees ORDER BY Salary DESC LIMIT 10;  
-- MySQL/PostgreSQL
```

◆ Execution Plan Types (Common Terms Explained)

Term	Meaning
Table Scan	Reads every row (slowest)
Index Scan	Reads every row in index (better, but not ideal)
Index Seek	Jumps directly to matching row using index (fastest)
Nested Loops Join	Good for small data sets
Hash Join	Efficient for large unsorted data
Merge Join	Best for pre-sorted data sets

 Important Notes:

- Always analyze query plans to tune performance.
- Create indexes based on actual query patterns.
- Test and compare query speed with and without indexes.
- Use tools like SQL Server Profiler, PostgreSQL EXPLAIN ANALYZE, or MySQL Workbench Visual EXPLAIN for deeper insights.

Summary : Key Optimization Concept

Concept	Description
Query Execution Plan	Shows how your query runs internally (like a GPS map)
Indexing	Improves speed for read-heavy queries, use wisely
Selective Columns	Use only needed columns, avoid <code>SELECT *</code>
Function-Free WHERE	Don't use functions on columns in WHERE clause
Proper Joins	Use indexed keys and choose right type of JOIN
Use LIMIT/TOP	Always fetch only what you need

❖ Oracle SQL: Optimization Concepts

◆ 🔎 1. Query Execution Plan in Oracle

◆ What is It in Oracle?

- In Oracle, the execution plan shows the steps Oracle takes to execute a SQL statement, including:
 - Access paths (table scan vs index scan),
 - Join methods,
 - Sort operations, etc.

◆ 📈 How to View It in Oracle?

Option 1: EXPLAIN PLAN

`EXPLAIN PLAN FOR`

`SELECT * FROM Employees WHERE Name = 'Vaibhav';`

`SELECT * FROM TABLE(DBMS_XPLAN.DISPLAY);`

Option 2: AUTOTRACE in SQL*Plus or SQL Developer

`SET AUTOTRACE ON;`

`SELECT * FROM Employees WHERE Name = 'Vaibhav';`

Output includes:

- The execution plan
- Number of consistent gets and physical reads (I/O stats)
- Indexes used (if any)

◆ 🌐 2. Indexing in Oracle SQL

- Oracle supports many types of indexes. The most common are:

◆ B-Tree Index (Default)

- Fast for equality and range queries.
- Used in WHERE, JOIN, ORDER BY.

`CREATE INDEX idx_name ON Employees(Name);`

◆ Bitmap Index

- Ideal for low-cardinality columns (e.g., gender, region).
- Not suitable for frequently updated columns.

`CREATE BITMAP INDEX idx_gender ON Employees(Gender);`

◆ Function-Based Index

- Indexes on expressions or functions.

`CREATE INDEX idx_upper_name ON Employees(UPPER(Name));`

And then query:

`SELECT * FROM Employees WHERE UPPER(Name) = 'VAIBHAV';`

◆ Composite Index

- Multi-column index.

`CREATE INDEX idx_dept_salary ON Employees(Department, Salary);`

📌 Oracle uses the leftmost column(s) for index lookups.

◆ 3. Performance Tuning in Oracle SQL

💡 Key Tips:

1. Use Bind Variables to Avoid Hard Parsing:

```
SELECT * FROM Employees WHERE Name = :name;
```

- Helps Oracle reuse execution plans.

2. Avoid Functions on Indexed Columns (unless using function-based index)

3. Use Hints (Advanced Tuning)

Oracle provides hints to guide the optimizer:

```
SELECT /*+ INDEX(Employees idx_name) */ * FROM Employees WHERE Name = 'Vaibhav';
```

4. Use Statistics and Gather Regularly

Oracle relies on table/index statistics for optimization.

```
EXEC DBMS_STATS.GATHER_TABLE_STATS('HR', 'EMPLOYEES');
```

5. Check Cost, Cardinality, Bytes in Plan Output

```
SELECT * FROM TABLE(DBMS_XPLAN.DISPLAY);
```

Look at:

- Cost: Estimated effort to run the query.
- Cardinality: Estimated number of rows returned.

◆ Oracle-Specific Tools for Tuning

Tool	Purpose
EXPLAIN PLAN	Shows estimated plan
AUTOTRACE	Shows actual statistics
DBMS_XPLAN	Better formatting of plan
SQL Tuning Advisor	Offers tuning suggestions in Oracle Enterprise
AWR / ADDM Reports	Historical performance diagnostics

⌚ Summary :

Concept	Oracle-Specific Feature or Note
Execution Plan	Use EXPLAIN PLAN + DBMS_XPLAN.DISPLAY
Index Types	B-Tree (default), Bitmap (for few distinct values), Function-based
Optimizer Use	Oracle Cost-Based Optimizer (CBO) chooses best plan using stats
Statistics Gathering	Crucial for performance — use DBMS_STATS regularly
Bind Variables	Improve plan reuse and parsing efficiency
SQL Hints	Manually influence execution plan if needed

◆ Understanding Data Export and Import

- Exporting Data: Extracting data from a database and saving it in a file format (like `.sql` or `.csv`) for backup, migration, or analysis. [CodeLucky](#)
- Importing Data: Loading data from external files into a database.

❖ MySQL

◆ Exporting Data

1. Using `mysqldump`:

- Exports the entire database or specific tables.

```
DBACLASS+9Kimis.org - We make games reality+9Coefficient+9
mysqldump -u username -p database_name > backup.sql
```

2. Exporting Specific Tables:

```
mysqldump -u username -p database_name table1 table2 > tables_backup.sql
```

3. Exporting Data to CSV:

```
SELECT * FROM table_name
INTO OUTFILE '/path/to/file.csv'
FIELDS TERMINATED BY ','
ENCLOSED BY ""
LINES TERMINATED BY '\n';
```

◆ Importing Data

1. Importing SQL Dump:

```
mysql -u username -p database_name < backup.sql
```

2. Importing CSV Data:

```
LOAD DATA INFILE '/path/to/file.csv'
INTO TABLE table_name
FIELDS TERMINATED BY ','
ENCLOSED BY ""
LINES TERMINATED BY '\n';
```

❖ Oracle SQL

◆ Exporting Data

1. Using Data Pump (`expdp`):

```
expdp username/password DIRECTORY=dir_name DUMPFILE=export.dmp
LOGFILE=export.log FULL=Y
```

2. Exporting Specific Schema:

```
expdp username/password SCHEMAS= schema_name DIRECTORY=dir_name
DUMPFILE= schema_export.dmp LOGFILE= schema_export.log
```

3. Exporting Table Data to CSV:

```
SPOOL /path/to/file.csv
SELECT * FROM table_name;
SPOOL OFF;
```

◆ Importing Data

1. Using Data Pump (`impdp`):

```
impdp username/password DIRECTORY=dir_name DUMPFILE=export.dmp  
LOGFILE=import.log FULL=Y
```

2. Importing Specific Schema:

```
impdp username/password SCHEMAS=schema_name DIRECTORY=dir_name  
DUMPFILE=schema_export.dmp LOGFILE=schema_import.log
```

3. Importing Data Using SQL*Loader:

Control File (`control.ctl`):

```
LOAD DATA  
INFILE '/path/to/file.csv'  
INTO TABLE table_name  
FIELDS TERMINATED BY ',' OPTIONALLY ENCLOSED BY ""  
TRAILING NULLCOLS  
(column1, column2, column3)
```

Command:

```
sqlldr username/password control=control.ctl log=loader.log
```

❖  PostgreSQL

◆ Exporting Data

1. Using `pg_dump`:

```
pg_dump -U username -W -F c -b -v -f backup.dump database_name
```

2. Exporting Table Data to CSV:

```
COPY table_name TO '/path/to/file.csv' DELIMITER ',' CSV HEADER;
```

◆ Importing Data

1. Restoring from Dump:

```
pg_restore -U username -W -d database_name -v backup.dump
```

2. Importing CSV Data:

```
COPY table_name FROM '/path/to/file.csv' DELIMITER ',' CSV HEADER;
```

❖  SQL Server

◆ Exporting Data

1. Using SQL Server Management Studio (SSMS):

- Navigate to Tasks > Export Data.
- Follow the wizard to export data to desired format (e.g., CSV, Excel).

2. Using `bcp` Utility:

```
bcp database_name.schema_name.table_name out "C:\path\to\file.csv" -c -t, -S  
server_name -U username -P password
```

◆ Importing Data

1. Using SSMS:

- Navigate to Tasks > Import Data.

- Follow the wizard to import data from various sources.
2. Using **bcp** Utility:
- ```
bcp database_name.schema_name.table_name in "C:\path\to\file.csv" -c -t, -S server_name -U username -P password
```
3. Using **BULK INSERT**:

```
BULK INSERT schema_name.table_name
FROM 'C:\path\to\file.csv'
WITH (
 FIELDTERMINATOR = ',',
 ROWTERMINATOR = '\n',
 FIRSTROW = 2
)
```

## ◆ Real-time Use Cases

1. ETL Data Queries
2. Report Generation
3. Data Validation

### ◆ 1. ETL DATA QUERIES

- ETL stands for:
  - Extract → Get data from source tables.
  - Transform → Modify or clean data.
  - Load → Insert into target tables (often for reporting or analytics).
- 🧠 Real-Life Scenario:
  - Let's say you work in a bank.
  - Customer and Transaction data are stored in raw tables.
  - You want to prepare a summary table of customer spending per month.

#### ✓ Step 1: Extract

- Extract data from source.  
-- Common for all SQL engines

```
SELECT * FROM transactions;
```

#### ✓ Step 2: Transform

- Use aggregation, filters, joins, etc.  
-- Monthly total amount spent by customer

```
SELECT
 customer_id,
 DATE_FORMAT(transaction_date, '%Y-%m') AS month,
 SUM(amount) AS total_spent
FROM transactions
GROUP BY customer_id, DATE_FORMAT(transaction_date, '%Y-%m');
```

#### ▀ Differences across SQLs:

| SQL Engine | Syntax Difference                      |
|------------|----------------------------------------|
| MySQL      | DATE_FORMAT(transaction_date, '%Y-%m') |
| PostgreSQL | TO_CHAR(transaction_date, 'YYYY-MM')   |
| SQL Server | FORMAT(transaction_date, 'yyyy-MM')    |
| Oracle SQL | TO_CHAR(transaction_date, 'YYYY-MM')   |

#### ✓ Step 3: Load

- Insert into the summary table.

```
INSERT INTO monthly_summary (customer_id, month, total_spent)
SELECT
 customer_id,
 TO_CHAR(transaction_date, 'YYYY-MM') AS month,
 SUM(amount)
```

```
FROM transactions
GROUP BY customer_id, TO_CHAR(transaction_date, 'YYYY-MM');
```

💡 Tip: In real projects, ETL tools like Informatica, Airflow, or SSIS automate these steps.

- ◆ 2. REPORT GENERATION

- Purpose: To prepare reports for management, clients, or dashboards.

🧠 Real-Life Scenario:

- You want to create a report of top 5 customers by spending for this month.

- MySQL, PostgreSQL

```
SELECT customer_id, SUM(amount) AS total_spent
FROM transactions
WHERE MONTH(transaction_date) = MONTH(CURDATE())
GROUP BY customer_id
ORDER BY total_spent DESC
LIMIT 5;
```

- SQL Server:

```
SELECT TOP 5 customer_id, SUM(amount) AS total_spent
FROM transactions
WHERE MONTH(transaction_date) = MONTH(GETDATE())
GROUP BY customer_id
ORDER BY total_spent DESC;
```

- Oracle SQL:

```
SELECT customer_id, SUM(amount) AS total_spent
FROM transactions
WHERE EXTRACT(MONTH FROM transaction_date) = EXTRACT(MONTH
FROM SYSDATE)
GROUP BY customer_id
ORDER BY total_spent DESC
FETCH FIRST 5 ROWS ONLY;
```

### ◆ 3. DATA VALIDATION

- Purpose: To check data quality, duplicates, nulls, business rule failures, etc.

#### 🧠 Real-Life Scenario:

You want to ensure that:

- No null values exist in `amount`.
- Each transaction has a valid customer.

#### A. Check for NULLs

```
SELECT * FROM transactions WHERE amount IS NULL;
```

#### B. Check Invalid Foreign Keys

```
-- Transactions where customer_id does NOT exist in customer table
SELECT *
FROM transactions t
LEFT JOIN customers c ON t.customer_id = c.customer_id
WHERE c.customer_id IS NULL;
```

#### C. Duplicate Records

```
-- Finding duplicate transaction_ids
SELECT transaction_id, COUNT(*)
FROM transactions
GROUP BY transaction_id
HAVING COUNT(*) > 1;
```

#### 💻 Final Output Example (Assume `transactions` table)

| customer_id | month   | total_spent |
|-------------|---------|-------------|
| 101         | 2025-06 | 12,500.00   |
| 102         | 2025-06 | 11,400.00   |
| 103         | 2025-06 | 10,900.00   |

---

#### 📋 Summary : Syntax Differences

| Action       | MySQL                      | PostgreSQL                | SQL Server             | Oracle SQL                           |
|--------------|----------------------------|---------------------------|------------------------|--------------------------------------|
| Date format  | <code>DATE_FORMAT()</code> | <code>TO_CHAR()</code>    | <code>FORMAT()</code>  | <code>TO_CHAR()</code>               |
| Limit rows   | <code>LIMIT 5</code>       | <code>LIMIT 5</code>      | <code>TOP 5</code>     | <code>FETCH FIRST 5 ROWS ONLY</code> |
| Current date | <code>CURDATE()</code>     | <code>CURRENT_DATE</code> | <code>GETDATE()</code> | <code>SYSDATE</code>                 |

---

## Summary :

| Database   | Export Tool/Command                                          | Import Tool/Command                                                | Notes                             |
|------------|--------------------------------------------------------------|--------------------------------------------------------------------|-----------------------------------|
| MySQL      | <code>mysqldump</code> ,<br><code>SELECT INTO OUTFILE</code> | <code>mysql</code> ,<br><code>LOAD DATA INFILE</code>              | Simple and widely used            |
| Oracle SQL | <code>expdp</code> ,<br><code>SPOOL</code>                   | <code>impdp</code> ,<br><code>SQL*Loader</code>                    | Powerful with advanced options    |
| PostgreSQL | <code>pg_dump</code> ,<br><code>COPY TO</code>               | <code>pg_restore</code> ,<br><code>COPY FROM</code>                | Flexible and efficient            |
| SQL Server | SSMS Export Wizard,<br><code>bcp</code>                      | SSMS Import Wizard,<br><code>bcp</code> , <code>BULK INSERT</code> | Integrated tools with GUI support |

## ◆ OLTP (Online Transaction Processing) vs OLAP (Online Analytical Processing)

| Feature             | OLTP                                                   | OLAP                                                 |
|---------------------|--------------------------------------------------------|------------------------------------------------------|
| 1. Purpose          | Handles day-to-day operations (insert, update, delete) | Helps analyze data for decision making               |
| 2. Users            | End users like cashiers, clerks, customers             | Managers, analysts, business users                   |
| 3. Operations Type  | Read + Write (Insert, Update, Del)                     | Mostly Read (Query, Analyze)                         |
| 4. Speed            | Very fast for small transactions                       | Slower (but optimized for complex queries)           |
| 5. Data Volume      | Handles small amounts of data per transaction          | Works with large volumes of historical data          |
| 6. Examples         | ATM withdrawal, online shopping, ticket booking        | Sales analysis, profit comparison, forecasting       |
| 7. Data Structure   | Highly normalized (many small, related tables)         | Denormalized (fewer, bigger tables like star schema) |
| 8. Query Complexity | Simple, fast queries                                   | Complex queries with aggregations (SUM, AVG, etc.)   |
| 9. Backup Frequency | Regular & frequent (daily or even hourly)              | Less frequent (weekly or monthly)                    |
| 10. Example Systems | Banking, e-commerce, railway reservation systems       | Business Intelligence tools, Data Warehouses         |
| 11. Data Type       | Current, real-time data                                | Historical data                                      |

### 🧠 Quick Example to Understand:

- OLTP = Swiggy/Zomato order system 🍔
  - You place an order → DB updates order table with your payment & food details.
  - Needs to be fast & accurate for thousands of customers at the same time.
- OLAP = Swiggy Business Dashboard 📊
  - Management checks: “*What was the total revenue last month?*”, “*Which city orders the most pizzas?*”.
  - Needs historical + summarized data for decisions.

# PL/SQL

---

## ◆ Basic Introduction to PL/SQL :

### ❖ What is PL/SQL?

- PL/SQL (Procedural Language/Structured Query Language) is Oracle's extension to SQL, used to write code blocks that include programming constructs like variables, loops, conditions, and exception handling.
  - It is Oracle Corporation's procedural extension to SQL.
  - Think of PL/SQL as a way to make SQL "smart" — it lets you add logic and flow control to your database tasks.
- ◆ 2. What Makes PL/SQL Special?
- SQL is declarative: you just ask "what to do?" (e.g., `SELECT`, `INSERT`).
  - PL/SQL is procedural: you also control "how to do it" (e.g., using `IF`, `LOOP`, etc.).
  - PL/SQL = SQL + Programming Logic

### ◆ Why do we use PL/SQL?

| Reason                                               | Explanation                                                                  |
|------------------------------------------------------|------------------------------------------------------------------------------|
| <input checked="" type="checkbox"/> Procedural Logic | You can write IF-ELSE, loops, and conditions.                                |
| <input checked="" type="checkbox"/> Performance      | You send one PL/SQL block, not multiple SQL queries -- faster and efficient. |
| <input checked="" type="checkbox"/> Reusability      | You can create procedures and functions to reuse code.                       |
| <input checked="" type="checkbox"/> Security         | You can hide complex logic inside procedures or packages.                    |
| <input checked="" type="checkbox"/> Error Handling   | You can handle runtime errors using <code>EXCEPTION</code> blocks.           |

## ◆ Components of PL/SQL

| Part                  | Description                                                  |
|-----------------------|--------------------------------------------------------------|
| 1. Block              | Group of statements executed together.                       |
| 2. Variables          | Temporary memory to store values.                            |
| 3. SQL Statements     | You can use <code>SELECT</code> , <code>INSERT</code> , etc. |
| 4. Control Statements | IF-ELSE, LOOP, FOR, WHILE etc.                               |
| 5. Exception Handling | Manages runtime errors gracefully.                           |

#### ◆ 🌐 Real-Life Scenario

- Problem: You want to check if a customer's bill is greater than ₹1000 and give them a discount.
- In SQL, you'd need manual checking for each row.
- In PL/SQL, you can write:

```
DECLARE
 bill_amount NUMBER := 1200;
BEGIN
 IF bill_amount > 1000 THEN
 DBMS_OUTPUT.PUT_LINE('Apply 10% Discount');
 ELSE
 DBMS_OUTPUT.PUT_LINE('No Discount');
 END IF;
END;
```

#### 💻 Output:

Apply 10% Discount

#### 📝 Important Notes: When to use PL/SQL over SQL?

| Use Case                         | SQL | PL/SQL |
|----------------------------------|-----|--------|
| Run one-time query               | ✓   | ✗      |
| Apply business rules             | ✗   | ✓      |
| Execute multiple statements      | ✗   | ✓      |
| Handle logic + error handling    | ✗   | ✓      |
| Reusable logic (functions/procs) | ✗   | ✓      |

## ◆ Difference between SQL vs PL/SQL

| ◆ Feature                | ✓ SQL                                                              | ✓ PL/SQL                                                         |
|--------------------------|--------------------------------------------------------------------|------------------------------------------------------------------|
| 1. Full Form             | Structured Query Language                                          | Procedural Language / Structured Query Language                  |
| 2. Type                  | Declarative language                                               | Procedural programming language                                  |
| 3. Purpose               | Used to query and manipulate data                                  | Used to write full programs with logic and control               |
| 4. Executes              | One statement at a time                                            | A block of code (can have multiple SQL statements)               |
| 5. Can it contain logic? | ✗ No logic like IF, LOOP, etc.                                     | ✓ Yes – has IF, FOR, WHILE, CASE, etc.                           |
| 6. Code Structure        | Simple statements like <code>SELECT</code> , <code>INSERT</code>   | Structured block: <code>DECLARE - BEGIN - EXCEPTION - END</code> |
| 7. Use of Variables      | ✗ Not possible                                                     | ✓ Yes, variables can be declared and used                        |
| 8. Reusability           | ✗ Every time you must write queries again                          | ✓ Write once, reuse as Procedures, Functions, or Packages        |
| 9. Performance           | Less efficient for complex operations                              | More efficient as it reduces multiple database calls             |
| 10. Error Handling       | ✗ Limited (only SQL error codes)                                   | ✓ Full error handling with <code>EXCEPTION</code> blocks         |
| 11. Interaction With DB  | Directly interacts with data                                       | Can call SQL within its block to interact with data              |
| 12. Used For             | CRUD operations – <code>SELECT</code> , <code>UPDATE</code> , etc. | Automating logic – validations, calculations, loops              |
| 13. Example              | <code>SELECT * FROM EMP;</code>                                    | See example below for full PL/SQL block                          |
| 14. Can call each other? | SQL cannot call PL/SQL                                             | PL/SQL can run SQL inside it                                     |
| 15. Best For             | Simple queries and data fetching                                   | Complex programs, business logic, automation                     |

## ◇ What is a PL/SQL Block?

- PL/SQL code is always written inside a block structure. A block is like a program with a beginning, logic, and end.
-  Use `SET SERVEROUTPUT ON;` in SQL\*Plus, Oracle SQL Developer, or LiveSQL to see output.
- Basic Structure:

```
DECLARE
 -- variable declarations (optional)
BEGIN
 -- executable statements
EXCEPTION
 -- error handling (optional)
END;
```

### ❖ 1. Anonymous Block in PL/SQL:

- It is a temporary block of PL/SQL code.
- It does not have a name.
- It cannot be stored in the database for reuse.
- Mostly used for quick testing, debugging, and ad-hoc operations.

#### Syntax

```
DECLARE
 -- Declare variables (optional)
BEGIN
 -- Write logic here
EXCEPTION
 -- Error handling (optional)
END;
```

#### Example

```
DECLARE
 name VARCHAR2(20) := 'Vaibhav';
BEGIN
 DBMS_OUTPUT.PUT_LINE('Hello ' || name || ', Welcome to PL/SQL!');
END;
```

#### Output:

Hello Vaibhav, Welcome to PL/SQL!

#### Important Notes:

1. It is not stored in the Oracle database.
2. It's like a scratchpad for running PL/SQL code quickly.
3. Can include variables, control structures, and SQL queries.
4. Ideal for testing small logic or debugging.
5. Every time you want to run it, you need to retype it.

#### Real-Life Analogy:

- Like writing a reminder on a sticky note — it's useful now, but not saved permanently.

### ❖ 2. Named Block in PL/SQL:

- It is a reusable PL/SQL block.
- It has a name and can be stored permanently in the Oracle database.
- Can be called/executed again and again.

- Mainly used for modular, large applications.
  - ◆ Types of Named Blocks:
    1. Procedures
    2. Functions
    3. Packages
    4. Triggers
-  Syntax of a Named Procedure:
- ```
CREATE OR REPLACE PROCEDURE greet_user IS
BEGIN
    DBMS_OUTPUT.PUT_LINE('Hello from a Named Block!');
END greet_user;
```

- To execute it:

```
BEGIN
    greet_user;
END;
```

-  Output:

Hello from a Named Block!

-  Important Notes:

1. Named blocks are stored in the database.
2. They are reusable, and callable by applications or users.
3. They are modular, which improves code organization and performance.
4. Can have parameters, logic, loops, and exception handling.
5. Best suited for complex logic, batch jobs, or application back-end code.

-  Real-Life Analogy:

- Like saving a Microsoft Word template – you can reuse it again and again without typing it from scratch.

-  Tips for Remembering:

- ◆ Anonymous Block:
 - Think "A" for Ad-hoc
 - Good for temporary tasks, learning, or debugging
- ◆ Named Block:
 - Think "N" for Name + Reusability
 - Use when logic needs to be used again or shared with others

Summary : Anonymous vs Named Block

Feature	Anonymous Block	Named Block
Has a Name?	 No name	 Has a name
Stored in DB?	 No (temporary)	 Yes (permanent)
Reusable?	 Not reusable	 Reusable
Called by Programs?	 No	 Yes (via EXEC or BEGIN block)
Used For	Quick test, learning, debugging	Application logic, reports, stored tasks
Example Use Case	Show welcome message	Calculate tax, send report, update salaries

◆ PL/SQL Variables

◆ What is a Variable in PL/SQL?

- A variable is a named memory location that stores a value during the execution of a PL/SQL block.
- You can use it to store, change, pass, or return data.

◆ Declaring a Variable

Syntax

```
variable_name data_type [:= initial_value];
```

Example

```
DECLARE
    message VARCHAR2(30) := 'Good Morning Vaibhav!';
BEGIN
    DBMS_OUTPUT.PUT_LINE(message);
END;
```

Output

```
Good Morning Vaibhav!
```

❖ Variable Scope in PL/SQL

- Scope refers to the region or block of the PL/SQL program where a variable can be accessed or is visible.
- Variables can be local to a block or global to multiple blocks, depending on where they are declared.

◆ Types of Scope

◆ 1. Local Scope (Inner Block)

- Variables are declared inside a block.
- Only accessible within that block.
- Other blocks cannot use it.

 Syntax:

```
DECLARE
    emp_name VARCHAR2(50) := 'Vaibhav';
BEGIN
    DECLARE
        salary NUMBER := 50000; -- Local to inner block
    BEGIN
        DBMS_OUTPUT.PUT_LINE('Salary: ' || salary);
    END;

    -- DBMS_OUTPUT.PUT_LINE(salary); ✗ Error: salary not visible here
    DBMS_OUTPUT.PUT_LINE('Employee: ' || emp_name); -- ✓ Accessible
END;
```

 Output:

```
Salary: 50000
Employee: Vaibhav
```

❖ Note: `salary` is not accessible outside the inner block.

◆ 2. Global Scope (Outer Block)

- Variables are declared in the outer block.
- It is accessible inside all inner blocks.

 Syntax:

```
DECLARE
    dept_name VARCHAR2(50) := 'Sales'; -- Global
BEGIN
    DECLARE
        location VARCHAR2(30) := 'Mumbai';
    BEGIN
        DBMS_OUTPUT.PUT_LINE('Department: ' || dept_name); -- ✓ Accessible
        DBMS_OUTPUT.PUT_LINE('Location: ' || location); -- ✓ Accessible
    END;
END;
```

 Output:

```
Department: Sales
Location: Mumbai
```

❖ Note: `dept_name` is accessible inside inner block — it has global scope within the entire BEGIN...END.

◆ 3. Shadowing (Variable Overriding)

- The inner block variable can hide a variable from the outer block with the same name.
- This is called variable shadowing.

 Syntax:

```
DECLARE
    name VARCHAR2(20) := 'Outer Vaibhav';
BEGIN
    DECLARE
        name VARCHAR2(20) := 'Inner Vaibhav'; -- Shadows outer 'name'
    BEGIN
```

```

        DBMS_OUTPUT.PUT_LINE('Name: ' || name); -- Inner variable used
      END;
    END;

```

 Output:

Name: Inner Vaibhav

📌 Note: Inner block's variable overrides the outer block's variable temporarily.

 Real-Life Scenario:

- In a billing system, you declare a total_amount variable in the main block (global), and declare a discount in a sub-block only when an offer is available (local).
- This helps manage which variables are used where, and prevents accidental misuse.

 Important Notes:

- Always use unique variable names to avoid confusion.
- Variables declared in inner blocks are not visible outside that block.
- PL/SQL does not allow global variables across programs unless you use packages.

 Summary

Scope Type	Declared In	Accessible In	Shadowing Allowed?
Local	Inside block/procedure	Only within that block	<input checked="" type="checkbox"/> Yes
Global	Outer block	Inner blocks of same block	<input checked="" type="checkbox"/> Yes
Package Scope	Package declaration	Across procedures/functions	<input checked="" type="checkbox"/> Yes

❖ Constants

- A constant is a variable whose value cannot change once it is declared.
- Constants are useful when you want to store fixed values like tax rates, discount percentages, or status codes that should not be altered during program execution.
- Must use the **CONSTANT** keyword.
- Must be initialized at the time of declaration.
- Cannot assign a new value to a constant once defined.
- Improves code readability, maintainability, and avoids accidental changes.

💻 Syntax:

```
DECLARE
    pi CONSTANT NUMBER := 3.14159;
    company_name CONSTANT VARCHAR2(50) := 'OpenAI';
BEGIN
    DBMS_OUTPUT.PUT_LINE('Company: ' || company_name);
    DBMS_OUTPUT.PUT_LINE('Pi value: ' || pi);
END;
```

🖨️ Output:

```
Company: OpenAI
Pi value: 3.14159
```

- Constants can be anchored using **%TYPE**:

```
DECLARE
    emp_id employees.employee_id%TYPE := 101;
    max_id CONSTANT employees.employee_id%TYPE := 999;
```

🧠 Real-Life Scenario:

- You're building a payroll system. The maximum allowed leave per year is fixed to 20. You can declare it as:
MAX_LEAVE CONSTANT NUMBER := 20;

⌚ Summary

Feature	Description
Keyword Used	CONSTANT
Value Change Allowed?	✗ No
Initialization Mandatory?	✓ Yes
Use Case	Fixed values like tax rate, limits, etc.
Anchor Support	✓ %TYPE can be used

◆ PL/SQL Data Types

 PL/SQL Data Types Table:

◆ Main Category	◆ Sub Category	◆ Data Type
1. Scalar	Numeric	NUMBER, BINARY_INTEGER, PLS_INTEGER, SIMPLE_INTEGER, DEC, DECIMAL, FLOAT, INTEGER, INT, NUMERIC, SMALLINT, REAL, DOUBLE PRECISION
	Character	CHAR, VARCHAR2, VARCHAR, NCHAR, NVARCHAR2, LONG
	Date/Time	DATE, TIMESTAMP, TIMESTAMP WITH TIME ZONE, TIMESTAMP WITH LOCAL TIME ZONE, INTERVAL YEAR TO MONTH, INTERVAL DAY TO SECOND
	Boolean	BOOLEAN
2. Composite	Record	RECORD
	Collection	TABLE, VARRAY, NESTED TABLE, ASSOCIATIVE ARRAY (INDEX BY table)
3. LOB	Binary/CLOB	BLOB, CLOB, NCLOB, BFILE
4. Reference	Cursor/Object Reference	REF CURSOR (Weak, Strong), REF Object
5. User-Defined	Object/Collection	OBJECT TYPE, COLLECTION TYPE

❖ 1. SCALAR DATA TYPES

- Scalar types hold a single value — unlike composite types (like RECORDs).
- Classified into 3 major categories:

- ◊ a) Numeric Data Types:
 - Used for storing numbers: integers, decimals, etc.

- ◊ 1. **NUMBER**
 - Most common numeric data type.

 Syntax:

`salary NUMBER(7,2);` -- 7 digits total, 2 after decimal

 Example:

```
DECLARE  
    salary NUMBER(7,2) := 12345.67;
```

```
BEGIN
    DBMS_OUTPUT.PUT_LINE('Salary: ' || salary);
END;
```

☰ Output:
Salary: 12345.67

❖ Note:

- Can store very large or very small numbers.
 - Precision and scale are optional.
- ◊ 2. **BINARY_INTEGER** (Deprecated)
- Used for integer arithmetic.
 - Now replaced by **PLS_INTEGER** (better performance).
- ◊ 3. **PLS_INTEGER**
- Native to PL/SQL, faster than NUMBER for integers.

☰ Syntax:
count_PLIS_INTEGER := 100;

💡 Example:

```
DECLARE
    counter PLS_INTEGER := 5;
BEGIN
    counter := counter * 2;
    DBMS_OUTPUT.PUT_LINE('Double: ' || counter);
END;
```

☰ Output:
Double: 10

❖ Note: No scale/decimal allowed. Good for loops and counters.

◊ 4. **SIMPLE_INTEGER**

- Faster and more memory efficient than **PLS_INTEGER**

- ✓ Steps:
- No NULLs (NULLs are not allowed).
 - Overflow checks are skipped for performance.

💡 Example:

```
DECLARE
    val SIMPLE_INTEGER := 10;
BEGIN
    val := val + 5;
    DBMS_OUTPUT.PUT_LINE(val);
END;
```

☰ Output:
15

❖ Note: Ideal for performance-critical applications.

◊ 5. Synonym Types (Same as NUMBER)

- These are subtypes of NUMBER for readability:
 - DEC, DECIMAL
 - FLOAT
 - INTEGER, INT
 - NUMERIC
 - SMALLINT
 - REAL
 - DOUBLE PRECISION

💡 Example:

```
DECLARE
    age INT := 27;
    marks FLOAT := 88.7;
BEGIN
    DBMS_OUTPUT.PUT_LINE('Age: ' || age || ', Marks: ' || marks);
END;
```

🖨️ Output:

```
Age: 27, Marks: 88.7
```

◊ b) Character Data Types

- Used to store text/strings.

◊ 1. CHAR

- Fixed-length strings.
- Right-padded with spaces.

💡 Example:

```
DECLARE
    code CHAR(5) := 'AB';
BEGIN
    DBMS_OUTPUT.PUT_LINE('Code:' || code || '.');
END;
```

🖨️ Output:

```
Code:AB .
```

📌 Note: Use only if fixed size is required (like gender = 'M').

◊ 2. VARCHAR2 (Recommended)

- Variable-length string.
- Max size up to 32,767 bytes in PL/SQL (4000 in SQL).

💡 Example:

```
DECLARE
    name VARCHAR2(20) := 'Vaibhav';
BEGIN
    DBMS_OUTPUT.PUT_LINE('Name: ' || name);
END;
```

🖨️ Output:

- ◊ 3. **VARCHAR**
 - ANSI-compatible version of **VARCHAR2**.
 - Treated the same in Oracle but should be avoided.
 - ✗ Not recommended for future-proof PL/SQL code.
- ◊ 4. **NCHAR, NVARCHAR**
 - Used for Unicode character storage (multi-language).

💡 Example:

```
DECLARE
    country NCHAR(5) := 'भारत';
BEGIN
    DBMS_OUTPUT.PUT_LINE(country);
END;
```

🖨️ Output:

भारत

📌 Note: Required when building apps in multiple languages.

- ◊ 5. **LONG**
 - Legacy type. Stores long strings (up to 2 GB).
 - ✗ Avoid using — deprecated. Use **CLOB** instead.
- ◊ c) Date/Time Data Types
 - Used to work with dates and time.
- ◊ 1. **DATE**
 - Stores date and time up to seconds.

💡 Example:

```
DECLARE
    today DATE := SYSDATE;
BEGIN
    DBMS_OUTPUT.PUT_LINE('Today: ' || today);
END;
```

🖨️ Output:

Today: 03-JUN-25 10:30:20

📌 Note: Always includes time part.

- ◊ 2. **TIMESTAMP**
 - More precise than **DATE**.
 - Stores fractional seconds.

💡 Example:

```
DECLARE
```

```
ts TIMESTAMP := SYSTIMESTAMP;
BEGIN
    DBMS_OUTPUT.PUT_LINE(ts);
END;
```

Output:

```
03-JUN-25 10.30.20.123456
```

◊ 3. **TIMESTAMP WITH TIME ZONE**

- Includes timezone information.

💡 Example:

```
DECLARE
    tz TIMESTAMP WITH TIME ZONE := SYSTIMESTAMP;
BEGIN
    DBMS_OUTPUT.PUT_LINE(tz);
END;
```

Output:

```
03-JUN-25 10.30.20.123456 +05:30
```

◊ 4. **TIMESTAMP WITH LOCAL TIME ZONE**

- Stores time in the server's local timezone.

◊ 5. **INTERVAL YEAR TO MONTH**

- Stores a span of months/years (e.g., age, subscription)

💡 Example:

```
DECLARE
    i1 INTERVAL YEAR(2) TO MONTH := INTERVAL '1-6' YEAR TO MONTH;
BEGIN
    DBMS_OUTPUT.PUT_LINE('Interval: ' || i1);
END;
```

Output:

```
Interval: +01-06
```

◊ 5. **INTERVAL DAY TO SECOND**

- Stores span in days, hours, minutes, seconds

💡 Example:

```
DECLARE
    i2 INTERVAL DAY TO SECOND := INTERVAL '2 04:05:06' DAY TO
SECOND;
BEGIN
    DBMS_OUTPUT.PUT_LINE('Interval: ' || i2);
END;
```

Output:

```
Interval: +02 04:05:06.000000
```

◊ d) BOOLEAN

- The **BOOLEAN** data type represents logical values — **TRUE**, **FALSE**, and **NULL**.
- It's used mainly for control structures, such as **IF**, **LOOP**, or logical decisions in procedures and functions.
- **BOOLEAN** is only available in PL/SQL – **X** not supported in SQL statements.
- It is not stored in database tables.
- You cannot select or insert a **BOOLEAN** column directly in SQL.
- It's mostly used in procedural code and logic flows.

▀ Syntax:

```
DECLARE
    is_valid BOOLEAN := TRUE;
BEGIN
    IF is_valid THEN
        DBMS_OUTPUT.PUT_LINE('The condition is TRUE');
    ELSE
        DBMS_OUTPUT.PUT_LINE('The condition is FALSE');
    END IF;
END;
```

▀ Output:

The condition is TRUE

💡 Example:

```
DECLARE
    flag1 BOOLEAN := TRUE;
    flag2 BOOLEAN := FALSE;
    flag3 BOOLEAN := NULL;
BEGIN
    IF flag1 THEN
        DBMS_OUTPUT.PUT_LINE('flag1 is TRUE');
    END IF;

    IF flag2 THEN
        DBMS_OUTPUT.PUT_LINE('flag2 is TRUE');
    ELSE
        DBMS_OUTPUT.PUT_LINE('flag2 is FALSE');
    END IF;

    IF flag3 IS NULL THEN
        DBMS_OUTPUT.PUT_LINE('flag3 is NULL');
    END IF;
END;
```

▀ Output:

flag1 is TRUE
flag2 is FALSE
flag3 is NULL

📌 Note:

- Comparing with = **TRUE** or = **FALSE** is not necessary.
- ✓ Just writing **IF my_flag THEN ...** is enough.

🧠 Real-Life Scenario:

- Let's say you're building an HR system:
 - Store employee name in VARCHAR2
 - Store salary in NUMBER(8,2)
 - Store date of joining in DATE
 - Store age difference using INTERVAL YEAR TO MONTH

 Important Notes:

- Always prefer VARCHAR2 over CHAR.
- Use PLS_INTEGER or SIMPLE_INTEGER for performance.
- For date/time: DATE is enough for most cases unless high precision or timezone is required.
- Avoid legacy types like LONG, BINARY_INTEGER.

 Summary

Category	Types	Key Usage
Numeric	NUMBER, PLS_INTEGER, SIMPLE_INTEGER, FLOAT, INT, REAL, etc.	For storing integers and decimals
Character	CHAR, VARCHAR2, VARCHAR, NCHAR, NVARCHAR2, LONG	For storing text
Date/Time	DATE, TIMESTAMP, INTERVALS	Storing time/date values and time spans

❖ 2. Composite Data Types in PL/SQL

- Composite types group multiple values (unlike scalar types which store single values).
 - There are 2 major categories:
 - Records – Like a row with multiple fields
 - Collections – Like arrays (single or multi values)
- ◊ 1. RECORD (like a row structure)
- RECORD stores multiple fields of different data types in one variable (like a table row).
- ◆ Ways to declare RECORD:
- Explicit Definition
 - %ROWTYPE (based on a table)
 - %TYPE (based on columns)
- ◊ a) Explicit Definition:
- When you manually declare a variable by specifying its data type, it is called an explicit definition.
 - You *explicitly write* the data type such as NUMBER, VARCHAR2, DATE, etc., instead of deriving it from a table column or another variable.
 - You control the data type and size.
 - Commonly used when you don't want to tie the variable to a table column.
 - Opposite of anchored declaration (%TYPE, %ROWTYPE).

 Syntax:

```
DECLARE
  TYPE emp_record_type IS RECORD (
    emp_id NUMBER,
```

```

    emp_name VARCHAR2(50),
    salary NUMBER(8,2)
);
emp_rec emp_record_type;
BEGIN
    emp_rec.emp_id := 101;
    emp_rec.emp_name := 'Vaibhav';
    emp_rec.salary := 55000;
    DBMS_OUTPUT.PUT_LINE(emp_rec.emp_name || ' earns ' || emp_rec.salary);
END;

```

 Output:

Vaibhav earns 55000

You've defined:

- A custom record type `emp_record_type` → 
- You explicitly specified the data types (`NUMBER`, `VARCHAR2(50)`, etc.) → 
- Then created a variable `emp_rec` of that type → 
- Assigned values manually → 
- Used `DBMS_OUTPUT.PUT_LINE` to show the output → 

 Real-Life Scenario:

- Suppose you're building a mini calculator or a data entry screen, and you want some working variables like total, name, status, etc.
- These don't need to be tied to any table – just define them explicitly.

 Important Notes:

- Use explicit definition when:
 - You want temporary storage variables.
 - No direct dependency on table columns.
 - Don't confuse it with default initialization (like: `v_count NUMBER := 0;`).
 - Explicit definition makes your code independent from table changes.
 - But you must manually update data types if business rules change (e.g., size of name becomes 100 chars).
- ◊ b) Using `%ROWTYPE`:
- It is used to declare a variable that can store an entire row of a table.
 - `%ROWTYPE` is used to declare a record that can hold an entire row of a table or a cursor with the same column structure and data types.
 - It saves you from declaring each column's data type individually.
 - Automatically adapts to table structure changes.
 - Can be used with tables or cursors.
 - All fields from the table are accessible using dot notation (e.g., `record.column_name`).

 Syntax:

```

DECLARE
    emp_record employees%ROWTYPE;
BEGIN
    SELECT * INTO emp_record FROM employees WHERE emp_id = 101;
END;

```

 Now `emp_record` can access all columns like:

`Emp_record.name`
`Emp_record.salary`

Emp_record.department_id

📝 Important Notes:

- Dynamically matches the structure of a table row.
- Auto-updates if table structure changes.
- If the table structure changes (e.g., column added), `%ROWTYPE` automatically adjusts — no need to update your code manually.
- The record holds all columns, even if you don't use them.
- `%TYPE` is not just for tables — it can refer to variables too.
- If table structure changes, no need to manually update variable types.

🧠 Real-Life Scenario:

- You're creating a report or logic that processes rows from a table like `employees`, `orders`, etc.
- Instead of defining every column separately, just use `%ROWTYPE` — quick, efficient, and clean!

◊ c) Using `%TYPE`

- It is used to declare a variable with the same data type as a column in a table.
- `%TYPE` is used to declare a variable in PL/SQL that automatically takes the data type of an existing table column or another variable.
- Keep your variable's data type in sync with the table column.
- Reduces manual effort to change variable types when the table changes.
- Works with both table columns and previously declared variables.
- Prevents data type mismatch errors in long-term use.

💻 Syntax:

```
DECLARE
    TYPE dept_rec_type IS RECORD (
        dept_id departments.department_id%TYPE,
        dept_name departments.department_name%TYPE
    );
    dept_rec dept_rec_type;
BEGIN
    dept_rec.dept_id := 10;
    dept_rec.dept_name := 'HR';
    DBMS_OUTPUT.PUT_LINE(dept_rec.dept_name);
END;
```

🧠 Real-Life Scenario:

- You're developing a payroll system and referencing `employee.salary` in many places.
- If the column data type changes (e.g., `NUMBER(10,2)` → `NUMBER(12,2)`),
- all `%TYPE` variables will adapt automatically. 

📝 Important Notes:

- Records cannot be directly compared or assigned to each other.
- You can nest records inside other records.

🧠 Real-Life Scenario:

- You want to hold full employee info temporarily during logic processing — use a RECORD instead of many separate variables.

Difference Between %ROWTYPE vs %TYPE

Feature	%TYPE	%ROWTYPE
Used for	Single column/variable	Entire row from table/cursor
Columns it holds	One	All columns in a row
Syntax	var column_name%TYPE	var table_name%ROWTYPE
Flexibility	Moderate	High – adapts to full row structure

◊ 2. COLLECTIONS (Array-like Structures)

- Collections are groups of elements (same data type) stored in one variable.

Three types of collections:

◊ a) INDEX-BY TABLE (Associative Array)

- Oldest, simplest array — indexed by number or string.
- An INDEX-BY TABLE (now known as Associative Array) is a key-value pair collection where:
 - The key can be a number or a string.
 - The value is of any scalar, record, or even object type.
 - Unlike other collections, it has no fixed size and is sparse (can have gaps).
- Indexed using either PLS_INTEGER or VARCHAR2.
- Best for look-up operations or temporary data storage.

Steps to Use:

1. Declare a TYPE using TABLE OF with INDEX BY.
2. Declare a variable of that TYPE.
3. Use index (number or string) to assign values.
4. Access using the same index.

Syntax:

```
DECLARE
  TYPE emp_table IS TABLE OF VARCHAR2(100) INDEX BY PLS_INTEGER;
  emp_names emp_table;
BEGIN
  emp_names(1) := 'Vaibhav';
  emp_names(2) := 'Amit';
  DBMS_OUTPUT.PUT_LINE(emp_names(1)); -- prints Vaibhav
END;
```

Output:

Vaibhav

Note:

- Dynamic size
- Sparse (not all indexes need to be filled)
- No need to initialize

Real-Life Scenario:

- Suppose you are calculating daily sales per store in memory, and want to store:
 'Pune' → ₹50,000, 'Mumbai' → ₹80,000
 No need for a table — just use an associative array with string index.

 Important Notes:

- No need to initialize size.
- Cannot be stored in DB or used in SQL queries.
- Ideal for temporary logic, like loops, lookups, caching values.
- Can be passed to procedures/functions.

◊ b) NESTED TABLE

- A Nested Table is a collection type in PL/SQL that can:
- Store multiple elements like an array (but with no upper bound).
- Be stored in the database as a column.
- Be sparse (can have gaps between elements).
- Be queried and modified using SQL.
- It behaves like a normal array in PL/SQL, but you can store it in a DB column.
- Allows non-contiguous elements (you can delete elements and leave gaps).
- More flexible than VARRAYs.
- Needs to be initialized before usage.

 Syntax:

```
DECLARE
    TYPE num_table IS TABLE OF NUMBER;
    nums num_table := num_table(10, 20, 30);
BEGIN
    FOR i IN 1..nums.COUNT LOOP
        DBMS_OUTPUT.PUT_LINE(nums(i));
    END LOOP;
END;
```

 Output:

```
10
20
30
```

 Note:

- Needs initialization.
- Can use **EXTEND**, **DELETE**, etc.
- Index starts from 1.

 Real-Life Scenario:

- Imagine storing multiple phone numbers or email addresses for a customer in a single row of a table.
- A nested table allows you to do this without creating a separate table and joining.

 Important Notes:

- Nested tables are unbounded.
- Elements can be modified, deleted, or extended.
- They are ideal for storing sets of data like orders, scores, emails, etc.
- Must use **TABLE()** keyword to query from SQL.

◊ c) VARRAY (Variable-Size Array):

- A VARRAY (Variable-Size Array) is a collection type in PL/SQL used to store a fixed number of ordered elements of the same type.
- VARRAYs have a maximum size (upper limit defined).
- They are dense (no gaps allowed between elements).
- Useful for small, ordered, fixed-size data.
- Elements are stored in contiguous memory (no gaps allowed).
- Can be used in both PL/SQL and SQL.
- Ideal for small lists that don't change much.
- Must declare the size during type creation.
- Can be stored as a column in a table.
- Maintains order of elements.
- Deleting individual elements not allowed in VARRAY.

💡 VARRAY Methods:

Method	Description
.COUNT	Number of elements
.FIRST	Index of first element
.LAST	Index of last element
.EXTEND(n)	Adds n empty elements
.DELETE	Not allowed in VARRAY

🧠 Real-Life Scenario:

- You want to store monthly sales data for 6 months per product.
- Since the size is known and fixed, a VARRAY(6) would be a perfect fit.

💻 Syntax:

```
DECLARE
    TYPE salary_array IS VARRAY(5) OF NUMBER;
    sal salary_array := salary_array(20000, 30000, 40000);
BEGIN
    FOR i IN 1..sal.COUNT LOOP
        DBMS_OUTPUT.PUT_LINE(sal(i));
    END LOOP;
END;
```

🖨️ Output:

```
20000
30000
40000
```

◆ Note:

- Size is fixed at declaration.
- Good for when the limit is known (e.g., 12 months).

🧠 Real-Life Scenario:

- RECORD: Temporary storage of a full employee or order row.
- Associative Array: Mapping employee ID → name or salaries.
- Nested Table: Holding a list of dynamically added items like project names.
- VARRAY: Store fixed monthly sales data (12 months).

📝 Important Notes:

- Always initialize Nested Table and VARRAY before use.
- Use `.COUNT`, `.DELETE`, `.EXTEND`, `.EXISTS` for collection operations.
- Associative Arrays are very useful in procedural logic.

█ Difference Between Associative Array vs Nested Table VARRAY

Feature / Type	Associative Array (INDEX-BY TABLE)	Nested Table	VARRAY
Index Type	<code>PLS_INTEGER</code> / <code>VARCHAR2</code> (string or number)	Sequential (1..N)	Sequential (1..N)
Fixed Size?	✗ No (dynamic)	✗ No (dynamic)	✓ Yes (must specify size at creation)
Sparse (Gaps Allowed)?	✓ Yes	✓ Yes	✗ No (always dense)
Can Store in DB?	✗ No	✓ Yes	✓ Yes
Used in SQL Queries?	✗ No	✓ Yes	✓ Yes
Maintains Order?	✗ No (order not guaranteed)	✓ Yes	✓ Yes
Performance	✓ Fast (in-memory)	⚠ Moderate	✓ Fast for small data
Deletion Allowed?	✓ Yes (any index)	✓ Yes (DELETE or TRIM)	✗ No
Initialization Required?	✓ Yes (by assigning values directly)	✓ Yes	✓ Yes
Can be used in PL/SQL?	✓ Yes	✓ Yes	✓ Yes

Summary:

Composite Type	Purpose	Example Use
RECORD	Row-like, multiple columns/fields	Store full employee row
Assoc. Array	Dynamic key-value pair storage	ID to Name mapping
Nested Table	Dynamic list (storable in DB)	Multiple items per customer
VARRAY	Fixed-size ordered list	Monthly salary/sales values

3. LOB (Large Object)

- LOBs are used to store very large data like text, images, files, audio, video, etc.
- ◊ 1. BLOB (Binary Large Object)
 - Used for storing binary data like:
 - Images (JPG, PNG)
 - Audio files (MP3)
 - PDFs, documents

Syntax:

```
CREATE TABLE images (
    id NUMBER,
    photo BLOB
);
```

Note:

- Stored inside the database.
- No character encoding applied.
- Maximum size: up to 4 GB.

Real-Life Scenario:

- Storing scanned documents, resumes, or product images.

◊ 2. CLOB (Character Large Object)

- ◆ Used for storing large text data like:
 - Articles
 - Books
 - Long descriptions

Syntax:

```
CREATE TABLE articles (
    id NUMBER,
    content CLOB
);
```

 Example:

```
DECLARE
    my_clob CLOB;
BEGIN
    SELECT content INTO my_clob FROM articles WHERE id = 1;
    DBMS_OUTPUT.PUT_LINE(DBMS_LOB.SUBSTR(my_clob, 100));
END;
```

 Output:

(First 100 characters of the content)

 Note:

- Supports up to 4 GB of text.
- Can use **DBMS_LOB** package for operations.

◊ 3. NCLOB (National Character LOB)

- Same as **CLOB**, but stores Unicode text.

 Syntax:

```
CREATE TABLE translations (
    id NUMBER,
    content NCLOB
);
```

 Note:

- Used when storing multilingual content (e.g., Hindi, Japanese).
- Supports global applications.

◊ 4. BFILE (Binary File)

- Used for storing external binary files.
- Key Difference: The data is not stored in the database, only a pointer to an OS file.

 Syntax:

```
CREATE TABLE audio_files (
    id NUMBER,
    song BFILE
);
```

 Example:

```
DECLARE
    music BFILE;
BEGIN
    SELECT song INTO music FROM audio_files WHERE id = 1;
    DBMS_OUTPUT.PUT_LINE(DBMS_LOB.GETLENGTH(music));
END;
```

 Note:

- Read-only from PL/SQL.
- File must be in a directory object registered in Oracle.

 Steps to use:

1. Create a DIRECTORY object.

```
CREATE OR REPLACE DIRECTORY my_files AS '/home/oracle/audio';
```

2. Grant access:

```
GRANT READ ON DIRECTORY my_files TO your_user;
```

Difference Between LOB Types

Feature	BLOB	CLOB	NCLOB	BFILE
Data Type	Binary	Text	Unicode	Binary
Stored In DB	<input checked="" type="checkbox"/> Yes	<input checked="" type="checkbox"/> Yes	<input checked="" type="checkbox"/> Yes	<input type="checkbox"/> No (external)
Max Size	~4 GB	~4 GB	~4 GB	~4 GB (pointer)
Read/Write	<input checked="" type="checkbox"/> Yes	<input checked="" type="checkbox"/> Yes	<input checked="" type="checkbox"/> Yes	<input type="checkbox"/> Read-only
Use Case	Images	Articles	Unicode Docs	External Media

Important Notes:

- Use **DBMS_LOB** package to manipulate LOBs (read/write/append/delete).
- Always use temporary LOBs when doing large LOB operations in PL/SQL.
- LOBs are usually not fetched entirely for performance — use chunks.

Real-Life Scenarios:

- Resume Upload Module → use **BLOB**
- Product Description in E-commerce → use **CLOB**
- Multilingual Blog Posts → use **NCLOB**
- Referencing Audio/Video files from server → use **BFILE**

Summary

Type	Used For	Stored In DB?	Notes
BLOB	Images, PDFs, Binary data	<input checked="" type="checkbox"/> Yes	Binary content
CLOB	Long Text, Documents	<input checked="" type="checkbox"/> Yes	Character content
NCLOB	Unicode content	<input checked="" type="checkbox"/> Yes	For multilingual apps
BFILE	External Files	<input type="checkbox"/> No	Pointer to server-side file

❖ 4. Reference Types in PL/SQL

- PL/SQL Reference Types are pointers to complex data structures like cursors and objects.
- Categories of Reference Types:
 1. REF CURSOR – Pointer to a result set (Dynamic SQL)
 2. REF OBJECT – Pointer to an object instance (like a row object)
- ◊ 1. REF CURSOR (Cursor Variables)
 - A REF CURSOR does not store data, it points to a query result set — dynamic and flexible.
 - Types of REF CURSOR
 - ◆ a) Strong REF CURSOR
 - Has fixed return structure (data type)
- ◆ b) Weak REF CURSOR
 - No return type check, more flexible

💻 Syntax:

```
DECLARE
    TYPE emp_cursor_type IS REF CURSOR RETURN employees%ROWTYPE;
    emp_cursor emp_cursor_type;
    emp_rec employees%ROWTYPE;
BEGIN
    OPEN emp_cursor FOR SELECT * FROM employees WHERE department_id =
10;
    LOOP
        FETCH emp_cursor INTO emp_rec;
        EXIT WHEN emp_cursor%NOTFOUND;
        DBMS_OUTPUT.PUT_LINE(emp_rec.first_name || ' earns ' || emp_rec.salary);
    END LOOP;
    CLOSE emp_cursor;
END;
```

💻 Output:

```
John earns 50000
Ravi earns 60000
...
```

📌 Note:

- Return type is fixed to employees%ROWTYPE
- Useful in static query scenarios

◆ b) Weak REF CURSOR

- No return type check, more flexible

💻 Syntax:

```
DECLARE
    TYPE any_cursor IS REF CURSOR;
    emp_cursor any_cursor;
    emp_name employees.first_name%TYPE;
BEGIN
    OPEN emp_cursor FOR SELECT first_name FROM employees;
    LOOP
        FETCH emp_cursor INTO emp_name;
        EXIT WHEN emp_cursor%NOTFOUND;
        DBMS_OUTPUT.PUT_LINE(emp_name);
    END LOOP;
```

```
CLOSE emp_cursor;
END;
```

💻 Output:

```
John
Ravi
...
```

📝 Important Notes:

- No return type specified
- Allows more dynamic SQL
- REF CURSORS are used to return data from procedures/functions to front-end apps (like Java, .NET, Python).
- Also widely used in report generation and pagination logic.

🧠 Real-Life Scenario:

- A procedure needs to return different columns based on the user role.
- You can use a weak REF CURSOR to return various result sets dynamically.

◊ 2. REF Object (Object References)

- These are pointers to object instances in an object-relational model (advanced usage).
- Steps to Use REF OBJECT:

✓ Step 1: Create Object Type

```
CREATE OR REPLACE TYPE emp_obj AS OBJECT (
    emp_id NUMBER,
    emp_name VARCHAR2(50)
);
```

✓ Step 2: Create Object Table

```
CREATE TABLE emp_table OF emp_obj;
```

✓ Step 3: Insert Data

```
INSERT INTO emp_table VALUES (101, 'Vaibhav');
```

✓ Step 4: Use REF to point to object

```
DECLARE
    emp_ref REF emp_obj;
    emp_instance emp_obj;
BEGIN
    SELECT REF(e) INTO emp_ref FROM emp_table e WHERE emp_id = 101;
    FETCH DREF(emp_ref) INTO emp_instance;
    DBMS_OUTPUT.PUT_LINE(emp_instance.emp_name);
END;
```

💻 Output:

```
Vaibhav
```

📌 Note:

- REF is the pointer
- DREF() is used to get the object's values

📝 Important Notes:

- REF CURSORS are much more common than REF OBJECTs.
- Use weak cursors when you don't know the structure in advance.

- Use strong cursors for better compile-time checks.
- REF OBJECT is useful in object-relational models (rare in general business apps).

Difference Between Strong & Weak REF CURSOR

Feature	Strong REF CURSOR	Weak REF CURSOR
Return Type	Fixed (e.g., ROWTYPE)	Not defined
Flexibility	Low	High
Compile-Time Check	 Yes	 No
Use Case	Static queries	Dynamic SQL, joins

Summary:

Reference Type	Description	Used For
REF CURSOR	Pointer to a result set	Dynamic queries, procedure return
Strong CURSOR	Fixed structure (compile-time check)	Fixed table queries
Weak CURSOR	No structure (more flexible)	Dynamic SQL, varying results
REF OBJECT	Pointer to object instance in DB	Advanced, object-relational use

❖ 5. User-Defined Types (UDTs) in PL/SQL

- PL/SQL allows you to create your own data types using:
 1. OBJECT TYPE – Custom structure like a class in OOP
 2. COLLECTION TYPE – Custom lists/arrays like arrays, sets

◊ 1. OBJECT TYPE

- Object Type is a user-defined composite type that groups related variables (attributes) and subprograms (methods), like a class in object-oriented programming.

Syntax:

```
CREATE OR REPLACE TYPE emp_type AS OBJECT (
    emp_id NUMBER,
    emp_name VARCHAR2(50),
    emp_salary NUMBER
);
```

✓ Steps to Use Object Type:

1. Create the object type
2. Create a table OF that object (optional)
3. Declare and use it in PL/SQL

💡 Example: Create and Use Object Type

```
DECLARE
    emp_obj emp_type;
BEGIN
    emp_obj := emp_type(101, 'Vaibhav', 55000);
    DBMS_OUTPUT.PUT_LINE('Name: ' || emp_obj.emp_name || ', Salary: ' ||
    emp_obj.emp_salary);
END;
```

🖨️ Output:

```
Name: Vaibhav, Salary: 55000
```

📌 Note:

- You can define methods (functions/procedures) inside object types as well.
- Useful in OOP-based design in PL/SQL or when interacting with object-relational DB tables.

🧠 Real-Life Scenario:

- You're building an HR system where each employee has a structure (ID, name, dept, etc.). Instead of using multiple variables, define a reusable `emp_type` object.

📝 Important Notes:

- Object Types help modularize and reuse code like OOP.
- Collections are powerful for loops, bulk processing, and in-memory operations.
- Nested Tables and VARRAYs can be stored in DB tables, but Associative Arrays are only for PL/SQL blocks.

📋 Summary:

Type	Use For	Stored in DB?	Dynamic?
OBJECT TYPE	Create reusable data structures	<input checked="" type="checkbox"/> Yes	<input checked="" type="checkbox"/> Yes
VARRAY	Small fixed-size lists	<input checked="" type="checkbox"/> Yes	<input checked="" type="checkbox"/> No
Nested Table	Unbounded list (SQL + PL/SQL)	<input checked="" type="checkbox"/> Yes	<input checked="" type="checkbox"/> Yes
Associative Array	Key-value pairs (PL/SQL only)	<input checked="" type="checkbox"/> No	<input checked="" type="checkbox"/> Yes

◆ Control Structures in PL/SQL

- Control structures in PL/SQL allow us to control the flow of execution in our programs.
- Just like the steering wheel in a car 🚗.
- They help you make decisions, repeat actions, and navigate conditions dynamically.

❖ Three Main Categories of Control Structures

◊ 1. Conditional Control – IF, CASE

- Used to execute certain parts of code based on specific conditions (TRUE/FALSE).
- Helps in decision-making inside the PL/SQL block.
- You can choose between multiple actions depending on variable values or expressions.
- 💡 Examples include: IF, IF-ELSE, ELSIF, and CASE.

◊ 1. IF Statement

- Used to execute code based on one or more conditions.
- Helps you make decisions like grading, salary calculation, or validations.
- You can use IF, ELSIF, and ELSE for multi-branch decisions.

💻 Syntax:

```
IF condition THEN
    -- statements
ELSIF another_condition THEN
    -- statements
ELSE
    -- statements
END IF;
```

💡 Example:

```
DECLARE
    marks NUMBER := 75;
BEGIN
    IF marks >= 90 THEN
        DBMS_OUTPUT.PUT_LINE('Grade: A');
    ELSIF marks >= 75 THEN
        DBMS_OUTPUT.PUT_LINE('Grade: B');
    ELSE
        DBMS_OUTPUT.PUT_LINE('Grade: C');
    END IF;
END;
```

💻 Output:

Grade: B

◊ 2. CASE Statement

- Handling multiple conditions more clearly than nested IFs
- Provides a cleaner alternative to multiple IF...ELSIF conditions.
- Used when checking the value of a single variable against many options.

💻 Syntax:

```
CASE variable
    WHEN value1 THEN statements
    WHEN value2 THEN statements
```

```
    ELSE statements  
END CASE;
```

💡 Example:

```
DECLARE  
    marks NUMBER := 75;  
BEGIN  
    CASE  
        WHEN marks >= 90 THEN  
            DBMS_OUTPUT.PUT_LINE('Grade: A');  
        WHEN marks >= 75 THEN  
            DBMS_OUTPUT.PUT_LINE('Grade: B');  
        ELSE  
            DBMS_OUTPUT.PUT_LINE('Grade: C');  
    END CASE;  
END;
```

💻 Output:

Very Good

- ◊ 2. Iterative Control – **LOOP**, **WHILE**, **FOR**, Nested Loops
 - Used to repeat a set of statements multiple times until a condition is met or a counter ends.
 - Allows you to repeat a block of code multiple times, either until a condition is met or a count ends.
 - Mainly used for loops and repetition tasks like traversing rows, calculating sums, etc.
 - Prevents code duplication and enhances readability during repetitive logic.
 - 💡 Examples include: **LOOP**, **WHILE**, **FOR**, and **Nested Loops**.
- ◊ 1. LOOP Statement (Simple Loop)
 - Repeats a set of statements indefinitely until manually exited.
 - Must include an **EXIT** condition, otherwise becomes an infinite loop.

💻 Syntax:

```
LOOP  
    -- statements  
    EXIT WHEN condition;  
END LOOP;
```

💡 Example:

```
DECLARE  
    i NUMBER := 1;  
BEGIN  
    LOOP  
        DBMS_OUTPUT.PUT_LINE('Value: ' || i);  
        i := i + 1;  
        EXIT WHEN i > 3;  
    END LOOP;  
END;
```

💻 Output:

Value: 1
Value: 2

Value: 3

- ◊ 2. WHILE Loop
 - Repeats statements as long as a condition is TRUE.
 - Condition is checked before entering the loop.

 Syntax:

```
WHILE condition LOOP
    -- statements
END LOOP;
```

 Example:

```
DECLARE
    i NUMBER := 1;
BEGIN
    WHILE i <= 3 LOOP
        DBMS_OUTPUT.PUT_LINE('WHILE Value: ' || i);
        i := i + 1;
    END LOOP;
END;
```

 Output:

```
WHILE Value: 1
WHILE Value: 2
WHILE Value: 3
```

- ◊ 3. FOR Loop
 - Looping a known number of times.
 - Executes code for a fixed number of iterations.
 - Ideal for known ranges.
 - The loop variable is auto-incremented and scope-limited to the loop.

 Syntax:

```
FOR counter IN [REVERSE] start..end LOOP
    -- statements
END LOOP;
```

 Example:

```
BEGIN
    FOR i IN 1..3 LOOP
        DBMS_OUTPUT.PUT_LINE('FOR Loop: ' || i);
    END LOOP;
END;
```

 Output:

```
FOR Loop: 1
FOR Loop: 2
FOR Loop: 3
```

- ◊ 4. Nested Loops
 - Looping inside another loop (common in matrix, table operations)
 - One loop inside another loop.

- Used for matrix-like structures or when comparing multiple sets of data.

 Example:

```
BEGIN
  FOR i IN 1..2 LOOP
    FOR j IN 1..2 LOOP
      DBMS_OUTPUT.PUT_LINE('i=' || i || ',j=' || j);
    END LOOP;
  END LOOP;
END;
```

 Output:

```
i=1, j=1
i=1, j=2
i=2, j=1
i=2, j=2
```

◊ 3. Sequential Control – **GOTO, NULL, EXIT, etc.**

- Controls the linear flow or sequence of execution within a block.
- Does not involve decision-making or loops, but alters how control moves (e.g., skip, exit, jump).
- Used to terminate, jump to labels, or perform no action at all.
-  Examples include: **GOTO, NULL, EXIT, RETURN**.

◊ 1. **GOTO**

- Used to transfer control to a labeled part of the PL/SQL block.
- Helps you jump from one part of the code to another.
- Can make code hard to read, so should be used carefully.
- GOTO should be avoided in structured programming unless absolutely necessary.

 Syntax:

```
<<label_name>>
-- statements
GOTO label_name;
```

 Example:

```
BEGIN
  DBMS_OUTPUT.PUT_LINE('Before GOTO');
  GOTO skip;
  DBMS_OUTPUT.PUT_LINE('This line will be skipped');

  <<skip>>
  DBMS_OUTPUT.PUT_LINE('After GOTO');
END;
```

 Output:

```
Before GOTO
After GOTO
```

 Note: The **DBMS_OUTPUT.PUT_LINE('This line will be skipped')** is not executed due to the jump.

◊ 2. NULL

- A placeholder that does nothing.
- Often used in IF statements where no action is needed for certain cases.

💡 Example:

```
BEGIN
    IF 1 = 2 THEN
        NULL;
    ELSE
        DBMS_OUTPUT.PUT_LINE('Condition is false');
    END IF;
END;
```

💻 Output:

```
Condition is false
```

◊ 3. EXIT

- Used to exit a loop prematurely based on a condition.
- Can be used with or without a label.

💡 Example:

```
DECLARE
    i NUMBER := 1;
BEGIN
    LOOP
        EXIT WHEN i > 3;
        DBMS_OUTPUT.PUT_LINE('EXIT Example: ' || i);
        i := i + 1;
    END LOOP;
END;
```

💻 Output:

```
EXIT Example: 1
EXIT Example: 2
EXIT Example: 3
```

◊ 4. RETURN

- Used to exit from a PL/SQL subprogram (procedure or function).
- Can be placed anywhere inside the block.
- In a function, it's also used to return a value to the caller.

💻 Syntax (for Procedure):

```
RETURN;
```

💻 Syntax (for Function)

```
RETURN value;
```

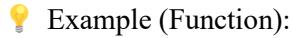
💡 Example (Procedure):

```
CREATE OR REPLACE PROCEDURE greet IS
BEGIN
    DBMS_OUTPUT.PUT_LINE('Hello Vaibhav!');
    RETURN;
    DBMS_OUTPUT.PUT_LINE('This will not be printed');
```

```
END;
```



Output:
Hello Vaibhav!



Example (Function):

```
CREATE OR REPLACE FUNCTION square(n NUMBER)
RETURN NUMBER IS
BEGIN
    RETURN n * n;
END;
```



Calling:

```
BEGIN
    DBMS_OUTPUT.PUT_LINE('Square is: ' || square(5));
END;
```



Output:

Square is: 25



Important Notes:

- In procedures, **RETURN** simply exits the block.
- In functions, **RETURN** also sends back a value.
- Only one **RETURN** is required in a function, but multiple can be used for different conditions.



Example:

```
BEGIN
    NULL; -- placeholder that does nothing
    DBMS_OUTPUT.PUT_LINE('Executed NULL without error');
END;
```



Important Notes:

- Use **EXIT** to break any loop.
- **IF** and **CASE** can't be used interchangeably always. **CASE** is better when checking values of a single variable.
- Avoid infinite **LOOP** without EXIT – it causes runtime hangs.

Summary Table

Control Structure	Used For	Exit Type
IF	Conditional logic	No exit needed
CASE	Multi-branch conditional	No exit needed
LOOP	Infinite loop until EXIT condition	Manual <code>EXIT</code>
WHILE	Loop while condition is true	Auto-check
FOR	Known range of iteration	Auto exit
Nested Loops	Matrix/table traversal	Manual per loop
NULL	Placeholder	—
EXIT	Exit from a loop	Immediate exit

◆ Operators in PL/SQL

- Operators are symbols or keywords used to perform operations on variables, constants, or expressions.
- These operations can be arithmetic, comparison, logical, string manipulation, and more.
- ◊ 1. Arithmetic Operators
 - These are used to perform basic mathematical operations on numeric data types.
 - Arithmetic operators return numeric results and follow mathematical precedence rules.

- ◆ Operators:
 - `+` → Adds two numbers
 - `-` → Subtracts one number from another
 - `*` → Multiplies two numbers
 - `/` → Divides one number by another
 - `**` → Raises a number to the power of another
 - `MOD(a, b)` → Returns the remainder of `a / b`

💻 Syntax:

```
result := a + b;
```

💡 Example:

```
DECLARE
    a NUMBER := 10;
    b NUMBER := 3;
BEGIN
    DBMS_OUTPUT.PUT_LINE('Addition: ' || (a + b));
    DBMS_OUTPUT.PUT_LINE('Power: ' || (a ** b));
    DBMS_OUTPUT.PUT_LINE('Modulus: ' || MOD(a, b));
END;
```

💻 Output:

```
Addition: 13
Power: 1000
Modulus: 1
```

- ◊ 2. Relational (Comparison) Operators

- Used to compare values and return a Boolean result: TRUE, FALSE, or NULL.
- These are often used inside IF statements and loops to make decisions.

- ◆ Operators:
 - `=` → Equal to
 - `!=` or `<>` → Not equal to
 - `>` → Greater than
 - `<` → Less than
 - `>=` → Greater than or equal to
 - `<=` → Less than or equal to

💻 Syntax:

```
IF a < b THEN
    -- do something
END IF;
```

💡 Example:

```
DECLARE
```

```

x NUMBER := 15;
y NUMBER := 20;
BEGIN
  IF x < y THEN
    DBMS_OUTPUT.PUT_LINE('x is less than y');
  END IF;
END;

```

 Output:

x is less than y

◊ 3. Logical Operators

- Used to combine multiple Boolean expressions.
- Mainly used in control structures like IF, CASE, LOOP, etc.
- ◆ Operators:
 - **AND** → TRUE if both conditions are true
 - **OR** → TRUE if any one condition is true
 - **NOT** → Reverses the logical result (TRUE becomes FALSE)

 Syntax:

```

IF condition1 AND condition2 THEN
  -- logic
END IF;

```

 Example:

```

DECLARE
  age NUMBER := 28;
  salary NUMBER := 45000;
BEGIN
  IF age > 25 AND salary > 40000 THEN
    DBMS_OUTPUT.PUT_LINE('Eligible for loan');
  END IF;
END;

```

 Output:

Eligible for loan

◊ 4. String Operator

- PL/SQL provides **||** as a concatenation operator to join strings.
- Useful in reporting, logging, or printing full names, messages, etc.

 Syntax:

```
full_name := first_name || ' ' || last_name;
```

 Example:

```

DECLARE
  fname VARCHAR2(20) := 'Vaibhav';
  lname VARCHAR2(20) := 'Dhakulkar';
BEGIN
  DBMS_OUTPUT.PUT_LINE('Full Name: ' || fname || ' ' || lname);
END;

```

 Output:

Full Name: Vaibhav Dhakulkar

◊ 5. Assignment Operator

- Used to assign a value to a variable.
- This operator does not return a value — it simply sets a value.
- ◆ Operator:
 - `:=` → Assigns the value on the right to the variable on the left

 Syntax:

`variable_name := expression;`

 Example:

```
DECLARE
    marks NUMBER;
BEGIN
    marks := 90;
    DBMS_OUTPUT.PUT_LINE('Marks: ' || marks);
END;
```

 Output: Marks: 90

◊ 6. Membership & Range Operators

- These operators check if a value belongs to a list or a range.
- Used in filtering and decision-making in IF/CASE structures.

◆ Operators:

- `IN` → Checks if a value is in a list of values
- `BETWEEN a AND b` → Checks if a value lies between `a` and `b` (inclusive)

 Syntax:

```
IF grade IN ('A', 'B', 'C') THEN ...
IF salary BETWEEN 20000 AND 50000 THEN ...
```

 Example:

```
DECLARE
    grade CHAR := 'B';
    salary NUMBER := 30000;
BEGIN
    IF grade IN ('A', 'B', 'C') THEN
        DBMS_OUTPUT.PUT_LINE('Valid Grade');
    END IF;

    IF salary BETWEEN 25000 AND 50000 THEN
        DBMS_OUTPUT.PUT_LINE('Middle Range Salary');
    END IF;
END;
```

 Output:

Valid Grade
Middle Range Salary

Real-Life Use Case:

- In a banking app, you could use:
 - **BETWEEN** to filter accounts by balance range.
 - **AND** to check customer eligibility.
 - **||** to prepare names and messages for reports
 - **:=** to assign limits and thresholds.

Summary Table

◆ Operator Type	◆ Used For	💡 Example
Arithmetic	Math operations	$a + b$, <code>MOD(a, b)</code>
Relational	Value comparison	$a = b$, $a < b$
Logical	Logical condition evaluation	$a > 10 \text{ AND } b < 50$
String	Joining string values	<code>First_name Last_name</code>
Assignment	Set values to variables	<code>x := 100</code>
Membership / Range	Value in list or range	<code>IN, BETWEEN</code>

◆ Subprogram in PL/SQL

◆ What is a Subprogram?

- A subprogram is a named PL/SQL block that can be called and reused to perform a specific task.
- It helps to modularize code and avoid repetition.

◆ Types of Subprograms:

1. Procedure – performs an action (does not have to return a value)
2. Function – performs a calculation and returns a value

❖ Procedure In PL/SQL

- A procedure is a named block of PL/SQL code that performs one or more actions, but it does not return a value directly.
- It can accept input (IN), output (OUT), or both (IN OUT) parameters.
- You can call procedures to perform tasks repeatedly (e.g., insert data, generate reports, update records).

Syntax: Create Procedure

```
CREATE OR REPLACE PROCEDURE procedure_name (
    param1 IN datatype,
    param2 OUT datatype
)
IS
    -- local variables
BEGIN
    -- executable statements
END;
```

Example:

```

CREATE OR REPLACE PROCEDURE greet_employee (
    emp_name IN VARCHAR2
)
IS
BEGIN
    DBMS_OUTPUT.PUT_LINE('Hello, ' || emp_name || '!');
END;

```



Call it:

```

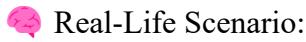
BEGIN
    greet_employee('Vaibhav');
END;

```



Output:

Hello, Vaibhav!



Real-Life Scenario:

- You can create a procedure to:
 - Insert daily attendance
 - Generate monthly report
 - Log error details to an error table

❖ Functions In PL/SQL

- A Function in PL/SQL is a named block of code that performs a specific task and returns a single value
- A function is similar to a procedure, but it must return a single value.
- It can accept input parameters, perform calculations or logic, and must return a value using the **RETURN** keyword.
- It is often used in expressions and SQL SELECT statements, unlike procedures.

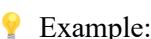


Syntax: Create Function

```

CREATE OR REPLACE FUNCTION function_name (
    param1 IN datatype
)
RETURN return_datatype
IS
    -- local variables
BEGIN
    -- executable statements
    RETURN value;
END;

```



Example:

```

CREATE OR REPLACE FUNCTION get_bonus (
    salary IN NUMBER
)
RETURN NUMBER
IS
    bonus NUMBER;
BEGIN
    bonus := salary * 0.10;
    RETURN bonus;
END;

```

 Call it:

```
DECLARE
    b NUMBER;
BEGIN
    b := get_bonus(50000);
    DBMS_OUTPUT.PUT_LINE('Bonus is: ' || b);
END;
```

 Output:

Bonus is: 5000

 Real-Life Scenario:

- You can create functions to:
 - Calculate tax based on salary
 - Return department name for a department ID
 - Validate input formats (e.g., phone number)

Difference Between Procedure and Function

Feature	Procedure	Function
Return Value?	 No	 Yes (must return a value)
Use in SQL Statement?	 No	 Yes (can be used inside SQL)
Call with BEGIN?	 Yes	 Yes
Use in SELECT?	 Not Allowed	 Allowed
Return with RETURN?	Optional (for exit)	Mandatory to return a value
Example Use	Insert, Update, Delete operations	Calculations, Validations, Lookups

 Important Notes:

- ◆ You can use parameters as:
 - IN — for input
 - OUT — for output
 - IN OUT — for both input and output
- ◆ Use `DBMS_OUTPUT.PUT_LINE()` to display output during debugging
- ◆ Procedures and functions can be stored in the database and called by other programs
- ◆ Modular Code: Both procedures and functions help in creating reusable, organized logic

Summary:

- Procedure → Performs a task but doesn't return value directly
- Function → Performs a task and returns a single value
- Use procedures for actions, like inserting data
- Use functions for calculations, like returning bonus or tax

❖ PL/SQL Parameters: IN, OUT, IN OUT

◊ 1. IN Parameter (Default)

- Used to pass a value *into* a procedure or function.
- Acts like a read-only variable inside the subprogram.
- You cannot change its value inside the procedure — trying to do so gives an error.
- It's the default mode.

💻 Syntax:

```
PROCEDURE greet_user(name IN VARCHAR2)
```

💡 Example:

```
CREATE OR REPLACE PROCEDURE greet_user(name IN VARCHAR2) IS
BEGIN
    DBMS_OUTPUT.PUT_LINE('Hello, ' || name || '!');
END;

BEGIN
    greet_user('Vaibhav');
END;
```

💻 Output:

```
Hello, Vaibhav!
```

📌 Note: `name` is just used — not modified.

◊ 2. OUT Parameter

- Used to return a value *out of* a procedure.
- The calling program sends nothing — the procedure fills the variable.
- Cannot be read *before* assigning a value inside the procedure.

💻 Syntax:

```
PROCEDURE square(num IN NUMBER, result OUT NUMBER)
```

💡 Example:

```
CREATE OR REPLACE PROCEDURE square(num IN NUMBER, result OUT
NUMBER) IS
BEGIN
    result := num * num;
END;

DECLARE
    output NUMBER;
BEGIN
    square(5, output);
    DBMS_OUTPUT.PUT_LINE('Square is: ' || output);
END;
```

💻 Output:

```
Square is: 25
```

📌 Note: `result` value is *sent back* to the calling block.

◊ 3. IN OUT Parameter

- Acts as both input and output.
- Value is passed in, used and modified, then passed back.
- Useful when you want to update the passed value.

💻 Syntax:

```
PROCEDURE update_salary(salary IN OUT NUMBER)
```

💡 Example:

```
CREATE OR REPLACE PROCEDURE update_salary(salary IN OUT NUMBER)
IS
BEGIN
    salary := salary + 1000;
END;

DECLARE
    emp_salary NUMBER := 15000;
BEGIN
    update_salary(emp_salary);
    DBMS_OUTPUT.PUT_LINE('Updated Salary: ' || emp_salary);
END;
```

💻 Output:

```
Updated Salary: 16000
```

🧠 Real-Life Scenario: Pass an employee's current salary, give hike, and return the updated amount.

◊ RETURN in Function

- Unlike procedures, functions return a value directly using RETURN.
- A function always has a return data type defined.
- The returned value can be used in SELECTs, assignments, etc.

💻 Syntax:

```
FUNCTION function_name(param1 IN type) RETURN return_type
```

💡 Example:

```
CREATE OR REPLACE FUNCTION get_area(radius IN NUMBER)
RETURN NUMBER IS
BEGIN
    RETURN 3.14 * radius * radius;
END;

DECLARE
    area NUMBER;
BEGIN
    area := get_area(5);
    DBMS_OUTPUT.PUT_LINE('Circle Area: ' || area);
END;
```

💻 Output:

```
Circle Area: 78.5
```

📌 Note: You cannot use OUT or IN OUT in function parameter list if you're calling it from SQL.

Important Notes:

- Always use **IN** when you just want to pass input data
- Use **OUT** to return results
- Use **IN OUT** when value needs to be updated inside and returned
- Functions must return exactly one value via **RETURN**

Difference Between Procedure vs Function (Summary)

Feature	Procedure	Function
Return value	 Uses OUT parameter	 Uses RETURN keyword
Usage in SQL	 Not allowed directly	 Allowed
Purpose	Perform an action	Perform and return a value
Parameters	IN, OUT, IN OUT	Mostly IN
Syntax	PROCEDURE p_name(...)	FUNCTION f_name(...) RETURN type

◆ What is a Cursor in PL/SQL?

- When an SQL statement is processed, oracle creates a memory area known as context area.
- A cursor is a pointer to this context area. It contains all information needed for processing the statements.
- In PL/SQL the context area is controlled by the cursor. A cursor contains information on a select statement & the rows of data accessed by it.
- A cursor is used to refer to a program to fetch and process the rows returned by the SQL statement one at a time.
- A cursor is a memory pointer in PL/SQL that helps in retrieving and processing query results row by row.
- In PL/SQL, when you run a **SELECT** query that returns multiple rows, you cannot directly assign all rows to variables — instead, you use cursors to process each row individually.
- A cursor acts like a bridge between a SQL query and PL/SQL code to read multiple records one at a time.
- A cursor is like a pointer or handle that helps you fetch and process multiple rows returned by a query one row at a time.

Think of it like:

A cursor is a for-loop for SQL results — it lets you go through each row of a query result step by step.

Why Do We Need a Cursor?

In PL/SQL, a normal **SELECT INTO** can only handle one row.

But if your query returns multiple rows, you need a cursor to handle them one at a time.

❖ Cursor Life Cycle (Explicit & Parameterized)

- ✓ 1. Declare — define the SELECT query
- ✓ 2. Open — start executing the query
- ✓ 3. Fetch — retrieve one row at a time
- ✓ 4. Close — release memory and lock

❖ Types of Cursors

- ◊ 1. Implicit Cursor
 - Created automatically by Oracle when you perform a DML (INSERT, UPDATE, DELETE) or a `SELECT INTO` query.
 - You do not declare it manually — Oracle handles everything for you in the background.
 - Use attributes like `SQL%ROWCOUNT`, `SQL%FOUND`, `SQL%NOTFOUND`, and `SQL%ISOPEN` to track status.
- ◆ Cursor Attributes (for both types)
 - `%FOUND` → Returns TRUE if a row was fetched
 - `%NOTFOUND` → Returns TRUE if no row was fetched
 - `%ROWCOUNT` → Number of rows fetched so far
 - `%ISOPEN` → Returns TRUE if the cursor is still open

💡 Example: `SELECT salary INTO v_sal FROM employees WHERE emp_id = 100;`

📌 Use when: You expect only one row or performing DML operations.

💻 Syntax:

`SELECT column INTO variable FROM table WHERE condition;`

💡 Example:

```
DECLARE
    v_name employees.first_name%TYPE;
BEGIN
    SELECT first_name INTO v_name
    FROM employees
    WHERE employee_id = 100;

    DBMS_OUTPUT.PUT_LINE('Name: ' || v_name);
END;
```

🖨 Output:

`Name: Steven`

📌 Note: Use `SQL%ROWCOUNT`, `SQL%FOUND`, `SQL%NOTFOUND`, and `SQL%ISOPEN` for status.

- ◊ 2. Explicit Cursor
 - You manually declare and manage the cursor to fetch multiple rows from a SELECT query.
 - Gives full control over the row-fetching process.
 - You need to perform 4 steps:
 1. DECLARE
 2. OPEN
 3. FETCH
 4. CLOSE

📌 Use when: Query returns multiple rows, and you want to manually handle data row-by-row.

 Syntax:

```
DECLARE
    CURSOR emp_cur IS
        SELECT first_name, salary FROM employees WHERE department_id = 10;
        v_name employees.first_name%TYPE;
        v_salary employees.salary%TYPE;
BEGIN
    OPEN emp_cur;
LOOP
    FETCH emp_cur INTO v_name, v_salary;
    EXIT WHEN emp_cur%NOTFOUND;
    DBMS_OUTPUT.PUT_LINE(v_name || ' earns ' || v_salary);
END LOOP;
CLOSE emp_cur;
END;
```

 Output:

```
Neena earns 12000
Lex earns 10000
...
```

 Note: Always close the cursor to free resources.

◊ 3. Cursor FOR Loop

- A simplified version of an explicit cursor.
PL/SQL automatically handles OPEN, FETCH, and CLOSE steps.
- Each row is fetched into a record variable inside the loop.
- Makes the code cleaner and less error-prone.
- No need to declare extra variables for fetched values — the loop variable acts as a record.

 Use when: You just need to iterate through records and don't need complex logic between fetch and close.

 Syntax:

```
FOR record_var IN cursor_name LOOP
    -- use record_var.column_name
END LOOP;
```

 Example:

```
DECLARE
    CURSOR emp_cur IS
        SELECT first_name, salary FROM employees WHERE department_id = 20;
BEGIN
    FOR emp IN emp_cur LOOP
        DBMS_OUTPUT.PUT_LINE(emp.first_name || ' earns ' || emp.salary);
    END LOOP;
END;
```

 Output:

```
Michael earns 13000
Pat earns 6000
...
```

◆ Note: `emp` acts as a row-type record variable automatically.

◊ 4. Parameterized Cursor

- A cursor that accepts input parameters (like functions) to make queries dynamic and reusable.
- You can pass different values each time the cursor is opened.
- Helps in filtering data without rewriting the query.

◆ Use when: You want to run the same cursor logic with different WHERE conditions dynamically.

💡 Example: Cursor with department ID or job title as parameter.

💻 Syntax:

```
CURSOR cursor_name(parameter datatype) IS  
    SELECT ... WHERE column = parameter;
```

💡 Example:

```
DECLARE  
    CURSOR emp_cur(dept_id NUMBER) IS  
        SELECT first_name, salary FROM employees WHERE department_id = dept_id;  
BEGIN  
    FOR emp IN emp_cur(30) LOOP  
        DBMS_OUTPUT.PUT_LINE(emp.first_name || ' earns ' || emp.salary);  
    END LOOP;  
END;
```

💻 Output:

```
David earns 7500  
John earns 7200  
...
```

◆ Note: Acts like a function with input arguments.

◊ 5. REF Cursor (Cursor Variables)

- A cursor variable that can point to different result sets at runtime.
- Unlike normal cursors, REF Cursors are dynamic — the structure of the result set can vary.
- Often used to return query results from procedures/functions to other PL/SQL blocks or front-end apps.

◆ Two types:

- Strong REF Cursor → Fixed return structure
- Weak REF Cursor → Flexible (any structure)

◆ Use when:

- Writing generic procedures that return varying result sets.
- Returning rows to host programs like Java, .NET, or Python.

💻 Syntax (Weak REF CURSOR):

```
DECLARE  
    TYPE ref_type IS REF CURSOR;  
    emp_ref ref_type;  
    v_name employees.first_name%TYPE;  
BEGIN
```

```

OPEN emp_ref FOR SELECT first_name FROM employees WHERE
department_id = 40;
LOOP
  FETCH emp_ref INTO v_name;
  EXIT WHEN emp_ref%NOTFOUND;
  DBMS_OUTPUT.PUT_LINE('Employee: ' || v_name);
END LOOP;
CLOSE emp_ref;
END;

```

Output:

Employee: Rose

Employee: Jack

...

Note:

- REF CURSORS can be passed to procedures/functions.
- Ideal for dynamic applications.

Difference Between Cursors

Cursor Type	Need to Declare	Suitable for	 Example Use
Implicit		Single row	SELECT INTO
Explicit		Multiple rows	Fetch employees manually
Cursor FOR Loop		(simplified)	Multiple rows
Parameterized			Dynamic WHERE
REF Cursor		(advanced)	Dynamic SQL
			Procedure/function passing

Real-Life Scenarios:

Cursor Type	Used In Scenario
Implicit Cursor	When updating employee salaries in a single command
Explicit Cursor	Reading employee records and applying custom logic row-by-row
Cursor FOR Loop	Generating monthly salary slips automatically
Parameterized Cursor	Creating reports for different departments by passing department_id dynamically
REF Cursor	Passing dynamic result sets from PL/SQL to Java front-end for data visualization

Important Notes:

- Always handle `%NOTFOUND` and `%FOUND` to avoid infinite loops.
- Cursors consume memory. Always `CLOSE` explicitly declared ones.
- Prefer `Cursor FOR Loop` for cleaner, readable code.
- Use REF cursors in advanced procedures/packages.

Summary: Cursor

- Cursors help fetch multiple rows from a query one at a time
 - You use:
 - Implicit Cursors → For automatic handling of single-row or DML queries
 - Explicit Cursors → When you need full control over row-by-row logic
 - FOR Loop Cursors → When you want cleaner code and don't need manual handling
 - Parameterized Cursors → When you need dynamic query filtering
 - REF Cursors → When working with dynamic SQL and interacting with applications.
-

What is Exception Handling in PL/SQL?

- An exception is an error or unexpected situation that occurs during the execution of a PL/SQL block.
- Oracle provides mechanisms to catch and handle these exceptions gracefully without crashing the program.
- Exception handling is done inside the **EXCEPTION** section of a PL/SQL block.

Syntax of Exception Handling:

```
DECLARE
    -- variable declarations
BEGIN
    -- executable statements
EXCEPTION
    WHEN exception_name THEN
        -- handle the exception
END;
```

Types of Exceptions in PL/SQL

1. Predefined Exceptions

- These are internally defined by Oracle.
- You can directly use them in your **WHEN** clause — no need to declare them.
- They're automatically triggered when a standard runtime error occurs.

Common Predefined Exceptions:

Exception Name	Triggered When...
NO_DATA_FOUND	A SELECT INTO returns no rows.
TOO_MANY_ROWS	A SELECT INTO returns more than one row.
ZERO_DIVIDE	Division by zero is attempted.
INVALID_NUMBER	Conversion from string to number fails.
VALUE_ERROR	Assignment fails due to incompatible values or sizes.

 Example: NO_DATA_FOUND

```

DECLARE
    v_name VARCHAR2(50);
BEGIN
    SELECT name INTO v_name FROM employees WHERE emp_id = 9999;
    DBMS_OUTPUT.PUT_LINE('Name: ' || v_name);
EXCEPTION
    WHEN NO_DATA_FOUND THEN
        DBMS_OUTPUT.PUT_LINE('Employee not found.');
END;

```

 Output:

Employee not found.

◊ 2. User-Defined Exceptions

- When Oracle doesn't provide a specific exception for your logic, you can create your own named exceptions.
- You can define your own custom exceptions using the EXCEPTION keyword.
- Use RAISE to trigger them manually.

 Syntax:

```

DECLARE
    e_invalid_salary EXCEPTION; -- define
    salary NUMBER := -5000;
BEGIN
    IF salary < 0 THEN
        RAISE e_invalid_salary; -- raise
    END IF;
EXCEPTION
    WHEN e_invalid_salary THEN
        DBMS_OUTPUT.PUT_LINE('Invalid salary!');
END;

```

 Output:

Invalid salary!

 Real-Life Scenario:

- If a user tries to enter a negative quantity or price, raise a custom error before inserting into the DB.

◊ 3. Using RAISE Statement

- RAISE is used to manually throw an exception.
- You can raise either:
 - A predefined exception
 - A user-defined exception
- Use inside IF conditions or logic checks to exit early.
- Often used after a validation check or business logic failure.

 Example:

```

DECLARE
    e_negative_salary EXCEPTION;
    salary NUMBER := -2000;
BEGIN
    IF salary < 0 THEN

```

```

        RAISE e_negative_salary;
    END IF;
EXCEPTION
    WHEN e_negative_salary THEN
        DBMS_OUTPUT.PUT_LINE('Salary cannot be negative!');
    END;

```

💡 Note: Always use **RAISE** when your business rule fails — e.g., "Employee age should be > 18".

◊ 4. Using **RAISE_APPLICATION_ERROR**

- This is a built-in procedure of the **DBMS_STANDARD** package.
- This is used to raise custom, application-specific errors with error codes and messages.
- You can return an Oracle error between **-20000** and **-20999**.
- Mostly used in:
 - Stored Procedures
 - Business Logic validation
 - Logging and auditing

💡 Example:

```

BEGIN
    RAISE_APPLICATION_ERROR(-20001, 'Something went wrong in your logic!');
END;

```

💻 Output:

ORA-20001: Something went wrong in your logic!

◆ IMP POINT: **RAISE_APPLICATION_ERROR** can be used in stored procedures to inform calling programs (Java, Python, etc.) about the failure.

🧠 Real-Life Use:

- When you want the application calling your PL/SQL procedure to know why it failed (like invalid stock code, unauthorized access, etc.)

◊ 5. Extra: **WHEN OTHERS**

- It is a catch-all exception — it catches any unhandled exception.
- Use it wisely — don't silently suppress real errors.

💻 Syntax: Using **WHEN OTHERS**

```

EXCEPTION
    WHEN ZERO_DIVIDE THEN
        DBMS_OUTPUT.PUT_LINE('Divide by zero error!');
    WHEN OTHERS THEN
        DBMS_OUTPUT.PUT_LINE('Some unknown error occurred.');

```

💡 Example:

```

BEGIN
    NULL; -- simulate unknown error
EXCEPTION
    WHEN OTHERS THEN
        DBMS_OUTPUT.PUT_LINE('An unknown error occurred.');
END;

```

🧠 Real-Life Scenarios:

Scenario	Use
Employee not found in database	Use <code>NO_DATA_FOUND</code>
Invalid bonus calculation (like < 0)	Raise a user-defined exception
Division or average salary error	Catch <code>ZERO_DIVIDE</code>
Complex business rule violation	Use <code>RAISE_APPLICATION_ERROR</code>

Difference Table:

Feature	Predefined Exception	User-defined Exception	RAISE	RAISE_APPLICATION_ERROR
Defined by Oracle?	<input checked="" type="checkbox"/> Yes	<input type="checkbox"/> No (you define)	<input type="checkbox"/>	<input type="checkbox"/>
Need to Declare?	<input type="checkbox"/> No	<input checked="" type="checkbox"/> Yes	<input checked="" type="checkbox"/>	<input type="checkbox"/>
Raises Exception?	<input checked="" type="checkbox"/> Automatically	<input checked="" type="checkbox"/> Manually (with RAISE)	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>
Custom Message Support	<input type="checkbox"/> No	<input type="checkbox"/> (but with logic)	<input type="checkbox"/>	<input checked="" type="checkbox"/> Yes
Used in Business Logic	<input type="checkbox"/> Sometimes	<input checked="" type="checkbox"/> Frequently	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>

Important Notes:

- Always handle exceptions to prevent program crashes.
- You can have multiple `WHEN` clauses for different exceptions.
- Use `OTHERS` clause to catch any unhandled exception.
- Always write exception handlers to avoid Oracle throwing unhandled errors
- Don't suppress with `WHEN OTHERS` — always log or handle
- Use logging tables to store exceptions during failures
- Use `RAISE_APPLICATION_ERROR` in production systems for clear communication

Summary:

- Exception handling in PL/SQL helps in writing error-resilient code
- Types:
 - Predefined → Oracle built-in
 - User-defined → Defined and raised manually
 - RAISE → To trigger exceptions
 - `RAISE_APPLICATION_ERROR` → Custom message with error number
- Use exception blocks to log errors, continue the process, or notify users properly.

◆ PL/SQL Packages

- A package is a schema object that groups related procedures, functions, variables, constants, cursors, and exceptions into a single unit.
- A PL/SQL Package is a collection of related subprograms and other elements grouped together.
- It has two parts:
 - Specification (Spec) – declares what's available publicly.
 - Body – contains the actual implementation.
- Think of it like a class in Java: the spec is like a class interface and the body is the implementation.

◆ Package Specification (Header)

- Declares procedures, functions, cursors, variables, constants that are public (visible outside).
- No logic is written here — only declarations.

💻 Syntax:

```
CREATE OR REPLACE PACKAGE emp_pkg IS
    PROCEDURE display_emp(emp_id IN NUMBER);
    FUNCTION get_salary(emp_id IN NUMBER) RETURN NUMBER;
END emp_pkg;
```

◆ Package Body

- Contains the implementation of what is declared in the spec.
- You can also define private procedures/functions here that are not accessible outside.

💻 Syntax:

```
CREATE OR REPLACE PACKAGE BODY emp_pkg IS

    -- Public Procedure
    PROCEDURE display_emp(emp_id IN NUMBER) IS
        v_name VARCHAR2(50);
    BEGIN
        SELECT ename INTO v_name FROM emp WHERE empno = emp_id;
        DBMS_OUTPUT.PUT_LINE('Employee Name: ' || v_name);
    END;

    -- Public Function
    FUNCTION get_salary(emp_id IN NUMBER) RETURN NUMBER IS
        v_sal NUMBER;
    BEGIN
        SELECT sal INTO v_sal FROM emp WHERE empno = emp_id;
        RETURN v_sal;
    END;

    -- Private Procedure (Not in spec)
    PROCEDURE internal_log IS
    BEGIN
        DBMS_OUTPUT.PUT_LINE('Log: Internal Use Only');
    END;

END emp_pkg;
```

 Output Example (Calling Package Elements):

```
BEGIN  
    emp_pkg.display_emp(7369);  
    DBMS_OUTPUT.PUT_LINE('Salary: ' || emp_pkg.get_salary(7369));  
END;
```

 Output:

```
Employee Name: SMITH  
Salary: 800
```

 Public vs Private Procedures in Package

Feature	Public Procedure	Private Procedure
Declared In	Package Specification	Package Body Only
Accessible Outside?	 Yes	 No
Use Case	Used by external code (SQL/Apps)	Helper subroutines inside package
Example Name	display_emp, get_salary	internal_log

 Note: If a procedure is not declared in the spec, it's private.

 Real-Life Scenario:

Let's say a company has:

- `display_emp` (Public): For showing employee names.
- `get_salary` (Public): For returning salaries.
- `internal_log` (Private): For logging internally, no one outside needs it.

They wrap it all into a package so that:

- Everything related to employee logic stays together.
- Security is maintained (no one can call `internal_log` directly).

 Important Notes:

- You can initialize global variables in the package body.
- Packages improve performance by loading all logic into memory on first call.
- You can overload procedures/functions in packages (same name, different parameters).
- Can also define package-level constants, cursors, and exceptions.

 Summary:

- A PL/SQL Package = Spec (what) + Body (how)
- Spec declares public stuff
- Body writes the logic and can include private subprograms
- Used for modularity, reusability, performance, and security

◆ What is a Trigger in PL/SQL?

- A trigger is a stored PL/SQL block that automatically executes or fires when a specific event (INSERT, UPDATE, DELETE) occurs on a table or view.
- Triggers are invoked by the oracle engine automatically whenever a specified event occurs.
- Triggers are stored into the database and invoked repeatedly when specific conditions match.
- It helps automate tasks.
- Cannot be manually called like procedures.
- Associated with a table or a view.
- Can be defined to fire before or after the DML operation.
- In a trigger, `:NEW` refers to the new values of the row (inserted or updated).
- Similarly, `:OLD` refers to the old values of the row (before update /delete).

◆ Trigger Structure (Syntax)

```
CREATE [OR REPLACE] TRIGGER trigger_name
{BEFORE | AFTER | INSTEAD OF}
{INSERT | UPDATE | DELETE}
ON table_name
[FOR EACH ROW]
[WHEN (condition)]
DECLARE
    -- Variable declarations
BEGIN
    -- Trigger logic
END;
/
```

❖ Types of Triggers:

1. BEFORE Trigger
 2. AFTER Trigger
 3. INSTEAD OF Trigger (Only for Views)
 4. Row-Level Trigger
 5. Statement-Level Trigger
 6. Compound Triggers (Oracle 11g+)
- ❖ 1. BEFORE Trigger
- Fires before the DML operation.
 - Mainly used for validation or modification of data before committing to the table.

💡 Example: Validate salary before insert

```
CREATE OR REPLACE TRIGGER trg_before_emp
BEFORE INSERT ON employees
FOR EACH ROW
BEGIN
    IF :NEW.salary < 10000 THEN
        RAISE_APPLICATION_ERROR(-20001, 'Salary must be at least 10,000');
    END IF;
END;
/
```

💻 Output

If someone tries to insert an employee with salary `5000`, Oracle will throw an error:

`ORA-20001: Salary must be at least 10,000`

Real-Life Scenario:

- ◆ Used in e-commerce platforms for data validation:
 - When a customer places an order, a **BEFORE INSERT** trigger checks if the product stock is available.
 - If stock is zero, the trigger prevents the insert into the orders table.
- ◊ 2. **AFTER Trigger**
 - Fires after the DML operation.
 - Used for logging, auditing, notifications, etc.

 Example: Log employee inserts

```
CREATE OR REPLACE TRIGGER trg_after_emp
AFTER INSERT ON employees
FOR EACH ROW
BEGIN
    INSERT INTO emp_audit(emp_id, action_time)
        VALUES (:NEW.emp_id, SYSDATE);
END;
/
```

Output

Every new insert in `employees` logs the action in `emp_audit`.

Real-Life Scenario:

- ◆ Used in banking systems to log transactions:
- After a money transfer (UPDATE on balance), an **AFTER UPDATE** trigger logs the old and new balances into a transaction log table for audit trail and regulatory compliance.

◊ 3. **INSTEAD OF Trigger** (Only for Views)

- Used when you want to perform actions on a view (which normally doesn't support DML).
- Define how the view should handle DML.

 Example: Insert into view `emp_view` that joins multiple tables

```
CREATE OR REPLACE TRIGGER trg_instead_view
INSTEAD OF INSERT ON emp_view
FOR EACH ROW
BEGIN
    INSERT INTO employees(emp_id, name) VALUES (:NEW.emp_id, :NEW.name);
    INSERT INTO departments(dept_id, dept_name) VALUES (:NEW.dept_id,
:NEW.dept_name);
END;
/
```

Output

Allows inserting into a complex view by inserting into base tables.

Real-Life Scenario:

- ◆ Used in HR reporting dashboards:
- The front-end reports are built using views that combine employee and department data. An **INSTEAD OF** trigger is used to allow inserts into these views so that both employee and department records are updated in their respective base tables.

◊ 4. Row-Level Trigger

- Defined using **FOR EACH ROW**.
- Executes once for each row affected.
- Access old/new values using **:OLD**, **:NEW**.

💡 Example: Audit salary changes row-by-row

```
CREATE OR REPLACE TRIGGER trg_salary_audit
AFTER UPDATE OF salary ON employees
FOR EACH ROW
BEGIN
    INSERT INTO salary_audit(emp_id, old_sal, new_sal, mod_date)
    VALUES (:OLD.emp_id, :OLD.salary, :NEW.salary, SYSDATE);
END;
/
```

💻 Output:

Each salary update will log the old and new values.

🧠 Real-Life Scenario:

- ◆ Used in inventory systems:
- Every time a product's quantity is updated in the inventory table, a **ROW-LEVEL AFTER UPDATE** trigger records the quantity change in an inventory audit table with user details and timestamp.

◊ 5. Statement-Level Trigger

- Default type (if **FOR EACH ROW** is not specified).
- Executes once per statement, not per row.

💡 Example: Log the number of rows deleted

```
CREATE OR REPLACE TRIGGER trg_del_log
AFTER DELETE ON employees
BEGIN
    INSERT INTO log_table(msg, log_time)
    VALUES ('Employees deleted', SYSDATE);
END;
/
```

💻 Output

Deletes 10 rows → trigger runs once, logs one message.

📝 Important Notes:

- Row-level → Runs row by row → uses **:NEW** and **:OLD**.
- Statement-level → Runs once per SQL statement → cannot use row values.

🧠 Real-Life Scenario:

◆ Used in access monitoring systems:

- When someone deletes a record from a secured table (e.g., user data), a **STATEMENT-LEVEL AFTER DELETE** trigger sends an alert email to the admin team without caring about which rows were deleted.

◊ 6. Compound Triggers (Oracle 11g+)

- Useful when you want multiple timing points (before, after) in one trigger.
- Avoids multiple trigger creation and avoids mutating table error.

💡 Example: Track rows inserted & log count

```
CREATE OR REPLACE TRIGGER trg_compound_example
FOR INSERT ON employees
COMPOUND TRIGGER

    -- Declare section
    row_count NUMBER := 0;

    BEFORE STATEMENT IS
        BEGIN
            DBMS_OUTPUT.PUT_LINE('Starting insert...');

        END BEFORE STATEMENT;

    AFTER EACH ROW IS
        BEGIN
            row_count := row_count + 1;
        END AFTER EACH ROW;

    AFTER STATEMENT IS
        BEGIN
            INSERT INTO insert_log(msg, row_count, log_time)
            VALUES ('Inserted into employees', row_count, SYSDATE);

        END AFTER STATEMENT;

    END trg_compound_example;
/
```

💻 Output:

If 5 employees are inserted, it logs `row_count = 5` after the entire statement.

🧠 Real-Life Scenario:

- ♦ Used in bulk data operations (e.g., data migration tools):
 - When migrating thousands of customer records, a **COMPOUND TRIGGER** is used to track total inserted rows and log the result only once at the end of the operation — improving performance and avoiding mutating errors.

📝 Important Notes on Triggers

- Triggers run automatically — you cannot call them manually.
- Triggers should not contain complex business logic; use procedures/functions for that.
- Avoid DML on the same table inside a row-level trigger (mutating table error).
- Use compound triggers to handle bulk processing efficiently.
- Always test triggers thoroughly as they silently run in the background.
- Avoid having too many triggers on one table — it may slow down DML operations.
- Don't create triggers just for fun — always have a valid use case.
- :OLD and :NEW are available only in row-level triggers.
- **INSTEAD OF** triggers are the only way to perform DML on complex views.
- Make sure to document triggers — they're often hard to trace in large systems.

Summary

Type	Timing	Scope	Use Case
BEFORE	Before DML	Row/Stmt	Validation, Data Transformation
AFTER	After DML	Row/Stmt	Auditing, Logging
INSTEAD OF	Replaces DML	Row Only	DML on Views
Row-Level	Each Row	Needs FOR EACH ROW	Fine-grained Control
Statement-Level	Entire DML	Default	Summary, Count, Control
Compound	Multi-timing	Row & Stmt	Combine logic, avoid mutating error

PL/SQL COLLECTIONS & RECORDS

- Collections are PL/SQL structures used to store multiple values (like arrays in other languages).
- They help in bulk processing, looping through values, and temporary storage during PL/SQL operations.

Collection Methods

- All PL/SQL collections support some common built-in methods that help you handle elements:

Method	Purpose	Supported By
COUNT	Returns total number of elements	All
EXTEND	Adds empty elements	Nested Table, VARRAY
DELETE	Removes elements	Nested Table, Assoc. Array
EXISTS(n)	Checks if element at index n exists	Assoc. Array
FIRST/LAST	Gives lowest/highest index value	All
PRIOR/NEXT	Used for iteration	All
TRIM	Removes elements from end	Nested Table, VARRAY

 Example:

```
names.DELETE(2); -- deletes 2nd element  
IF names.EXISTS(2) THEN ...
```

 Important Notes:

- Not all methods are available for all collection types.
- Always check `.EXISTS()` before accessing index in associative arrays.

 Collection Types

 1. RECORD

- A PL/SQL Record is a group of related data items (like a row in a table) of different data types.
- You can use it to store a single row of data from a table or cursor.
- Record = Row; Collection = Column Group

 Syntax:

```
DECLARE  
    TYPE emp_rec_type IS RECORD (  
        emp_id NUMBER,  
        emp_name VARCHAR2(50),  
        salary NUMBER  
    );  
    emp_rec emp_rec_type;  
BEGIN  
    emp_rec.emp_id := 101;  
    emp_rec.emp_name := 'Vaibhav';  
    emp_rec.salary := 60000;  
  
    DBMS_OUTPUT.PUT_LINE('ID: ' || emp_rec.emp_id || ', Name: ' ||  
        emp_rec.emp_name || ', Salary: ' || emp_rec.salary);  
END;  
/
```

 Example: Storing a single employee's details

 Output:

ID: 101, Name: Vaibhav, Salary: 60000

 Real-Life Scenario:

- Used in procedures/functions to hold one row of data fetched from the `employees` table for further processing.

 Important Notes:

- Record fields can be of different data types.
- Can be based on `table%ROWTYPE` to automatically match table structure.

 2. NESTED TABLE

- A nested table is an unordered set of homogeneous elements (same type), just like a regular table inside PL/SQL.
- Can grow dynamically and be stored in database columns.

 Syntax:

```
DECLARE
```

```

TYPE name_table IS TABLE OF VARCHAR2(100);
names name_table := name_table('Vaibhav', 'Ajay', 'Ravi');
BEGIN
    names.EXTEND;
    names(4) := 'Karan';

    FOR i IN 1..names.COUNT LOOP
        DBMS_OUTPUT.PUT_LINE('Name: ' || names(i));
    END LOOP;
END;
/

```

Example: Storing and printing multiple names

Output:

```

Name: Vaibhav
Name: Ajay
Name: Ravi
Name: Karan

```

Real-Life Scenario:

- Used in student registration systems to store multiple subjects selected by each student in a single row.

Important Notes:

- Supports methods like **EXTEND**, **DELETE**, **COUNT**.
- Needs constructor syntax to initialize (`:= type_name(...)`).
- Can be used in SQL if stored in database columns.

◊ 3. VARRAY (Variable-size Array)

- A VARRAY is a bounded collection — meaning it has a fixed maximum size.
- Best when the maximum number of elements is known in advance.

Syntax:

```

DECLARE
    TYPE score_array IS VARRAY(5) OF NUMBER;
    scores score_array := score_array(78, 85, 92);
BEGIN
    FOR i IN 1..scores.COUNT LOOP
        DBMS_OUTPUT.PUT_LINE('Score: ' || scores(i));
    END LOOP;
END;
/

```

Example: Storing exam scores (maximum 5)

Output:

```

Score: 78
Score: 85
Score: 92

```

Real-Life Scenario:

- Used in quiz apps or test systems to store maximum 5 test scores of each student.

Important Notes:

- Size must be declared while defining the type.
- Can be stored in tables.

- Does not support the DELETE method.
- ◊ 4. ASSOCIATIVE ARRAY (INDEX-BY Table)
 - Also called index-by tables, these are key-value pairs, where keys can be numbers or strings.
 - Useful when you need dynamic access via key and don't care about order.

 Syntax:

```

DECLARE
    TYPE emp_table_type IS TABLE OF VARCHAR2(50) INDEX BY
    BINARY_INTEGER;
    emp_names emp_table_type;
BEGIN
    emp_names(101) := 'Vaibhav';
    emp_names(102) := 'Ajay';
    emp_names(103) := 'Ravi';

    FOR i IN 101..103 LOOP
        DBMS_OUTPUT.PUT_LINE('EmpID: ' || i || ', Name: ' || emp_names(i));
    END LOOP;
END;
/

```

 Example: Associating employee IDs with names

 Output:

```

EmpID: 101, Name: Vaibhav
EmpID: 102, Name: Ajay
EmpID: 103, Name: Ravi

```

 Real-Life Scenario:

- Used in billing systems to map product IDs to names/prices without using joins or lookup tables.

 Important Notes:

- Cannot be stored in DB tables.
- Extremely fast and efficient in memory.
- Supports dynamic key types (like strings too).

◊ 5. Nested Collections (Multi-level)

- A collection inside another collection — advanced concept but helpful in scenarios like hierarchical data (e.g., departments → employees).

 Syntax:

```

TYPE emp_list IS TABLE OF VARCHAR2(30);
TYPE dept_emp_table IS TABLE OF emp_list;

```

 Real-Life Scenario: A department stores a list of employees in each row as a nested list.

 Important Notes:

- Not commonly used in beginner projects.
- Needs special handling using loops and constructors.

Difference Between All Collection Types

Feature	RECORD	NESTED TABLE	VARRAY	ASSOCIATIVE ARRAY
Heterogeneous types	✓	✗	✗	✗
Fixed Size	✗	✗	✓	✗
Can be stored in DB	✗	✓	✓	✗
Indexed by	Field name	Integer	Integer	Integer/String
Supports methods	✗	✓ (EXTEND, DELETE)	✓ (EXTEND)	✓ (DELETE, EXISTS)

Important Notes:

- Use `%TYPE` and `%ROWTYPE` to make collections more flexible and maintainable.
- Always initialize collections before using.
- Collections are stored in memory, so use wisely in large data operations.
- Prefer `BULK COLLECT + FORALL` when performance matters.

Summary

Type	Use Case	Key Feature
RECORD	Store one row of different fields	Like a table row
NESTED TABLE	Dynamic array	Can grow/shrink, used in tables
VARRAY	Fixed size array	Predictable size, SQL support
ASSOCIATIVE ARRAY	Key-value structure	Memory-only, super fast
<code>%ROWTYPE</code>	Auto match table structure	Dynamic and reusable

◆ DYNAMIC SQL

- Dynamic SQL lets you build and run SQL statements at runtime.
- It's useful when the structure of SQL (table/column names, conditions) is not known until execution.
- Use dynamic SQL when static SQL won't work (e.g., dynamic table names, dynamic DML, etc.)

🧠 Real-Life Scenario:

- In a multi-client system, dynamically creating or updating tables like `client1_users`, `client2_users`, etc.

◊ 1. EXECUTE IMMEDIATE

- Most commonly used for dynamic SQL in PL/SQL.
- Supports DML, DDL, and anonymous PL/SQL blocks.
- It allows you to run a SQL statement dynamically at runtime.
- Commonly used when SQL structure (like table name or condition) changes based on logic or user input.
- It supports DML (INSERT, UPDATE, DELETE), DDL (CREATE, DROP), and PL/SQL blocks.
- It's a simple and faster alternative to the `DBMS_SQL` package for most use cases.

💻 Basic Syntax:

```
EXECUTE IMMEDIATE 'SQL statement';
```

💡 Example – Create Table Dynamically:

```
DECLARE
    table_name VARCHAR2(30) := 'temp_users';
BEGIN
    EXECUTE IMMEDIATE 'CREATE TABLE ' || table_name || '(id NUMBER, name
VARCHAR2(50))';
    DBMS_OUTPUT.PUT_LINE('Table created dynamically!');
END;
/
```

💻 Output:

```
Table created dynamically!
```

🧠 Real-Life Scenario:

- Useful in automation scripts where table names are passed dynamically.

◊ 2. Using BIND VARIABLES with EXECUTE IMMEDIATE

- Prevents SQL injection and improves performance by reusing parsed statements.
- Bind variables are placeholders (`:1`, `:2`, etc.) that let you pass values safely into your dynamic SQL.
- They prevent SQL injection and allow Oracle to reuse parsed SQL for performance.
- Values are passed using the `USING` clause after the SQL statement.
- You can bind IN, OUT, and IN OUT variables, but only by position, not by name.

💻 Syntax with USING:

```
EXECUTE IMMEDIATE 'UPDATE employees SET salary = :1 WHERE emp_id =
:2' USING new_salary, emp_id_val;
```

💡 Example:

```

DECLARE
    new_salary NUMBER := 5000;
    emp_id NUMBER := 101;
BEGIN
    EXECUTE IMMEDIATE
        'UPDATE employees SET salary = :1 WHERE emp_id = :2'
        USING new_salary, emp_id;
    DBMS_OUTPUT.PUT_LINE('Salary updated dynamically.');
END;
/

```

 Output:
Salary updated dynamically.

 Note: Only positional binds (:1, :2) are supported.

◊ 3. RETURNING INTO Clause

- Used to fetch values from DML statements dynamically.
- It allows you to capture output from a DML statement (INSERT, UPDATE, DELETE) during dynamic execution.
- You use **RETURNING INTO** to get values like inserted IDs or deleted column values immediately.
- Only works if the SQL returns exactly one row.
- This is supported in **EXECUTE IMMEDIATE** but not in **DBMS_SQL**.

 Syntax:

```

EXECUTE IMMEDIATE
    'DELETE FROM employees WHERE emp_id = :1 RETURNING name INTO :2'
    USING emp_id_val RETURNING INTO emp_name;

```

 Important Notes:

- Only works for single-row operations.
- **RETURNING INTO** is not supported in **DBMS_SQL**.
-

◊ 4. EXECUTE IMMEDIATE with INTO Clause (for SELECT)

- Used when you dynamically select a value into a variable.
- Used to dynamically run a SELECT query and fetch its result into a variable.
- It only works when the SELECT returns one and only one row.
- You can combine it with **USING** if the query has placeholders.
- Great for dynamic counts, max/min, single lookups, etc.

 Syntax:

```
EXECUTE IMMEDIATE 'SELECT COUNT(*) FROM employees' INTO v_count;
```

 Output:

Total count: 87

 Important Notes:

- Works only when the query returns exactly one row.

- ◊ 5. EXECUTE IMMEDIATE for Anonymous PL/SQL Block
 - You can build dynamic PL/SQL blocks and run them on the fly.
 - You can dynamically build and run entire PL/SQL logic blocks using EXECUTE IMMEDIATE.
 - Useful when the logic changes based on input or needs to be stored externally (e.g., logs, templates).
 - Helps in executing multiple statements or logic from strings.

 Example:

```
BEGIN
  EXECUTE IMMEDIATE '
    BEGIN
      DBMS_OUTPUT.PUT_LINE("Hello from dynamic block");
    END;';
  END;
/
```

 Output:

Hello from dynamic block

- ◊ 6. SQL Injection & Prevention
 - SQL Injection happens when user input is directly concatenated into a dynamic SQL string, allowing malicious SQL to run.
 - It's a major security threat if not handled.
 - Prevention: Always use bind variables, input validation, and Oracle's DBMS_ASSERT functions.
 - Bind variables ensure user input is treated as data, not code.

 Dangerous:

```
EXECUTE IMMEDIATE 'DELETE FROM employees WHERE name = "'||user_input||'"';
```

 Safe:

```
EXECUTE IMMEDIATE 'DELETE FROM employees WHERE name = :1' USING
user_input;
```

 Note: Use DBMS_ASSERT.simple_sql_name() to validate object names.

 Important Notes:

- Always avoid concatenating raw input.
- Bind parameters protect against attacks.

- ◊ 7. DBMS_SQL Package
 - Advanced method when you need flexibility and metadata control.
 - A lower-level API for fully dynamic SQL execution, offering more control and flexibility than EXECUTE IMMEDIATE.
 - Useful when you don't know the number or names of columns beforehand.
 - Supports describe columns, bind by name, fetch rows, get metadata dynamically.
 - More verbose and complex to use, but powerful in advanced use cases like report builders or ETL tools.

Syntax Example:

```
DECLARE
    cur_id  INTEGER;
    rows    NUMBER;
BEGIN
    cur_id := DBMS_SQL.OPEN_CURSOR;
    DBMS_SQLPARSE(cur_id, 'DELETE FROM employees WHERE department_id
= 90', DBMS_SQL.NATIVE);
    rows := DBMS_SQL.EXECUTE(cur_id);
    DBMS_SQLCLOSE_CURSOR(cur_id);
    DBMS_OUTPUT.PUT_LINE(rows || ' rows deleted');
END;
/
```

Output:

3 rows deleted

Real-Life Scenario:

- Used in dynamic reporting tools where column names are not known until runtime.

Important Notes:

- More verbose than **EXECUTE IMMEDIATE**.
- Supports dynamic column parsing, binding, and fetching.
- Use **EXECUTE IMMEDIATE** for 90% of use cases.
- Use **DBMS_SQL** only when:
 - You don't know the number/type of columns
 - You want to fetch metadata dynamically
 - Avoid string concatenation — it's slow and insecure.
 - Use **BULK COLLECT + FORALL** with dynamic SQL for better performance.

Difference Between EXECUTE IMMEDIATE vs DBMS_SQL

Feature	EXECUTE IMMEDIATE	DBMS_SQL
Simplicity	<input checked="" type="checkbox"/> Easy	<input type="checkbox"/> Complex
Positional Bind Support	<input checked="" type="checkbox"/> Yes	<input checked="" type="checkbox"/> Yes
Named Bind Support	<input type="checkbox"/> No	<input checked="" type="checkbox"/> Yes
Metadata Control	<input type="checkbox"/> No	<input checked="" type="checkbox"/> Yes (describe columns)
Performance	<input checked="" type="checkbox"/> Faster	<input type="checkbox"/> Slightly slower
Use Case	DML, DDL, PL/SQL blocks	Dynamic queries/columns

Real-Life Scenario Summary:

Use Case	Dynamic SQL Role
Multi-client database architecture	Create/alter tables dynamically
Data migration utilities	Load data into dynamic structures
Dynamic reporting systems	Fetch unknown columns/tables
Admin tools or ETL framework	Execute stored SQL scripts dynamically

Summary

Concept	Description
EXECUTE IMMEDIATE	For simple DDL, DML, SELECT, anonymous blocks
USING Clause	For binding variables to dynamic SQL
RETURNING INTO	To fetch values from DML queries dynamically
INTO Clause	For fetching single-row SELECT results
Dynamic PL/SQL Blocks	Run runtime PL/SQL logic dynamically
SQL Injection Prevention	Always use bind variables or DBMS_ASSERT
DBMS_SQL	Advanced option for unknown SQL at runtime

Bulk Operations

- designed to boost performance when dealing with large volumes of data.
- All four are directly part of or related to PL/SQL Bulk Operations.
- They are widely used in ETL, reporting, automation scripts, data migration, and high-volume batch processing.

◊ 1. BULK COLLECT

- **BULK COLLECT** is used to fetch multiple rows from a query into a collection variable (like VARRAY, Nested Table, Associative Array) in one go.
- **BULK COLLECT** is a PL/SQL feature used to fetch multiple rows from a SELECT statement into a collection (like Nested Table, VARRAY, or Associative Array) in a single context switch between SQL and PL/SQL engines.
- Normally, each row fetched requires communication between SQL and PL/SQL, which slows down performance for large datasets.
- With **BULK COLLECT**, all the rows are fetched at once, reducing the overhead and significantly improving performance for read-heavy operations.
- It's especially useful when processing large volumes of data like logs, orders, or transaction records.
- Collections populated using BULK COLLECT can be processed using loops, FORALL, or further business logic.

- Reduces the number of times PL/SQL must fetch rows one-by-one → improves performance significantly.
- ➔ Bulk Data Fetching : Used to fetch multiple rows from the database into a collection in one go.

 Syntax:

```
SELECT column1, column2
BULK COLLECT INTO collection1, collection2
FROM table_name
WHERE condition;
```

 Example:

```
DECLARE
    TYPE name_tab IS TABLE OF employees.ename%TYPE;
    TYPE sal_tab IS TABLE OF employees.sal%TYPE;
    v_names name_tab;
    v_sals sal_tab;
BEGIN
    SELECT ename, sal
    BULK COLLECT INTO v_names, v_sals
    FROM employees
    WHERE department_id = 10;

    FOR i IN v_names.FIRST .. v_names.LAST LOOP
        DBMS_OUTPUT.PUT_LINE(v_names(i) || ' earns ' || v_sals(i));
    END LOOP;
END;
/
```

 Output:

```
SMITH earns 3000
ADAMS earns 2500
```

 Real-Life Scenario:

- Used in report generation, data migration, or when fetching large datasets like 1000+ records in one query.

 Important Notes:

- Can be used with SELECT INTO, FETCH INTO, or RETURNING INTO.
- Supports LIMIT clause for controlling memory usage in loops.

◊ 2. FORALL

- **FORALL** is used to perform DML operations (INSERT, UPDATE, DELETE) on collection elements in bulk, instead of looping row-by-row.
- **FORALL** is used for bulk DML operations (INSERT, UPDATE, DELETE) based on data stored in a PL/SQL collection.
- Instead of executing a DML operation row by row using a loop, **FORALL** sends all the operations to the SQL engine in a single call, drastically reducing processing time.
- It is designed for high-speed write-heavy processing — like batch updates, mass insertions, or deletions.
- Unlike **FOR** loops, **FORALL** does not support SELECT operations — it is strictly for modifying data.

- It works best with index-by tables and nested tables, and each row is referenced by its index in the collection.
- Improves performance drastically by reducing context switches between PL/SQL and SQL engines
- → Bulk DML Processing: Used to perform DML (INSERT, UPDATE, DELETE) on multiple rows using a collection..

 Syntax:

```
FORALL index IN collection.FIRST .. collection.LAST
    DML_statement_using_collection(index);
```

 Example:

```
DECLARE
    TYPE id_array IS TABLE OF employees.emp_id%TYPE;
    v_ids id_array := id_array(101, 102, 103);
BEGIN
    FORALL i IN v_ids.FIRST .. v_ids.LAST
        DELETE FROM employees WHERE emp_id = v_ids(i);

    DBMS_OUTPUT.PUT_LINE('Bulk delete done.');
END;
/
```

 Output:

Bulk delete done.

 Real-Life Scenario:

- Used in ETL scripts, bulk data cleanup, or mass updates/inserts, e.g., marking users inactive in bulk.

 Important Notes:

- Works only with index-by or nested tables (not VARRAY).
- Can't use SELECT in FORALL; only DML.
- Use SAVE EXCEPTIONS to continue on error.

◊ 3. SAVE EXCEPTIONS

- Used with **FORALL** to handle errors gracefully during bulk DML. If one DML fails, the others still proceed.
- When using **FORALL**, if a single row operation fails (e.g., constraint violation), the entire operation stops unless you use **SAVE EXCEPTIONS**.
- **SAVE EXCEPTIONS** allows PL/SQL to continue processing the remaining rows even if one or more fail.
- After execution, the details of the failed rows (error index and error code) can be accessed using **SQL%BULK_EXCEPTIONS** collection.
- This approach is extremely valuable in production scenarios where skipping a few bad rows is better than halting the whole job.
- It ensures fault-tolerant, resilient bulk operations and is heavily used in real-time ETL and automation pipelines.
- Allows identifying which rows failed while continuing execution for valid rows.
- → Error Handling in FORALL: Used with **FORALL** to catch and continue on errors during bulk DML, instead of stopping execution.

 Syntax:

```
FORALL i IN collection.FIRST .. collection.LAST SAVE EXCEPTIONS DML_statement;
```

- ◆ Then handle the error:

```
EXCEPTION
  WHEN OTHERS THEN
    FOR i IN 1 .. SQL%BULK_EXCEPTIONS.COUNT LOOP
      DBMS_OUTPUT.PUT_LINE('Error on row ' || 
SQL%BULK_EXCEPTIONS(i).ERROR_INDEX);
      DBMS_OUTPUT.PUT_LINE(SQLERRM(-
SQL%BULK_EXCEPTIONS(i).ERROR_CODE));
    END LOOP;
```

 Example:

```
DECLARE
  TYPE emp_ids IS TABLE OF employees.emp_id%TYPE;
  v_ids emp_ids := emp_ids(101, 102, 999); -- 999 doesn't exist
BEGIN
  FORALL i IN v_ids.FIRST .. v_ids.LAST SAVE EXCEPTIONS
    DELETE FROM employees WHERE emp_id = v_ids(i);

  DBMS_OUTPUT.PUT_LINE('Bulk delete attempted.');
EXCEPTION
  WHEN OTHERS THEN
    FOR i IN 1 .. SQL%BULK_EXCEPTIONS.COUNT LOOP
      DBMS_OUTPUT.PUT_LINE('Failed at index: ' || 
SQL%BULK_EXCEPTIONS(i).ERROR_INDEX);
    END LOOP;
  END;
/

```

 Output:

```
Bulk delete attempted.
Failed at index: 3
```

 Real-Life Scenario:

- Useful in financial systems or record processing where skipping errors is better than failing the entire batch (e.g., retrying later).

 Important Notes:

- Works only with FORALL.
- Use **SQL%BULK_EXCEPTIONS** to analyze failed rows.

◊ 4. LIMIT Clause with BULK COLLECT

- Used to control how many rows you fetch at once in a loop, reducing memory overload.
- The **LIMIT** clause is used when fetching large amounts of data in chunks using a cursor + BULK COLLECT.
- It restricts the number of rows fetched into memory during each iteration, helping manage memory consumption effectively.
- Without LIMIT, fetching thousands of rows in one go can cause memory overflow and slow performance.
- Typically used in loops: fetch a batch of rows, process them, then fetch the next batch until all data is consumed.

- Very useful for processing millions of rows in manageable segments, ensuring both performance and resource optimization.
- → Controlled Bulk Fetching: Used to fetch data in chunks when processing very large data sets, for memory efficiency.

 Example:

```

LOOP
  FETCH c1 BULK COLLECT INTO v_names LIMIT 100;
  EXIT WHEN v_names.COUNT = 0;
  -- Process 100 records at a time
END LOOP;

```

 Important Notes:

- Especially useful when fetching large datasets in chunks.
- Prevents memory overflow and improves performance.

Difference Between: BULK COLLECT vs FORALL

Feature	BULK COLLECT	FORALL
Use	Fetching rows from SELECT	Executing DML (INSERT/UPDATE/DELETE)
Works With	SELECT, FETCH	INSERT, UPDATE, DELETE
Performance Gain	On fetching	On DML
Error Handling	Via EXCEPTION block	Use SAVE EXCEPTIONS for partial fails

 Important Notes:

- Always use collection types with bulk operations.
- Combine BULK COLLECT + FORALL for best performance.
- Use %ROWTYPE or RECORD if fetching full rows.
- Avoid using BULK COLLECT without **LIMIT** for very large datasets.
- Use **DBMS_OUTPUT.PUT_LINE** or logs to track processing.

Real-Life Scenario:

Real-Life Task	Bulk Operation Used
Archiving old records	BULK COLLECT (fetch) + FORALL (insert/delete)
Mass update of pricing in e-commerce	FORALL with SAVE EXCEPTIONS
Generating monthly reports in HR	BULK COLLECT with LIMIT clause
Syncing large tables in ETL pipelines	BULK COLLECT + FORALL

Summary

Concept	Purpose	Use Case
BULK COLLECT	Fetch multiple rows into collection in one go	Report generation, large fetch
FORALL	Perform bulk DML using collection values	Mass delete/update/insert
SAVE EXCEPTIONS	Allow some DML failures without halting all	Resilient operations with logs
LIMIT Clause	Restrict rows fetched per iteration	Avoid memory overflow in loops

Advanced Concepts & Utilities

UTL_FILE (File Handling in PL/SQL)

- UTL_FILE is a tool provided by Oracle that lets your PL/SQL program read from or write to text files (like .txt files).
 - These files are stored on the database server (not your personal computer).
 - You can use this when you want to save data from your table into a file, or read data from a file into your program.
 - UTL_FILE is an Oracle PL/SQL package used to read from and write to operating system text files from within the database.
 - It allows interaction with files located on the server's file system, making it useful for exporting logs, audit data, or reports.
 - Requires DIRECTORY object and write/read permissions.
- ◆ Why do we need it?
- Suppose you want to export a list of all employee names to a file for HR.
 - Or maybe save all errors into a log file during a batch job.
 - UTL_FILE lets you do that directly from the PL/SQL code without needing any external script.

Syntax:

```
DECLARE
    file_handler UTL_FILE.FILE_TYPE; -- Like opening a file in notepad
BEGIN
    -- Open the file named 'output.txt' for writing ('W')
    file_handler := UTL_FILE.FOPEN('MY_DIR', 'output.txt', 'W');

    -- Write one line into that file
    UTL_FILE.PUT_LINE(file_handler, 'Hello from PL/SQL!');

    -- Close the file (like saving and closing notepad)
    UTL_FILE.FCLOSE(file_handler);
END;
```

Example:

- You have a table of employees and want to write all names to a file:
 - ❖ This helps when HR asks for a quick export in a simple .txt file.

Output:

- A new file called **output.txt** will be created inside a special folder on the server.

- That file will contain:
`Hello from PL/SQL!`

 Real-Life Scenario:

- Error Logs: Save all failed records or error messages in a text file after a data load.
- Audit Report: Write every login record or transaction into a file for auditing purposes.
- Daily Report: Send a `.txt` file with daily sales to the sales team from the backend.

 Important Notes:

- You can't access your local laptop files — the file is created on the Oracle database server.
- You must ask your DBA (Database Admin) to create a folder (called DIRECTORY) and give you access to it.
- Without this permission, your program will show an error like:
`ORA-29283: Invalid file operation.`

 In Short:

- `UTL_FILE` is like a notepad for PL/SQL — you can open a file, write lines, and save it — all using PL/SQL code! 

 DBMS_SCHEDULER (Advanced Job Scheduling)

- It is a tool in Oracle used to automatically run programs at a fixed time or repeatedly (like every hour, day, week, etc.).
- You can think of it like an alarm clock for your PL/SQL code — you set a time, and it will execute your code automatically.
- Used for creating, managing, and monitoring scheduled jobs within the database.
- Replaces the older `DBMS_JOB` with more advanced features like logging, repeat intervals, dependencies, etc.

 Why do we need it?

- Suppose you have a cleanup job that must run every night at 11 PM.
 - Or you want to send daily sales reports to your team.
-  This tool helps automate those jobs without manually running them.

 Syntax:

```

BEGIN
  DBMS_SCHEDULER.CREATE_JOB (
    job_name      => 'test_job',
    job_type      => 'PLSQL_BLOCK',
    job_action    => 'BEGIN DBMS_OUTPUT.PUT_LINE("Hello World"); END;',
    start_date    => SYSTIMESTAMP,
    repeat_interval => 'FREQ=DAILY; BYHOUR=9; BYMINUTE=0', -- Every day at
9 AM
    enabled       => TRUE
  );
END;

```

 Example:

- You want to print “Hello World” every day at 9 AM.
- This job will start at the given time and run automatically.

 Output:

- The above code runs every day at 9:00 AM and prints:
Hello World
- You can replace this with any logic like inserting data, generating reports, etc.

 Real-Life Scenario:

- Auto-send email reminders daily
- Generate a monthly billing report
- Clean up temporary files or data from tables weekly

 Important Notes:

- Supports external jobs (e.g., OS-level scripts).
- Can be monitored via ALL_SCHEDULER_JOBS.
- It is more powerful than **DBMS_JOB** (which is older).
- You can monitor, pause, or stop jobs easily using this tool.
- Uses standard scheduling format like:
FREQ=HOURLY, FREQ=DAILY, FREQ=WEEKLY etc.

 DBMS_JOB (Basic Job Scheduler - Legacy)

- It is an older method used to schedule jobs in Oracle.
- It's simple and good for small or one-time jobs.
- Works inside the database, and is useful for beginners or simple tasks.
- Older packages used to schedule and run background jobs like procedures or anonymous blocks.
- Simpler but less powerful than DBMS_SCHEDULER.

 Why do we need it?

- You want to schedule a stored procedure to run every hour.
- You want a small job to clean a table daily — DBMS_JOB can do that quickly.

 Syntax:

```
DECLARE
    job_no NUMBER;
BEGIN
    DBMS_JOB.SUBMIT(job_no,
        'BEGIN DBMS_OUTPUT.PUT_LINE("Hello from DBMS_JOB"); END;',
        SYSDATE,
        'SYSDATE + 1/24' -- every 1 hour
    );
END;
```

 Example:

- Print a message every 1 hour using a scheduled job.

 Output:

Hello from DBMS_JOB
Repeated every hour automatically.

 Real-Life Scenario:

- Run a cleanup job to delete logs daily
- Update data in a table every night at midnight

 **Important Notes:**

- Still supported, but DBMS_SCHEDULER is preferred.
- DBMS_JOB is still available, but Oracle recommends using DBMS_SCHEDULER now.
- It's easy to use but less flexible (can't do external scripts, for example).

❖ AUTONOMOUS TRANSACTIONS

- A normal transaction in PL/SQL (like inserting or updating data) can only commit or rollback once when your block ends.
- But with Autonomous Transaction, you can commit/rollback immediately — even while the main code is still running.
- Used when you need to commit/rollback a transaction independently from the main transaction block.
- Helpful in logging, auditing, or error-tracking without affecting the main logic.

◆ Why do we need it?

- Imagine your code fails and you want to log an error in a table, but the full block is rolling back.
- You still want the error log to stay saved — so you use Autonomous Transaction to insert it separately and commit it immediately.

 **Syntax:**

```
DECLARE
    PRAGMA AUTONOMOUS_TRANSACTION; -- Special keyword
BEGIN
    INSERT INTO error_logs VALUES ('Something went wrong');
    COMMIT; -- This will happen separately from the main block
END;
```

 **Example:**

- Inside an exception block, log error messages without affecting main logic.

 **Output:**

The log will be inserted even if your main transaction fails.

 **Real-Life Scenario:**

- Logging error messages during failures
- Save log or audit information even if your main task fails
- Insert a record for tracking access, regardless of main logic

 **Important Notes:**

- Must use COMMIT or ROLLBACK inside.
- Commonly used in exception sections.
- You must always use COMMIT or ROLLBACK inside the autonomous block.
- Works well for logs, tracking, audit, or recording exceptions.
- Declared using:

PRAGMA AUTONOMOUS_TRANSACTION;

❖ SEQUENCES

- A sequence is a special object in Oracle that automatically generates numbers (like 1, 2, 3, ...).
- You mostly use it when you need a unique value each time — like for primary keys.
- Objects used to generate unique numbers, often for primary keys.
- Can be used in triggers, INSERTs, and programmatically.

- ◆ Why do we need it?
 - To give each new row in a table a unique ID.
 - Avoids manual work or duplicates.

 Syntax:

```
CREATE SEQUENCE emp_seq
START WITH 1
INCREMENT BY 1;
Use it like:
INSERT INTO employees (id, name) VALUES (emp_seq.NEXTVAL, 'Vaibhav');
```

 Example:

Adding employee IDs automatically:
 ID: 1 → Vaibhav
 ID: 2 → Rahul

 Output:

A new ID is created automatically each time you insert.

 Real-Life Scenario:

- Auto-generating Order IDs
- Assigning unique Order Numbers
- Generating Invoice IDs
- Tracking user registrations

 Important Notes:

- Use CURRVAL only after NEXTVAL has been used in session.
- Use NEXTVAL to get the next number.
- Use CURRVAL to get the current value used in the session.
- Sequences can also jump values (due to caching).

❖ SYNONYMS

- Alias for a database object (table, view, sequence, procedure).
- Makes it easier to refer to an object without needing the schema name.
- A synonym is like a nickname or shortcut name for a table, view, or object.
- It makes code shorter and helps you avoid writing full object names.

◆ Why do we need it?

- To simplify code when table names are long
- To access objects from other schemas easily

 Syntax:

```
CREATE SYNONYM emp FOR hr.employees;
Now you can just:
SELECT * FROM emp;
```

 Example:

Instead of:
 SELECT * FROM hr_department_data_view_2025;
 Use:
 SELECT * FROM dept_data;

 Output:

Same data, shorter code.

Real-Life Scenario:

- Shortening access to cross-schema tables
- Developers access tables from other users
- Simplify code for views or stored procedures
- Centralize access to common database objects

Important Notes:

- Can be public or private.
- You can create public or private synonyms
- Synonyms don't store data — they just refer to other objects

◆ Alias

- Temporary name given to a column or table in a SQL query.
- Exists only during that query execution.

◆ Synonym

- A permanent alternate name for a database object (table, view, sequence, etc.).
- Stored in the database dictionary.

❖ OBJECT TYPES (User-Defined Types)

- They let you create your own custom data type — a combination of variables (attributes) and actions (methods).
- Think of it like defining your own "employee" object with name, ID, and a display method.
- Used to define custom complex data structures in the database.
- Enables OOP-style programming (Encapsulation, Methods).

◆ Why do we need it?

- To model real-world entities like Employee, Customer, Product.
- Useful in object-oriented PL/SQL programs.

Syntax:

```
CREATE OR REPLACE TYPE emp_type AS OBJECT (
    emp_id NUMBER,
    emp_name VARCHAR2(100),
    MEMBER PROCEDURE show_info
);
```

Then create a body:

```
CREATE OR REPLACE TYPE BODY emp_type AS
    MEMBER PROCEDURE show_info IS
        BEGIN
            DBMS_OUTPUT.PUT_LINE('ID: ' || emp_id || ', Name: ' || emp_name);
        END;
    END;
```

Example:

Creating an object for an employee and showing info:

```
DECLARE
    e emp_type := emp_type(101, 'Vaibhav');
BEGIN
    e.show_info;
END;
```

Output:

ID: 101, Name: Vaibhav

 Real-Life Scenario:

- Passing complex data to/from procedures
- Represent customers, products, vehicles with properties
- Use object tables for storing structured data
- Enable reuse and better design of complex systems

 Important Notes:

- Can be used in tables or collections.
- Object types are used in Object-Relational databases
- Can include methods like in OOP (Object-Oriented Programming)

❖ PERFORMANCE TUNING (PL/SQL & SQL)

- It is the process of improving the speed and efficiency of your PL/SQL or SQL queries and procedures.
- Helps in reducing load time, improving user experience, and avoiding server overload.
- Process of optimizing database operations to ensure maximum performance and efficiency.
- Involves identifying slow SQLs, optimizing queries, reducing context switches, etc.

◆ Why do we need it?

- Slow queries affect app performance.
- Helps in reducing cost and improving scalability.

◆ Common Tuning Areas:

- SQL Query Tuning:
 - Use **EXPLAIN PLAN** to analyze slow queries
 - Create proper **INDEXES**
 - Avoid unnecessary **DISTINCT, GROUP BY, ORDER BY**
- PL/SQL Tuning:
 - Use **BULK COLLECT, FORALL** to handle large data
 - Avoid unnecessary context switches
 - Use **PRAGMA UDF** for faster functions in SQL

◆ Key Techniques:

- Use **BULK COLLECT / FORALL**
- Use **INDEXES** wisely
- Avoid unnecessary loops
- Use proper joins and WHERE clauses
- Analyze execution plans

 Example of Tuning:

Before:

```
FOR i IN (SELECT * FROM employees) LOOP
    -- process row
END LOOP;
```

After:

```
SELECT * BULK COLLECT INTO emp_list FROM employees;
FOR i IN emp_list.FIRST .. emp_list.LAST LOOP
    -- process row
END LOOP;
```

 Real-Life Scenario:

- Optimizing ETL jobs, monthly reports
- You notice a report taking 10 minutes — tuning brings it to 10 seconds
- A slow procedure during month-end slows everything — tuning improves performance

 **Important Notes:**

- Use EXPLAIN PLAN, DBMS_PROFILER, and AWR reports
- Avoid implicit cursors inside loops
- Always test after tuning
- Use tools like [SQL Developer](#), [Toad](#), or [AWR Reports](#)
- Indexes, proper joins, and good logic are the key

 **SUMMARY:**

Topic	Use Case	Tool/Package
UTL_FILE	File reading/writing	UTL_FILE
DBMS_SCHEDULER	Job scheduling with advanced features	DBMS_SCHEDULER
DBMS_JOB	Legacy job scheduling	DBMS_JOB
AUTONOMOUS TRANSACTION	Independent logging or commits	PRAGMA + COMMIT
SEQUENCES	Unique key generation	Sequence object
SYNONYMS	Object aliasing	CREATE SYNONYM
OBJECT TYPES	OOP style custom types	CREATE TYPE
PERFORMANCE TUNING	Improving overall DB performance	EXPLAIN PLAN, AWR, etc.

◆ Debugging with Messages

❖ DBMS_OUTPUT – Basic Debugging with Messages

◊ Debugging in PL/SQL

- Debugging helps you find and fix issues in your PL/SQL code (procedures, functions, triggers, etc.). It's like investigating what's going wrong step-by-step.

◆ What Exactly Is It?

- **DBMS_OUTPUT** is a built-in Oracle debugging utility.
- It lets you print out messages, values of variables, or processing checkpoints to the output screen.
- Very useful during development to check how your code behaves at each step

◆ IMP POINT: It doesn't interrupt your code like errors. It silently shows messages — you must enable it in tools like SQL Developer or SQL*Plus.

💻 Syntax:

```
BEGIN
    DBMS_OUTPUT.PUT_LINE('This is a debug message');
END;
```

You can also display values:

```
DECLARE
    salary NUMBER := 40000;
BEGIN
    DBMS_OUTPUT.PUT_LINE('Salary is: ' || salary);
END;
```

💡 Example:

Let's debug a bonus calculation:

```
DECLARE
    emp_id NUMBER := 101;
    bonus NUMBER;
BEGIN
    bonus := 2000;
    DBMS_OUTPUT.PUT_LINE('Bonus assigned to emp ' || emp_id || ' is ' || bonus);
END;
```

🖨️ Output:

Bonus assigned to emp 101 is 2000

🧠 Real-Life Scenario:

- You're calculating discounts for customers and you want to track the applied discount
- You're troubleshooting a PL/SQL trigger and want to see which condition passed or failed

📝 Important Notes:

- Always use **SET SERVEROUTPUT ON** in SQL*Plus or enable it in SQL Developer
- Doesn't work well in live systems — not visible to end users
- Good for development/testing only

❖ SQL Developer Debugger – Advanced, Visual Debugging

◆ What Is It?

- A graphical debugger tool inside Oracle SQL Developer
- Helps you pause and inspect PL/SQL code line-by-line
- You can view variable values, change them, and simulate logic flows

◆ Why Is It Powerful?

- Unlike **DBMS_OUTPUT**, it doesn't just show text.
- You can:
 - Set breakpoints
 - Watch variables as they change
 - Step into loops, conditionals, nested procedures

✓ Steps to Use:

1. Open your PL/SQL procedure
2. Right-click > Debug
3. Add breakpoints (click left of line numbers)
4. Run the debugger
5. Use tabs like "Watches", "Stack", "Variables" to observe live changes

💡 Example:

Let's debug a procedure:

```
CREATE OR REPLACE PROCEDURE calc_bonus(emp_id NUMBER) IS
    salary NUMBER := 50000;
    bonus NUMBER;
BEGIN
    IF salary > 30000 THEN
        bonus := salary * 0.1;
    ELSE
        bonus := salary * 0.05;
    END IF;
    DBMS_OUTPUT.PUT_LINE('Bonus: ' || bonus);
END;
```

- ➡ You set breakpoints on **IF** and **DBMS_OUTPUT** lines
- ➡ You run the debugger and watch how bonus is calculated

🧠 Real-Life Scenario:

- You get a bug in salary update procedure, but no errors shown
- You debug and realize the bonus line never runs because salary logic was skipped

📝 Important Notes:

- You may need DEBUG CONNECT SESSION privilege
- Best used for testing before deployment
- Not usable in live systems — only for developers

❖ Logging (Custom Debug Logs) – Production-Level Debugging

- ◆ What Is Logging?
 - Instead of printing on screen, you insert debug messages into a custom table.
 - These messages can be retrieved later and help in understanding what went wrong.
- ◆ Why Use Logging?
 - **DBMS_OUTPUT** doesn't work in production
 - Logging keeps permanent records
 - You can include timestamps, error messages, usernames, inputs

Syntax:

Step 1. – Create a table:

```
CREATE TABLE debug_log (
    log_time TIMESTAMP,
    module_name VARCHAR2(100),
    message   VARCHAR2(4000)
);
```

Step 2. – Insert logs:

```
INSERT INTO debug_log VALUES (SYSTIMESTAMP, 'Add_Employee', 'Insert started');
```

Step 3. – Review logs later:

```
SELECT * FROM debug_log WHERE module_name = 'Add_Employee';
```

Example:

```
DECLARE
    emp_name VARCHAR2(100) := 'Vaibhav';
BEGIN
    INSERT INTO debug_log VALUES (SYSTIMESTAMP, 'Emp_Insert', 'Started insert for ' || emp_name);
    INSERT INTO employees(name) VALUES (emp_name);
    INSERT INTO debug_log VALUES (SYSTIMESTAMP, 'Emp_Insert', 'Insert completed');
END;
```

Output:

log_time	module_name	message
08-JUN-2025 10:12AM	Emp_Insert	Started insert for Vaibhav
08-JUN-2025 10:12AM	Emp_Insert	Insert completed

Important Notes:

- Add indexes to log tables for performance
- Archive or delete old logs to save space
- Can combine with **EXCEPTION WHEN OTHERS THEN** block for error logs

 Real-Life Scenario:

- You want to track why data is failing during night batch jobs
- Logs tell you exactly where failure happened

 Important Notes:

- Use **DBMS_OUTPUT** when you're developing or practicing
- Use the Debugger for professional-level tracing and fixes
- Use Logging when code is live or in UAT/prod, and you can't see the output directly

 Summary Table:

Technique	What It Does	When To Use
DBMS_OUTPUT	Prints text to screen	Simple testing & variable checks
SQL Debugger	Step-by-step debugging	Full control over execution flow
Logging	Saves messages to a table	For production-level, trackable debugging

◊ 1. EXCEPTION HANDLING (Error Trapping)

- Exception handling is a key PL/SQL feature that helps manage runtime errors in a controlled way, allowing the program to handle errors gracefully without abrupt termination.
- This ensures smoother user experience and better program stability.

◆ You can trap specific exceptions like:

- **NO_DATA_FOUND**
- **TOO_MANY_ROWS**
- **ZERO_DIVIDE**
- or a generic one like **WHEN OTHERS THEN**

 Syntax:

```
BEGIN
    -- Code that might raise exception
EXCEPTION
    WHEN NO_DATA_FOUND THEN
        DBMS_OUTPUT.PUT_LINE('No record found.');
    WHEN OTHERS THEN
        DBMS_OUTPUT.PUT_LINE('Error: ' || SQLERRM);
END;
```

 Example:

Handling a division by zero error using exceptions.

 Output:

Displays a user-friendly message instead of abrupt termination.

 Real-Life Scenario:

- When a product price is divided by quantity and quantity is zero.

 Important Notes:

- Use specific exceptions when possible; use **WHEN OTHERS** carefully.

◊ 2. SQLERRM & SQLCODE (Error Message & Code)

- These are built-in functions used inside exception blocks to retrieve the error message and error number respectively.

 Syntax:

```
BEGIN
    -- Risky operation
    EXCEPTION
        WHEN OTHERS THEN
            DBMS_OUTPUT.PUT_LINE('Error Code: ' || SQLCODE);
            DBMS_OUTPUT.PUT_LINE('Error Message: ' || SQLERRM);
    END;
```

 Real-Life Scenario:

- Logging these values in an audit/error log table.

 Important Notes:

- **SQLERRM** gives detailed error message
- **SQLCODE** returns negative value for error

◊ 3. AUTONOMOUS TRANSACTIONS – For Independent Logging

- Allows you to perform independent DML operations like logging errors or messages, even if the main transaction fails or rolls back.

 Syntax:

```
CREATE OR REPLACE PROCEDURE log_error(p_msg VARCHAR2) IS
    PRAGMA AUTONOMOUS_TRANSACTION;
BEGIN
    INSERT INTO error_log (log_time, message) VALUES (SYSDATE, p_msg);
    COMMIT;
END;
```

 Example:

Use this inside your EXCEPTION block to make sure logs are saved even if the main transaction is rolled back.

 Output:

Data gets inserted into the **error_log** table even if outer block fails.

 Real-Life Scenario:

- Logging failed payment transactions.

 Important Notes:

- Requires an explicit **COMMIT**
- Can't read or write from parent transaction's context

◊ 4. TRACING with DBMS_TRACE

- Oracle-supplied package that traces line-by-line execution of PL/SQL code to help you detect the exact flow and where the issue happens.

 Syntax:

```
EXEC  
DBMS_TRACE.SET_PLSQL_TRACE(DBMS_TRACE.TRACE_ALL_CALLS);
```

 Real-Life Scenario:

- Debugging a complex stored procedure to track flow and logic.

 Important Notes:

- Needs configuration
- Mostly used in dev/test environments

◊ 5. DBMS_DEBUG / DBMS_DEBUG_JDWP

- Packages that allow debugging of stored procedures from remote tools (e.g., via SQL Developer or other IDEs).

 Example:

Setting breakpoints, stepping through code, inspecting variables.

 Real-Life Scenario:

- Used by teams where developers debug code from IDEs without direct database console access.

 Important Notes:

- Requires extra setup
- More relevant in team and enterprise environments

◊ 6. Logging to Files using UTL_FILE

- Instead of logging to tables, logs can be written to text files using the UTL_FILE package.

 Syntax:

```
DECLARE  
    file_handler UTL_FILE.FILE_TYPE;  
BEGIN  
    file_handler := UTL_FILE.FOPEN('LOG_DIR', 'debug.log', 'A');  
    UTL_FILE.PUT_LINE(file_handler, 'Error at step 1');  
    UTL_FILE.FCLOSE(file_handler);  
END;
```

 Real-Life Scenario:

- Used in scheduled jobs or ETL processes to write logs for each run.

 Important Notes:

- Directory must be created and granted to the user
- Handle file I/O exceptions

◊ 7. Custom Logging Procedures

- Creating a reusable procedure like `log_step()` to insert logs with procedure name, step name, and timestamp.

 Syntax:

```
PROCEDURE log_step(p_proc VARCHAR2, p_step VARCHAR2, p_msg  
VARCHAR2) IS
```

```
BEGIN  
    INSERT INTO app_log VALUES (SYSDATE, p_proc, p_step, p_msg);  
    COMMIT;  
END;
```

 Real-Life Scenario:

- Used across all procedures to maintain standard logs.

 Important Notes:

- Helps in auditing and debugging
- Pair with autonomous transactions for safety

◊ 8. SHOW ERRORS Command

- Command to view compilation errors of procedures, packages, functions, etc.

 Syntax:

```
SHOW ERRORS PROCEDURE my_procedure;
```

 Real-Life Scenario:

- Fixing compile-time errors in procedures/packages quickly.

 Important Notes:

- Available in SQL*Plus, SQL Developer, and other tools.

◊ 9. DBMS_PROFILER – Code Performance Analyzer

- Helps to measure how much time each line of PL/SQL code takes to execute.

 Real-Life Scenario:

- Useful for identifying slow-performing queries or blocks in procedures.

 Important Notes:

- Used more in the optimization phase than debugging, but helps identify performance issues.

◊ 10. PLSQL_WARNINGS

- Enabling this setting during the session allows you to see compile-time warnings.

 Syntax:

```
ALTER SESSION SET PLSQL_WARNINGS = 'ENABLE:ALL';
```

 Real-Life Scenario:

- Detecting dead code, unhandled exceptions, or performance risks at compile-time.

 Important Notes:

- Helps to proactively catch bugs.

 Summary:

- PL/SQL provides a wide range of debugging and error-tracking tools, from simple output to full-on trace and profiling systems.
- As you build more complex applications, combining these techniques gives better insight into what your code is doing and why something went wrong.

◆ PL/SQL Real-Time Applications

◊ 1. Reusable Packages

◆ What Are Packages?

- Packages in PL/SQL are collections of related procedures, functions, variables, and cursors grouped together in two parts:
 - Specification: Declaration (just names, no logic)
 - Body: Definition (actual logic inside)
- They help in modular programming, so you don't need to write the same code again and again. Also, it hides private logic from users and improves performance by loading once into memory.

💡 Example:

```
-- PACKAGE SPEC
CREATE OR REPLACE PACKAGE emp_pkg IS
    PROCEDURE insert_emp(p_id NUMBER, p_name VARCHAR2);
    FUNCTION get_emp_count RETURN NUMBER;
END emp_pkg;
/

-- PACKAGE BODY
CREATE OR REPLACE PACKAGE BODY emp_pkg IS
    PROCEDURE insert_emp(p_id NUMBER, p_name VARCHAR2) IS
        BEGIN
            INSERT INTO employees VALUES(p_id, p_name);
        END;

    FUNCTION get_emp_count RETURN NUMBER IS
        v_count NUMBER;
        BEGIN
            SELECT COUNT(*) INTO v_count FROM employees;
            RETURN v_count;
        END;
END emp_pkg;
/
```

◆ How to Use:

```
BEGIN
    emp_pkg.insert_emp(101, 'Vaibhav');
    DBMS_OUTPUT.PUT_LINE('Employee Count: ' || emp_pkg.get_emp_count);
END;
```

💻 Output:

```
Employee Count: 11
```

🧠 Real-Life Scenario:

- Imagine you're building an HR software where employee insert and count functions are used in many places. Instead of repeating the same logic, you write it once in a package and reuse it across the system.

◊ 2. PL/SQL Jobs using **DBMS_SCHEDULER**

◆ What Are They?

- Jobs are background tasks that run automatically at a fixed time — just like setting an alarm.
- These are helpful for doing tasks like taking backup, sending emails, generating reports without human action.
- In PL/SQL, we use **DBMS_SCHEDULER** to schedule these jobs.

💡 Example:

```
BEGIN
    DBMS_SCHEDULER.create_job (
        job_name      => 'daily_report_job',
        job_type      => 'PLSQL_BLOCK',
        job_action    => 'BEGIN generate_report(); END;',
        start_date    => SYSTIMESTAMP,
        repeat_interval => 'FREQ=DAILY; BYHOUR=23;',
        enabled       => TRUE
    );
END;
/
```

🧠 Real-Life Scenario:

- Let's say you have an online shopping app. You want to send a sales report to your team every night at 11 PM. You write a report procedure and schedule it using **DBMS_SCHEDULER**.

◊ 3. Backend Processes (with Triggers or Procedures)

◆ What Are Backend Processes?

- Backend processes are tasks that happen behind the scenes when users interact with your application.
- They run automatically based on some event, like data insert/update.
- Example: Stock update after placing an order.
- ◆ They are written using triggers or procedures.

💡 Example: Trigger to update stock

```
CREATE OR REPLACE TRIGGER trg_stock_update
AFTER INSERT ON orders
FOR EACH ROW
BEGIN
    UPDATE stock
    SET quantity = quantity - :NEW.qty
    WHERE product_id = :NEW.product_id;

    INSERT INTO order_log(order_id, status)
    VALUES(:NEW.order_id, 'Processed');
END;
/
```

🧠 Real-Life Scenario:

- In an e-commerce site like Flipkart, when someone buys a product, the stock gets updated automatically in the database. This is done using backend logic like this trigger.

◊ 4. Complex Data Logic in PL/SQL

◆ What Is Complex Logic?

- PL/SQL allows writing decision-making or business rules using IF-ELSE, CASE, loops, etc.
- You can use it to calculate tax, discount, commission, salary, EMI, etc.

💡 Example: Tax Calculation Function

```
CREATE OR REPLACE FUNCTION calculate_tax(salary NUMBER) RETURN
NUMBER IS
    tax NUMBER;
BEGIN
    IF salary <= 250000 THEN
        tax := 0;
    ELSIF salary <= 500000 THEN
        tax := salary * 0.05;
    ELSE
        tax := salary * 0.1;
    END IF;
    RETURN tax;
END;
/
```

◆ Usage:

```
BEGIN
    DBMS_OUTPUT.PUT_LINE('Tax: ' || calculate_tax(300000));
END;
```

💻 Output:

```
Tax: 15000
```

🧠 Real-Life Scenario:

- In a payroll system, this function will be used for every employee to calculate the tax based on their salary automatically before generating payslips.

◊ 5. Real-Time Case Study: Insurance Premium Calculation

- Let's build a small PL/SQL project example for Insurance Domain.
- Objective: Calculate premium based on vehicle age and risk type (HIGH/LOW) and store it in a table.

💡 Example:

Tables:

```
CREATE TABLE vehicles (
    id NUMBER,
    owner_name VARCHAR2(100),
    vehicle_age NUMBER,
    risk_category VARCHAR2(10)
);
```

```
CREATE TABLE premiums (
    id NUMBER,
    vehicle_id NUMBER,
    premium_amount NUMBER
);
```

Example:

```
Procedure:  
CREATE OR REPLACE PROCEDURE calculate_premium IS  
    CURSOR c1 IS SELECT id, vehicle_age, risk_category FROM vehicles;  
    premium NUMBER;  
BEGIN  
    FOR rec IN c1 LOOP  
        IF rec.vehicle_age < 3 THEN  
            premium := 5000;  
        ELSIF rec.vehicle_age BETWEEN 3 AND 10 THEN  
            premium := 7000;  
        ELSE  
            premium := 10000;  
        END IF;  
  
        IF rec.risk_category = 'HIGH' THEN  
            premium := premium + 3000;  
        END IF;  
  
        INSERT INTO premiums VALUES(premiums_seq.NEXTVAL, rec.id, premium);  
    END LOOP;  
END;  
/
```

Real-Life Scenario:

- This logic can be used in a real insurance system like PolicyBazaar or HDFC Ergo, where customers' vehicle details are entered and premiums are calculated dynamically.

Important Notes:

- PL/SQL is not just about writing queries — it's about building real business logic.
- Reusable packages save time and avoid duplicate code.
- Scheduler jobs help automate daily tasks like backups or reports.
- Backend logic keeps the system intelligent and responsive.
- Case studies help combine everything you've learned into a working system.

Summary:

Concept	Description	Real-World Usage
Packages	Group functions/procedures together	Used across modules (reusable code)
Scheduler Jobs	Auto run tasks like reports	Nightly backups, auto emails
Backend Processes	Auto run when data changes	Stock update, logging
Complex Logic	Decision-making logic	Tax, salary, discounts
Project Case Study	End-to-end PL/SQL usage	Insurance premium calc system

⭐ Additional Concepts You Might Encounter (already included indirectly):

◆ Data Modeling Concepts:

❖ What is Data Modeling?

- Data Modeling is the process of designing how data will be stored, related, and accessed in a database.
- It's like creating a blueprint of your database.
- In Simple Words:
 - Just like a civil engineer draws a map of a house, a data modeler draws a map of how tables, relationships, and data should be designed in a system.

🎯 Purpose:

- Define structure of data (tables, columns, relationships)
- Remove redundancy
- Improve performance
- Help business understand data flow

Ἑ Types of Data Models:

Type	Description
Conceptual	High-level model, for business understanding (ER Diagram)
Logical	More detailed, defines attributes, primary/foreign keys
Physical	Actual database structure with data types, constraints, etc.

❖ Normalization

Normalization is the process of organizing data in a database to:

- Reduce data redundancy
- Ensure data integrity
- Break large tables into smaller related tables

◆ Why Normalize?

画卷 Let's say you have this table:

EMP_ID	EMP_NAME	DEPT_NAME	DEPT_LOCATION
101	Raj	HR	Mumbai
102	Neha	HR	Mumbai
103	Amit	IT	Pune

Here, "HR" and "Mumbai" are repeated — this causes data redundancy.

- ◆ Types (Forms) of Normalization:
 1. 1NF (First Normal Form)
 2. 2NF (Second Normal Form)
 3. 3NF (Third Normal Form)
 4. BCNF (Boyce-Codd Normal Form)

◆ 1NF (First Normal Form)

Rules:

- Atomic values only (no multivalued columns)
- Unique column names
- Rows are uniquely identified

 Example:

EMP_ID	EMP_NAME	PHONE_NUMBERS
101	Raj	9999999999,88888888

● This is not 1NF. It contains multi-valued fields.

● Convert to 1NF:

EMP_ID	EMP_NAME	PHONE_NUMBER
101	Raj	9999999999
101	Raj	8888888888

◆ 2NF (Second Normal Form)

Rules:

- Should be in 1NF
- No partial dependency (applies to composite keys)

 Example:

EMP_ID	PROJECT_ID	EMP_NAME	PROJECT_NAME
101	101	Raj	Project A

Here, **EMP_NAME** depends only on **EMP_ID**, not full key (**EMP_ID + PROJECT_ID**) — so break it:

- ◆ Convert to 2NF:
- EMPLOYEE Table → (EMP_ID, EMP_NAME)
- PROJECT Table → (PROJECT_ID, PROJECT_NAME)
- EMP_PROJECT → (EMP_ID, PROJECT_ID)

◆ 3NF (Third Normal Form)

Rules:

- Should be in 2NF
- No transitive dependency (non-key column depending on another non-key)

💡 Example:

EMP_ID	EMP_NAME	DEPT_ID	DEPT_NAME
--------	----------	---------	-----------

DEPT_NAME depends on **DEPT_ID**, not directly on **EMP_ID**.

◆ Convert to:

- EMPLOYEE → (EMP_ID, EMP_NAME, DEPT_ID)
- DEPARTMENT → (DEPT_ID, DEPT_NAME)

◆ BCNF (Boyce-Codd Normal Form)

- Advanced version of 3NF
- Every determinant must be a candidate key

Use when tables have more than one candidate key and 3NF still has redundancy

⌚ Summary Table for Normalization

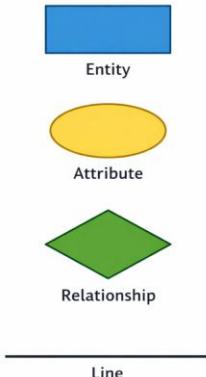
Form	Goal	Rule	Example Fix
1NF	Atomic Columns	No multi-values	Split multi-phone rows
2NF	Remove Partial Dependency	Full primary key dependency	Break composite key tables
3NF	Remove Transitive Dep.	Non-key depending on other non-key	Move dept to new table
BCNF	Candidate Key Clarity	Every determinant is a key	Handle unusual key structures

❖ ER Diagrams (Entity-Relationship Diagrams)

- ◆ What is an ER Diagram?
- An ER (Entity-Relationship) Diagram is a graphical representation of:
 - Entities (tables)
 - Attributes (columns)
 - Relationships (keys & foreign keys)

It shows how tables are related in a database.

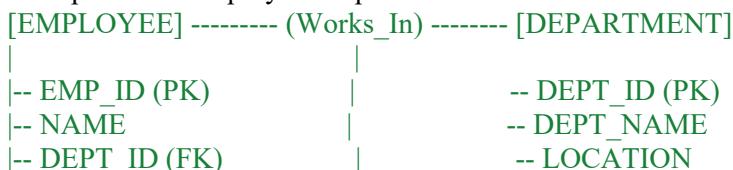
ER Diagram Components



Basic Components:

Symbol	Meaning
Rectangle	Entity (like TABLE)
Ellipse	Attribute (like COLUMN)
Diamond	Relationship (like "works in")
Line	Connects entity to relationship

💡 Example Case: Employee–Department



- ◆ This ER diagram represents:
 - Employee belongs to a department
 - DEPT_ID in EMPLOYEE is a foreign key to DEPARTMENT
- ◆ Cardinality:
 - One to One (1:1) – e.g., each employee has one laptop
 - One to Many (1:N) – e.g., one department has many employees
 - Many to Many (M:N) – e.g., students and courses
- ◆ Primary Key, Foreign Key in ER
 - Primary Key (PK) – Uniquely identifies a record (like EMP_ID)
 - Foreign Key (FK) – Points to a PK in another table (like DEPT_ID in employee)

- ❖ Use arrows or FK notation in ER diagrams to represent links
- ❖ Real-Time Scenario of Data Modeling
 - ◆ Suppose you're building a School Management System:
 - Entities:
 - Student
 - Course
 - Teacher
 - Department
 - Relationships:
 - A student enrolls in many courses (M:N)
 - A teacher teaches many courses (1:N)
 - Each course belongs to one department (1:N)
 - So you design:
 - STUDENT (SID, Name, Age)
 - COURSE (CID, Title, Dept_ID)
 - TEACHER (TID, Name)
 - DEPARTMENT (Dept_ID, Dept_Name)
 - ENROLLMENT (SID, CID) → handles many-to-many

Summary

Concept	Description
Data Modeling	Blueprint of your database structure
Normalization	Breaks data into smaller tables to reduce redundancy
1NF	Remove multivalued columns
2NF	Remove partial dependency
3NF	Remove transitive dependency
ER Diagram	Visual map of tables and relationships
PK & FK	Connect tables with unique references

◆ Performance Tuning in PL/SQL and SQL

- Learn how to make your queries and PL/SQL code run faster, avoid unnecessary processing, and optimize database usage
- This is very important in real-world projects like banking, insurance, and ecommerce apps.

❖ What is Performance Tuning?

- Performance tuning is the process of making your SQL queries or PL/SQL code work faster and use fewer resources like CPU, memory, and disk I/O.
- A poorly written query can take minutes or hours, while an optimized query does the same work in seconds.
- In real-time applications, slow queries affect customers, dashboards, reports, and even crash systems.
- In Simple Words:
 - Performance tuning means making your SQL/PLSQL code faster and lighter.
 - Just like a car needs tuning for better mileage, your database queries need tuning to avoid slow performance.
 - It saves time, cost, and improves user experience in real-time systems like mobile apps, dashboards, or financial systems.

◆ Major Areas of PL/SQL and SQL Performance Tuning

Area	Description
Indexes	Help speed up WHERE conditions
Execution Plan	Shows how Oracle executes your query
Bulk Processing	Avoids row-by-row (slow) operations
Bind Variables	Prevent parsing and reusing queries
Avoiding Cursors or Loops when possible	Use set-based logic
Limiting SELECT	Avoid SELECT * and unnecessary joins
Proper Exception Handling	Prevent errors and improve control

❖ 1. Use Indexes Properly

◆ What is an Index?

- It's like an address book. If you need to search for **Vaibhav** in a phonebook, it's quicker to jump to "V" than checking every page.

💻 Syntax to Create Index:

```
CREATE INDEX idx_emp_name ON employees(emp_name);
```

💡 Example:

```
-- Before index  
SELECT * FROM employees WHERE emp_name = 'Vaibhav';  
-- After the index is created on emp_name, it will be much faster.  
👉 Real-time Use: Search features in apps (e.g., employee lookup, order status check) use indexes for fast results.
```

- ◊ 2. Check and Understand Execution Plan
 - ◆ What is Execution Plan?
 - It shows how Oracle executes your SQL — whether it uses full table scan, index scan, nested loops etc.

- ◆ How to check:

```
EXPLAIN PLAN FOR
SELECT * FROM employees WHERE emp_id = 100;
```

```
SELECT * FROM TABLE(DBMS_XPLAN.DISPLAY);
```

📌 Use this in performance tuning to identify bottlenecks.

- ◊ 3. Avoid Using Loops for Row-by-Row Processing

- ◆ Problem:

```
-- BAD PRACTICE
FOR rec IN (SELECT * FROM employees) LOOP
  UPDATE employees SET salary = salary * 1.1 WHERE emp_id = rec.emp_id;
END LOOP;
```

This is very slow when thousands of records are involved.

- ◆ Better Option (SET-BASED Operation):

```
-- GOOD PRACTICE
```

```
UPDATE employees SET salary = salary * 1.1;
```

📌 Real-time: Use this in salary hikes, tax updates, mass record changes.

- ◊ 4. Use BULK COLLECT and FORALL

- ◆ Problem with Normal Loop:

```
FOR rec IN (SELECT * FROM emp) LOOP
  INSERT INTO emp_log VALUES (rec.emp_id, rec.name);
END LOOP;
```

It inserts one row at a time — very slow.

- ◆ Optimized Version:

```
DECLARE
  TYPE emp_table IS TABLE OF emp%ROWTYPE;
  l_data emp_table;
BEGIN
  SELECT * BULK COLLECT INTO l_data FROM emp;
```

```
FORALL i IN 1 .. l_data.COUNT
  INSERT INTO emp_log VALUES (l_data(i).emp_id, l_data(i).name);
END;
```

📌 This is 100x faster for inserting/updating/deleting large records.

◊ 5. Use Bind Variables

◆ Problem:

-- BAD: Hard-coded value forces query parsing every time
`SELECT * FROM employees WHERE emp_id = 101;`

◆ Better:

-- GOOD: Bind variable (used in apps or PL/SQL blocks)
`SELECT * FROM employees WHERE emp_id = :emp_id;`

❖ Saves parsing time and memory in apps like Oracle Forms, APEX, Java apps.

◊ 6. Avoid SELECT * — Always Use Required Columns

◆ Problem:

`SELECT * FROM employees;`

This fetches unnecessary columns, wastes time and memory.

◆ Better:

`SELECT emp_id, emp_name FROM employees;`

❖ Especially useful in reports, dashboards where speed matters.

◊ 7. Avoid Too Many Joins/Subqueries

◆ Problem:

`SELECT * FROM emp WHERE dept_id IN (SELECT dept_id FROM dept WHERE location = 'Pune');`

◆ Better (using JOIN):

`SELECT e.emp_id, e.name
FROM emp e
JOIN dept d ON e.dept_id = d.dept_id
WHERE d.location = 'Pune';`

❖ Joins are faster and better optimized by the Oracle optimizer.

◊ 8. Analyze and Gather Statistics

- Oracle uses table statistics to decide execution plans.

```
BEGIN  
  DBMS_STATS.GATHER_TABLE_STATS('HR', 'EMPLOYEES');  
END;
```

❖ Use it after major updates/inserts to keep queries fast.

◊ 9. Use Parallel Query for Large Data

`SELECT /*+ parallel(4) */ * FROM big_sales_data;`

❖ Divides the work into 4 threads, speeds up large data processing.

💬 Real-World Scenario: Tuning a Report Query

◆ Problem:

Monthly report query is taking 10 minutes. You observe:

- It uses `SELECT *`
- It scans full 10 million row table
- Joins are written with subqueries

- ◆ Fix:
 - Add index on filtering column (e.g. `sale_date`)
 - Replace `SELECT *` with only needed columns
 - Convert subqueries to JOINs
 - Gather stats after index creation

Summary: Performance Tuning

Technique	What it does	When to use
Indexes	Speeds up search/filter	WHERE clause on big table
Execution Plan	Shows how query runs	Analyze slow queries
BULK COLLECT/FORALL	Fast batch processing	INSERT/UPDATE many rows
Bind Variables	Avoid re-parsing	Dynamic SQL in apps
Avoid <code>SELECT*</code>	Fetch only needed data	Reports, dashboards
Set-based Logic	Avoid row-by-row	Mass updates or reports
Join Optimization	Better than subqueries	Large table joins
Stats Gathering	Helps optimizer	After data changes
Parallel Query	Speeds up large queries	Millions of rows

◆ PL/SQL Metadata Tables

- Metadata tables (also called data dictionary views) are system-generated views in Oracle that help you explore the structure and status of your own objects — like tables, columns, procedures, triggers, views, etc.
- Think of it like:
 - ▀ "Instead of opening SQL Developer and right-clicking on a table, just run a query on these views and get all the hidden details in seconds!"

◊ 1. `USER_TABLES`

- ◆ Purpose:
 - Shows the list of all tables owned by the current user.
 - It doesn't show columns, only table-level details.
- ◆ Key Columns:
 - `TABLE_NAME`: Name of your table
 - `TABLESPACE_NAME`: Where the table is stored
 - `NUM_ROWS`: Number of rows (based on statistics)
 - `LAST_ANALYZED`: When stats were last collected

💡 Example:

```
SELECT table_name, tablespace_name, num_rows, last_analyzed
FROM USER_TABLES;
```

☰ Output:

TABLE_NAME	TABLESPACE_NAME	NUM_ROWS	LAST_ANALYZED
EMPLOYEES	USERS	2000	17-JUN-2025 10:30am
DEPARTMENTS	USERS	12	16-JUN-2025 09:00am

🧠 Real-Time Use Case:

- You're building a report on employee data, and need to confirm if the table is created, where it's stored, and how large it is.

◊ 2. **USER_TAB_COLUMNS**

- ◆ Purpose:
- Shows the structure of your tables — i.e., column names, data types, and nullability.
- ◆ Key Columns:
 - TABLE_NAME
 - COLUMN_NAME
 - DATA_TYPE
 - DATA_LENGTH
 - NULLABLE

💡 Example:

```
SELECT table_name, column_name, data_type, data_length, nullable  
FROM USER_TAB_COLUMNS  
WHERE table_name = 'EMPLOYEES';
```

☰ Output:

TABLE_NAME	COLUMN_NAME	DATA_TYPE	DATA_LENGTH	NULLABLE
EMPLOYEES	EMP_ID	NUMBER	22	N
EMPLOYEES	NAME	VARCHAR2	50	Y
EMPLOYEES	SALARY	NUMBER	22	Y

🧠 Real-Time Use Case:

- You are writing a PL/SQL procedure to insert data, and want to check which columns are NOT NULL or what data types they are.

◊ 3. **ALL_OBJECTS**

- ◆ Purpose:
- Gives info about all objects (tables, views, procedures, packages, triggers, etc.) that you can access — both yours and those granted by others.
- ◆ Key Columns:
 - OBJECT_NAME: Name of the object
 - OBJECT_TYPE: TABLE / PROCEDURE / PACKAGE / TRIGGER / VIEW
 - STATUS: VALID / INVALID
 - CREATED, LAST_DDL_TIME: Created or modified date

💡 Example:

```
SELECT object_name, object_type, status, created
FROM ALL_OBJECTS
WHERE object_type IN ('PROCEDURE', 'PACKAGE');
```

💻 Output:

OBJECT_NAME	OBJECT_TYPE	STATUS	CREATED
EMP_PKG	PACKAGE	VALID	16-JUN-2025 10:30
CALC_BONUS	PROCEDURE	INVALID	15-JUN-2025 12:00

🧠 Real-Time Use Case:

- You want to compile only invalid packages, or see when your procedure was last updated.

◊ 4. USER_VIEWS

◆ Purpose:

- Lists all views created by the current user.
- It even contains the entire SQL code used to build the view.

◆ Key Columns:

- **VIEW_NAME**: Name of the view
- **TEXT**: The actual SQL that defines the view
- **TEXT_LENGTH**: Length of the view SQL

💡 Example:

```
SELECT view_name, text_length
FROM USER_VIEWS;
To view the SQL behind a specific view:
SELECT text
FROM USER_VIEWS
WHERE view_name = 'EMP_SAL_VIEW';
```

💻 Output (text):

```
SELECT e.emp_id, e.name, d.dept_name
FROM employees e JOIN departments d ON e.dept_id = d.dept_id;
```

🧠 Real-Time Use Case:

- When debugging or modifying dashboards or reporting views, you want to extract and improve existing view queries.

◊ 5. USER_PROCEDURES & USER_SOURCE

- **USER_PROCEDURES**: Lists all your PL/SQL procedures, functions, and packages.
- **USER_SOURCE**: Gives source code (line-by-line) of those objects.

💡 Example 1: Get List of PL/SQL Units

```
SELECT object_name, procedure_name, object_type, status  
FROM USER_PROCEDURES;
```

OBJECT_NAME	PROCEDURE_NAME	OBJECT_TYPE	STATUS
EMP_PKG	CALC_SALARY	PACKAGE BODY	VALID

💡 Example 2: View Code of Procedure

```
SELECT line, text  
FROM USER_SOURCE  
WHERE name = 'EMP_PKG'  
ORDER BY line
```

LINE	TEXT
1	PROCEDURE calc_salary IS
2	BEGIN
3	-- logic to calculate salary
4	END;

🧠 Real-Time Use Case:

- You want to download, edit, or version control your procedures and packages, especially before a release or code migration.

◊ 6. USER_TRIGGERS

- ◆ Purpose:
- Lists all triggers created by the user, and what events or tables they are associated with.
- ◆ Key Columns:
 - **TRIGGER_NAME**
 - **TABLE_NAME**
 - **TRIGGERING_EVENT** (INSERT / UPDATE / DELETE)
 - **STATUS**
 - **DESCRIPTION**

 Example:

```
SELECT trigger_name, table_name, triggering_event, status  
FROM USER_TRIGGERS;
```

TRIGGER_NAME	TABLE_NAME	TRIGGERING_EVENT	STATUS
TRG_LOG_ORD	ORDERS	INSERT	ENABLED

 Real-Time Use Case:

- You are debugging an audit trigger, or checking why insert is failing — and you want to confirm if any trigger is fired on that table.

◊ 7. USER_SEQUENCES

- ◆ Purpose:
- Lists all sequences (auto-increment counters) created by you.
- ◆ Key Columns:
 - SEQUENCE_NAME
 - MIN_VALUE, MAX_VALUE
 - INCREMENT_BY
 - LAST_NUMBER (next value to be generated)

 Example:

```
SELECT sequence_name, last_number, increment_by  
FROM USER_SEQUENCES;
```

SEQUENCE_NAME	LAST_NUMBER	INCREMENT_BY
EMP_SEQ	101	1

 Real-Time Use Case:

- You are assigning IDs during insert operations and want to check the current value of the sequence or reset it.

 In Real-Time Projects

- Here's how developers use metadata tables daily:
 - Data Engineers check `USER_TAB_COLUMNS` before data load
 - PL/SQL developers monitor `USER_SOURCE` before deployments
 - Report developers fix `USER_VIEWS` if dashboard fails
 - ETL teams use `ALL_OBJECTS` to recompile invalid packages

Summary : Full Comparison

Metadata View	Description	Best Used For
<code>USER_TABLES</code>	Lists all user-created tables	Check if a table exists, row stats
<code>USER_TAB_COLUMNS</code>	Shows all columns of your tables	Profile table design, null checks
<code>ALL_OBJECTS</code>	All objects (tables, packages, etc.)	Check status, created date
<code>USER_VIEWS</code>	Views and their SQL definition	Debug or edit complex views
<code>USER_SOURCE</code>	Full PL/SQL source code (line by line)	Extract code, backup, versioning
<code>USER_PROCEDURES</code>	List of PL/SQL procedures/functions	See if code is compiled or valid
<code>USER_TRIGGERS</code>	Lists all triggers created by user	Debug inserts, updates, deletes
<code>USER_SEQUENCES</code>	Shows all sequences and values	ID generation and tracking

◆ PL/SQL Code Reusability Best Practices

- ◊ 1. Use Packages for Grouping Logic
 - ◆ What is it?
 - A package is a collection of related procedures, functions, variables, cursors grouped together under a single name.
 - ◆ Why?
 - Promotes modular programming
 - Makes code easier to reuse and maintain
 - Supports encapsulation (public vs private procedures)
 - Improves performance (stored in memory)

▀ Syntax Example:

```
-- Package Specification
CREATE OR REPLACE PACKAGE emp_pkg IS
    PROCEDURE add_employee(p_id NUMBER, p_name VARCHAR2);
    FUNCTION get_employee(p_id NUMBER) RETURN VARCHAR2;
END emp_pkg;
/

-- Package Body
CREATE OR REPLACE PACKAGE BODY emp_pkg IS
    PROCEDURE add_employee(p_id NUMBER, p_name VARCHAR2) IS
        BEGIN
            INSERT INTO employees (emp_id, emp_name) VALUES (p_id, p_name);
        END;

    FUNCTION get_employee(p_id NUMBER) RETURN VARCHAR2 IS
        v_name VARCHAR2(100);
        BEGIN
            SELECT emp_name INTO v_name FROM employees WHERE emp_id = p_id;
            RETURN v_name;
        END;
END emp_pkg;
/
```

◆ How to reuse:

```
BEGIN
    emp_pkg.add_employee(101, 'Vaibhav');
    DBMS_OUTPUT.PUT_LINE(emp_pkg.get_employee(101));
END;
```

◊ 2. Use Procedures & Functions Instead of Repeating Code

- ◆ When to use:
 - Procedure → perform an action (insert, update, delete)
 - Function → return a value (like a calculated result)

💡 Example

```
CREATE OR REPLACE FUNCTION calc_bonus (p_salary NUMBER)
RETURN NUMBER IS
BEGIN
```

```
    RETURN p_salary * 0.10;
END;
/
```

- ◆ Usage:
`SELECT emp_name, calc_bonus(salary) FROM employees;`
- ◆ Instead of writing the bonus logic again and again, use this function in reports or triggers.

◊ 3. Use Constants and Global Variables

- ◆ Why?
- Avoids hardcoding values
- Easier to update from one place
- Promotes maintainability

💡 Example

```
CREATE OR REPLACE PACKAGE constants_pkg IS
    bonus_rate CONSTANT NUMBER := 0.10;
END;
```

Use in other procedures:

```
sal * constants_pkg.bonus_rate;
```

◊ 4. Use Cursors & Cursor Variables

- Use explicit cursors or cursor FOR loops to reuse select logic, especially when looping through data in reports or batch jobs.

💡 Example

```
CURSOR emp_cur IS
    SELECT emp_id, salary FROM employees;
```

```
FOR rec IN emp_cur LOOP
    -- Reuse logic for all employees
END LOOP;
```

◊ 5. Write Generic & Parameterized Code

- ◆ Why?
- Helps reuse the same code for different values
- Easy to plug and play in different modules

💡 Example

```
CREATE OR REPLACE PROCEDURE update_salary (p_id NUMBER, p_new_sal
NUMBER) IS
BEGIN
    UPDATE employees SET salary = p_new_sal WHERE emp_id = p_id;
END;
```

Use this in multiple reports, job schedulers, forms, etc.

◊ 6. Use Exception Blocks for Reusability

- Handle errors once and reuse the same standard error format across your app.

 Example

```
BEGIN
    -- logic
EXCEPTION
    WHEN OTHERS THEN
        log_error(SQLERRM); -- Reusable error logger
END;
```

Create a reusable error logger procedure:

```
CREATE OR REPLACE PROCEDURE log_error(p_msg VARCHAR2) IS
BEGIN
    INSERT INTO error_log(error_message, log_time)
    VALUES (p_msg, SYSDATE);
END;
```

◊ 7. Document and Comment the Code

- A best practice is to:
 - Add comments explaining inputs/outputs
 - Describe business logic
 - Mention version info or authors
- This makes code easier to reuse and maintain.

 Example

```
-- Procedure: update_salary
-- Author: Vaibhav D
-- Purpose: To update salary of employee
```

◊ 8. Avoid Copy-Pasting Logic Everywhere

- Instead of copying insert/update logic in multiple places:
 - Put it inside a procedure/package
 - Call it wherever needed
- This reduces bugs and inconsistencies

◊ 9. Use Modular Approach

- Break large code blocks into:
 - Helper procedures
 - Utility functions
 - Packages with private/public methods
- Modular code = clean, organized, testable

◊ 10. Plan for Future Use

- Always write PL/SQL with future use in mind:
 - Parameterize values
 - Avoid hard-coded table/column names
 - Keep logic separate from presentation

Summary Table

Best Practice	Benefit
Use Packages	Group logic, improve performance
Write Procedures/Functions	Avoid code repetition
Use Constants	Maintain values from one place
Use Cursors	Iterate reusable queries
Generic Parameters	Increase flexibility
Handle Exceptions Reusably	Central error management
Comment and Document Code	Easier to reuse and debug later
Avoid Hardcoding	Cleaner and scalable code
Modular Design	Easy to test and maintain